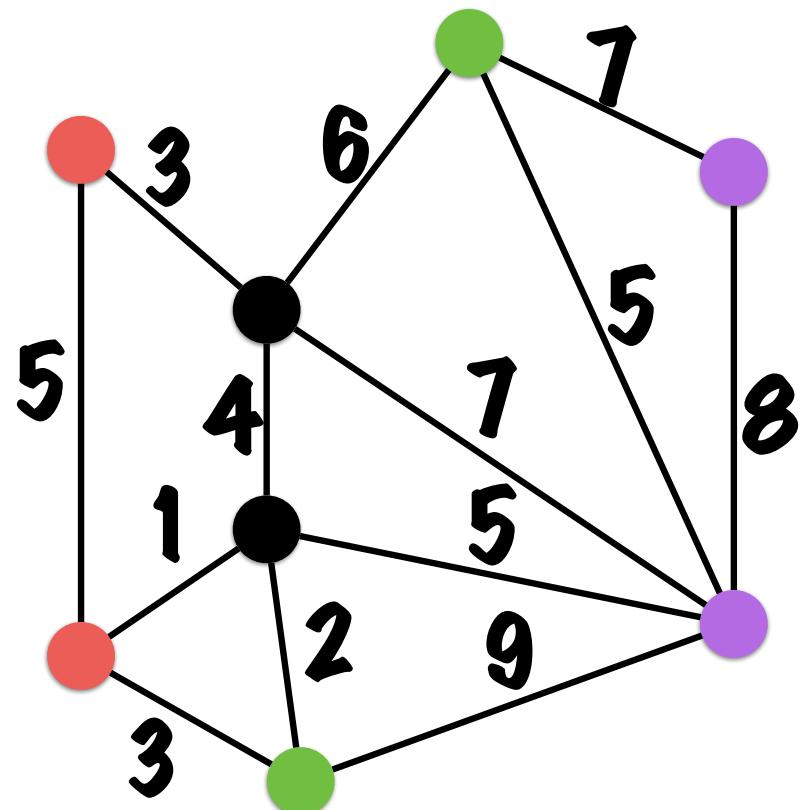


Steiner forest

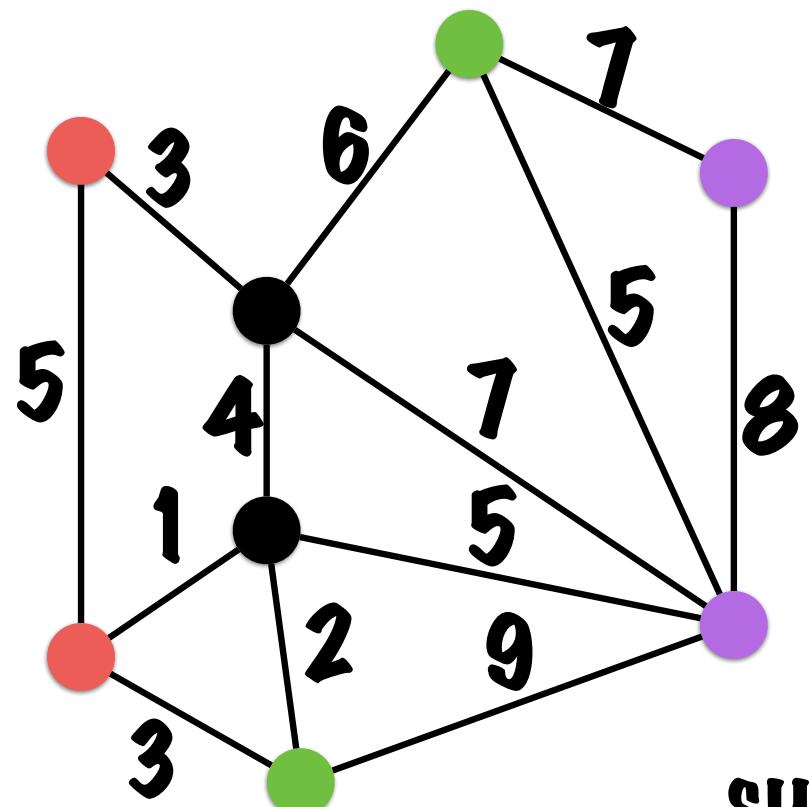


Appears in...



VLSI
**Optical and wireless
communication**
**Transportation and
distribution systems**
Network design

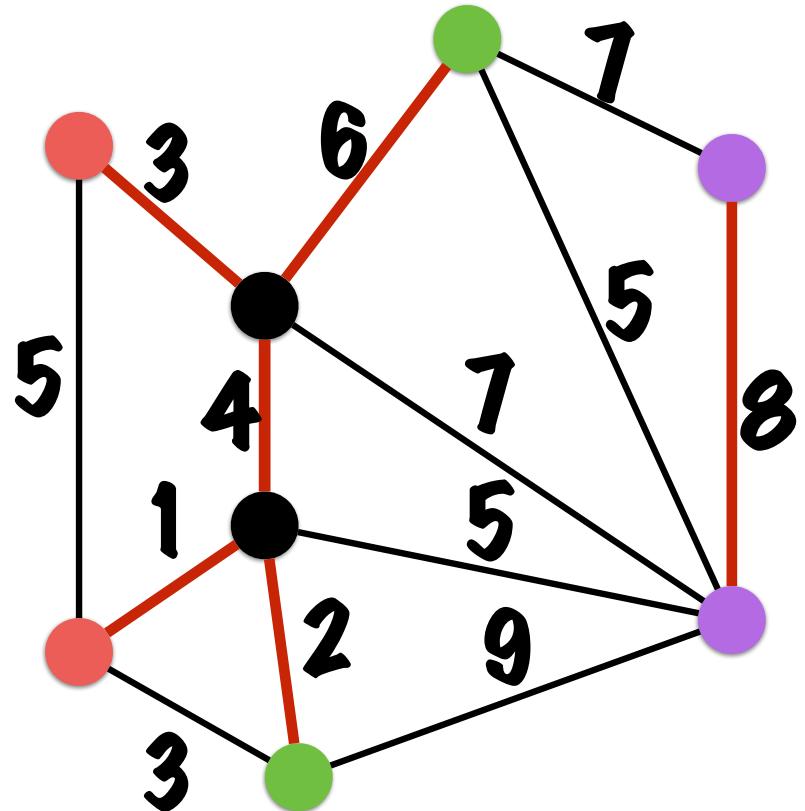
Connecting subsets



Input
Graph with edge costs
subsets of vertices (terminals)
 $\{\text{red, red}\}, \{\text{green, green}\}, \{\text{purple, purple}\}$

Connecting subsets

Input

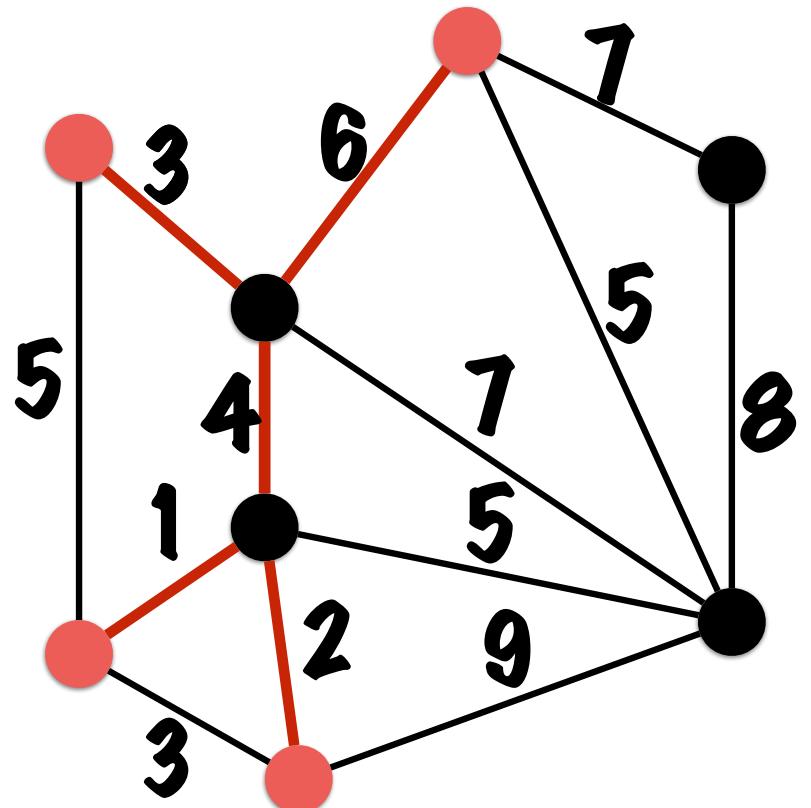


Graph with edge costs
subsets of terminals
 $\{\text{Red}, \text{Red}\}, \{\text{Green}, \text{Green}\}, \{\text{Purple}, \text{Purple}\}$

Output
choose edges s.t.
each subset is connected

Goal
minimize cost of edges

A special case: Steiner tree



Input

Graph with edge costs
one subset of terminals

{ 1, 2, 3, 4 }

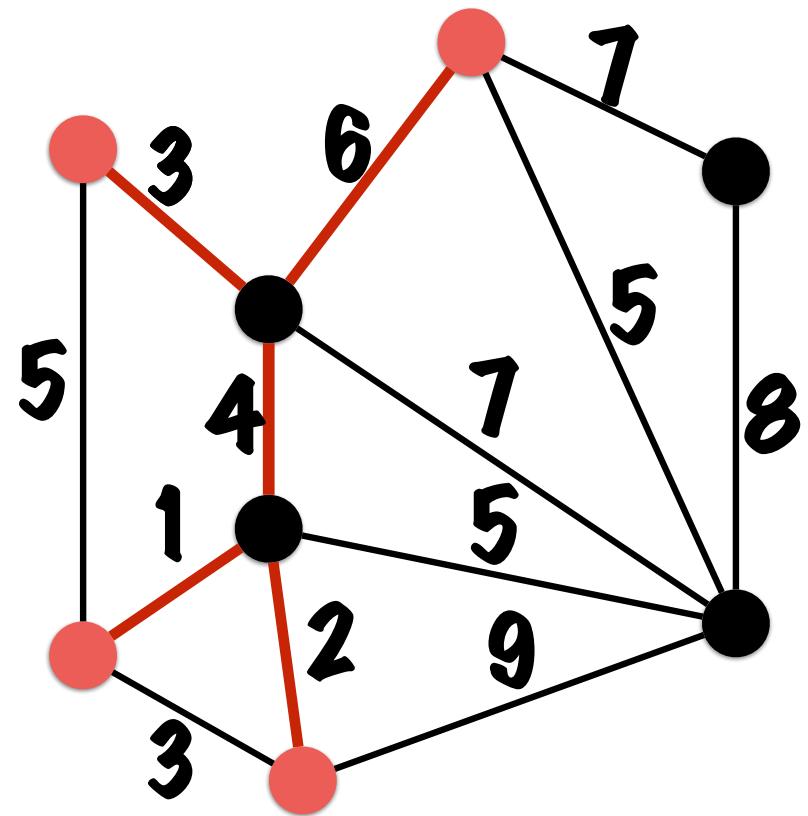
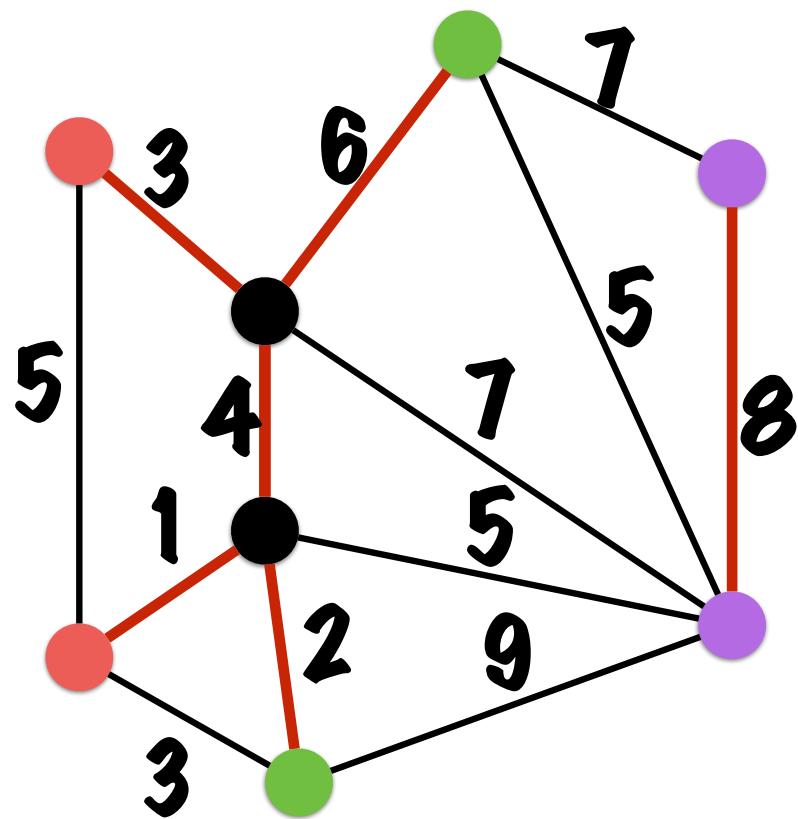
Output

choose edges s.t.
the subset is connected

Goal

minimize cost of edges

Steiner forest vs. Steiner tree



Steiner forest



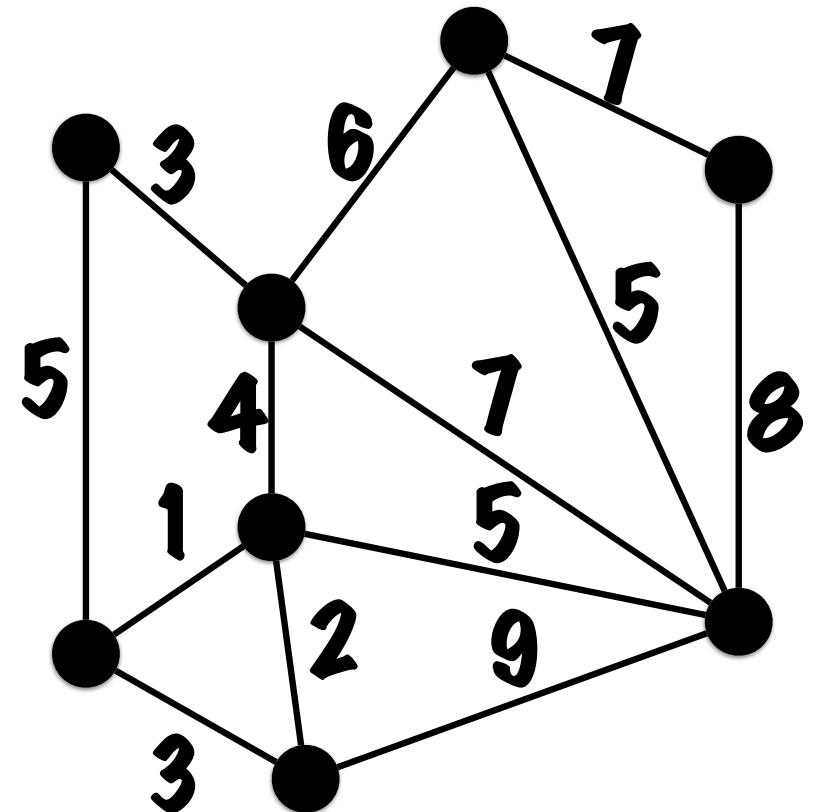
Steiner forest



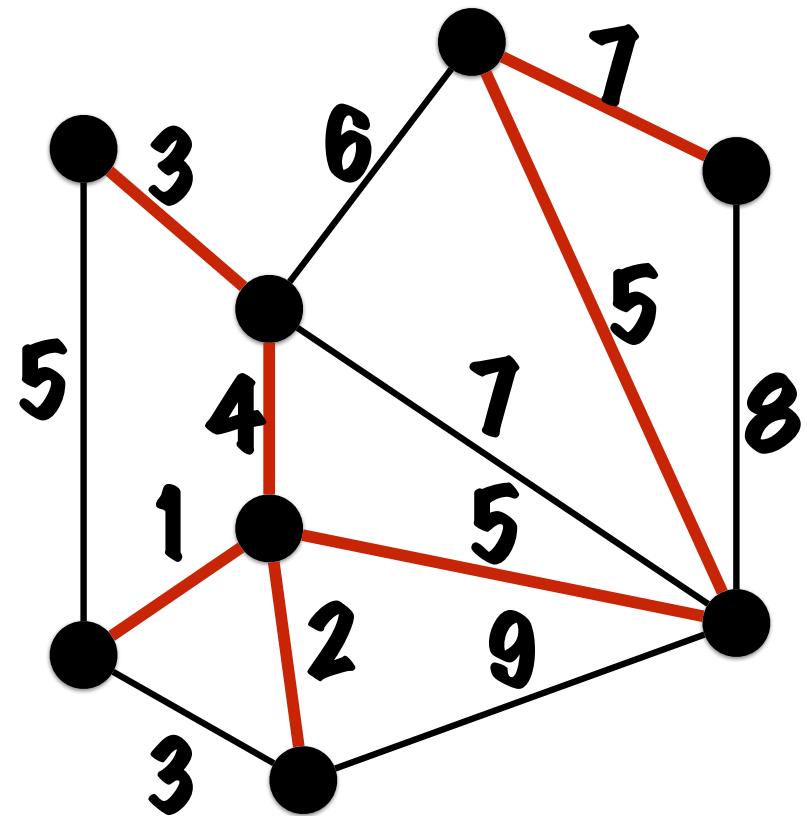
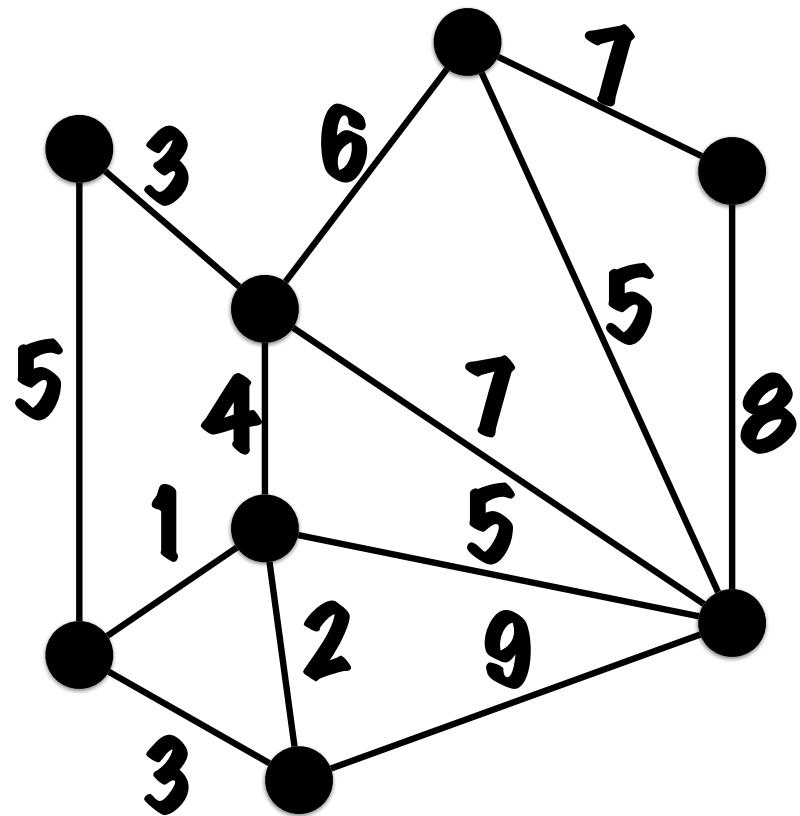
Steiner tree

Observe:

if the subset = all nodes
then Steiner tree =
minimum spanning tree



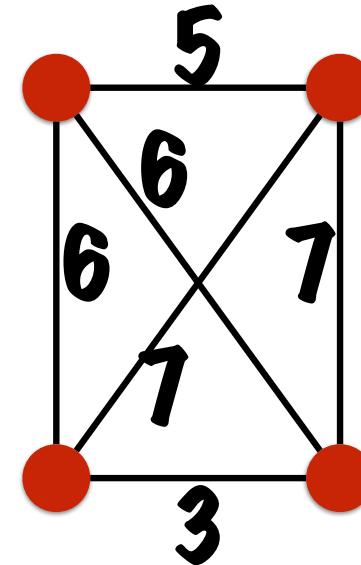
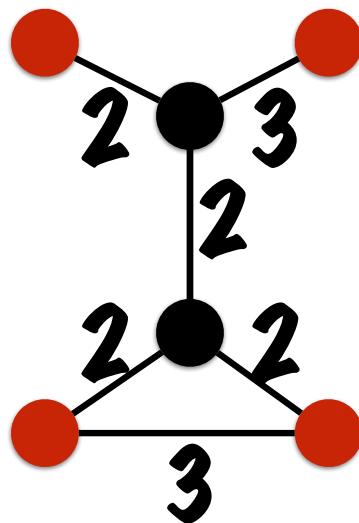
MST



Steiner tree approx. algorithm

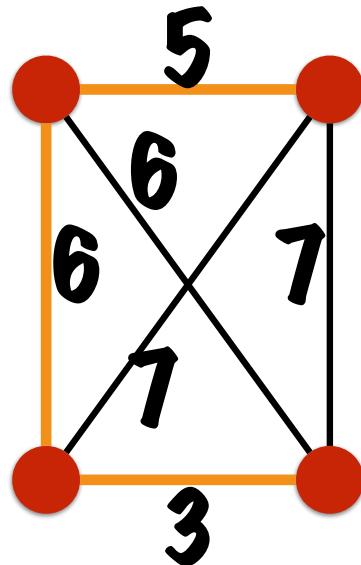
1. Define new graph:

- Vertex set = S
- Edge set = complete graph
- Weight of $\{u,v\}$ = shortest path length in G



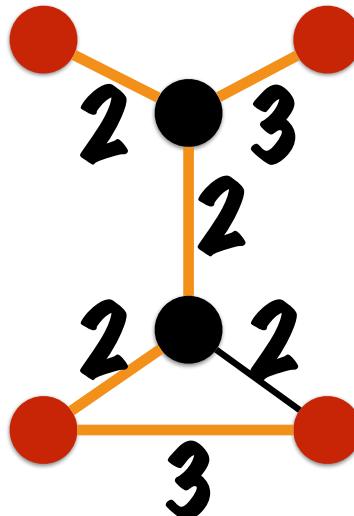
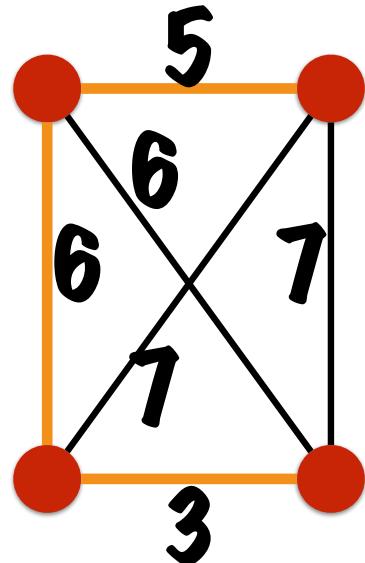
Steiner tree algorithm

2. Compute minimum spanning tree on new graph



Steiner tree algorithm

3. Output corresponding set of original edges



Here
Output=12
OPT=11

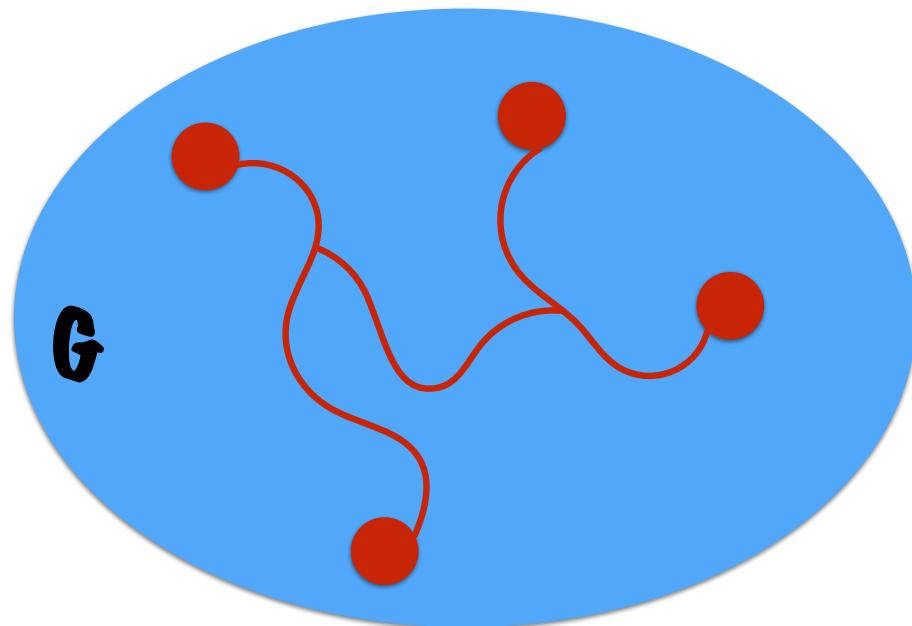
Steiner tree algorithm

- 1. new complete graph on subset**
- 2. minimum spanning tree on that graph**
- 3. output corresponding set of original edges**

Theorem: it's a 2-approximation

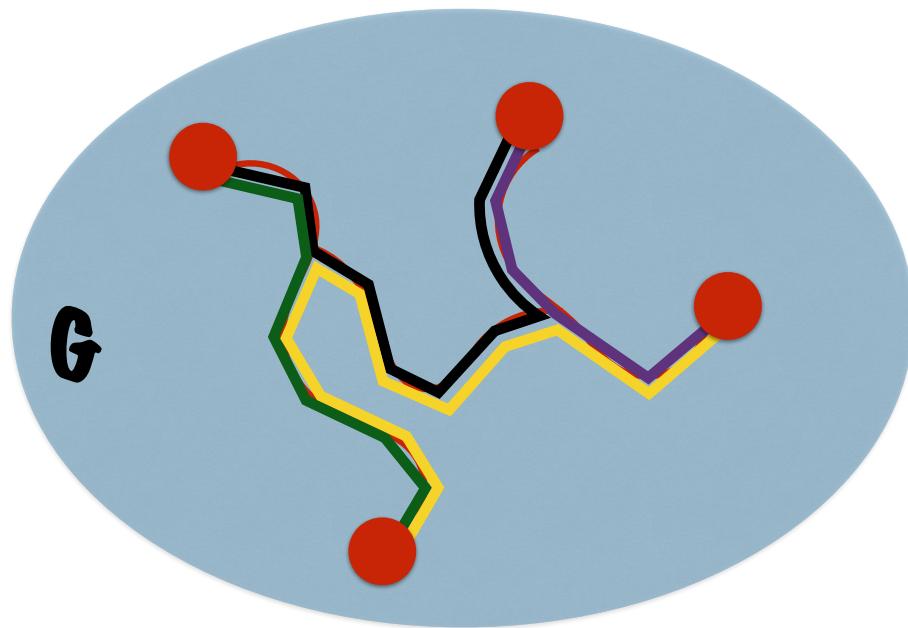
Theorem: it's a 2-approximation

Consider unknown OPT tree



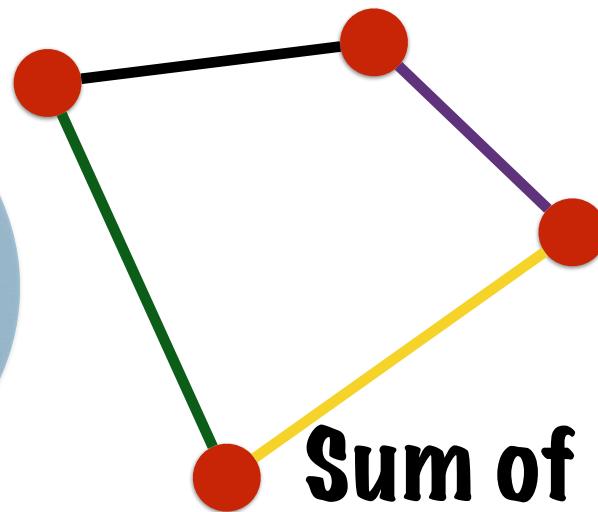
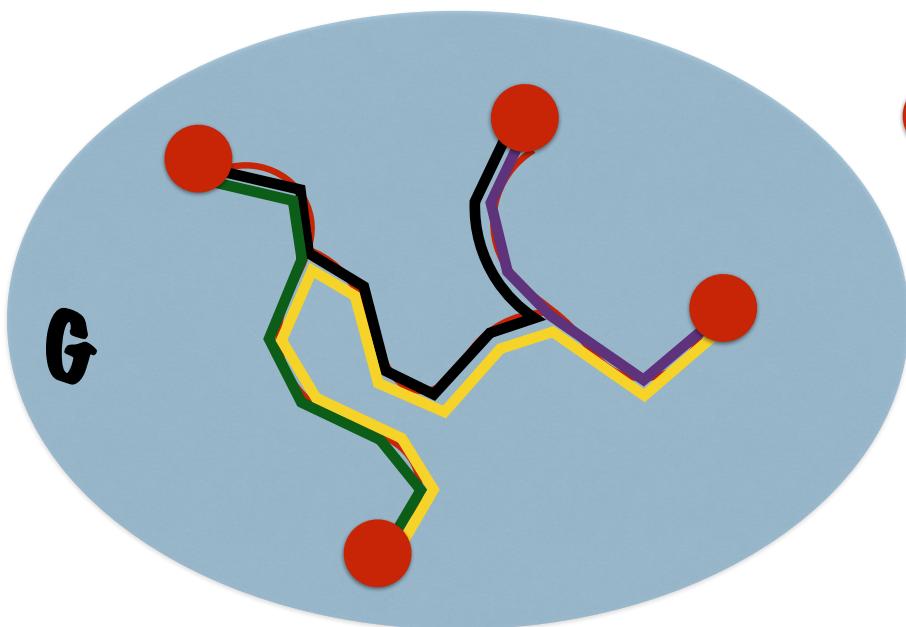
Theorem: it's a 2-approximation

Go around the OPT tree
to define terminal-to-terminal paths



Theorem: it's a 2-approximation

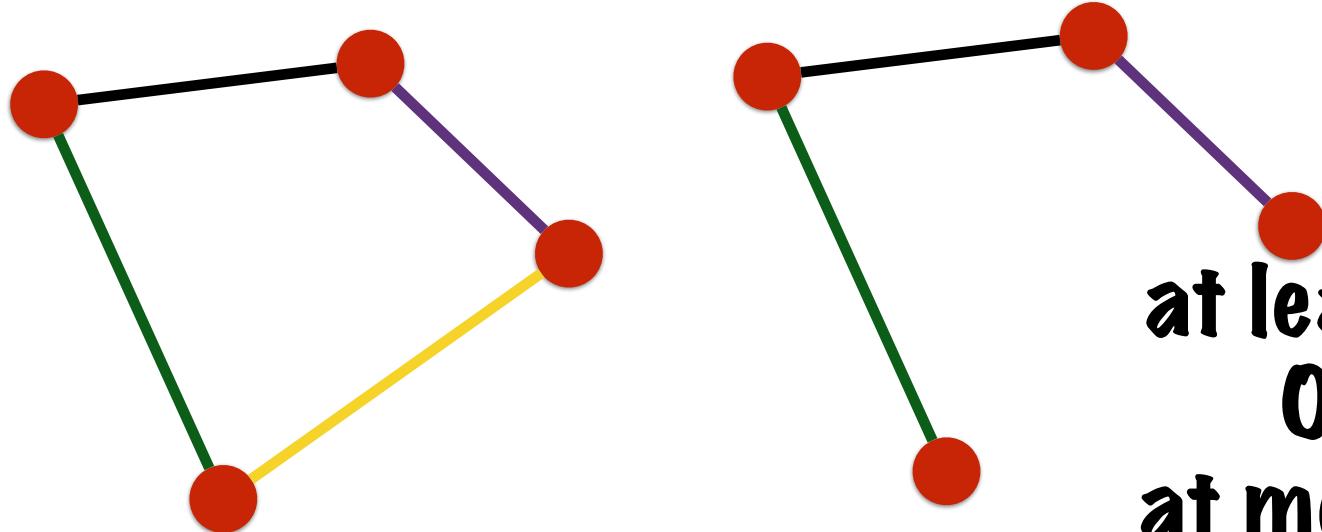
Map terminal-to-terminal paths
to edges in complete graph of terminals



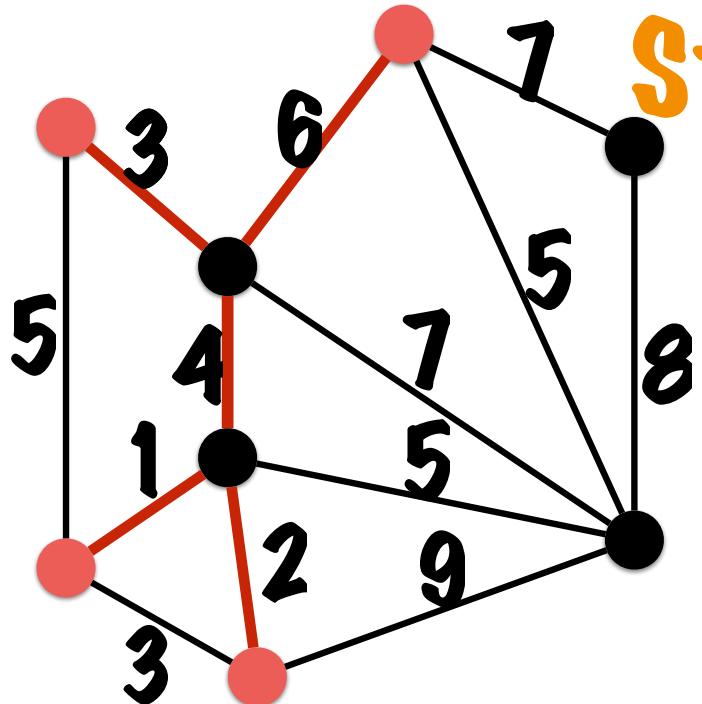
**Sum of path lengths
= 2 OPT**

Theorem: it's a 2-approximation

Remove one cycle edge



2OPT is
at least MST cost
Output cost is
at most MST cost
QED



Steiner tree



**Jakob
Steiner**



Henry Pollak

H.N. Gilbert

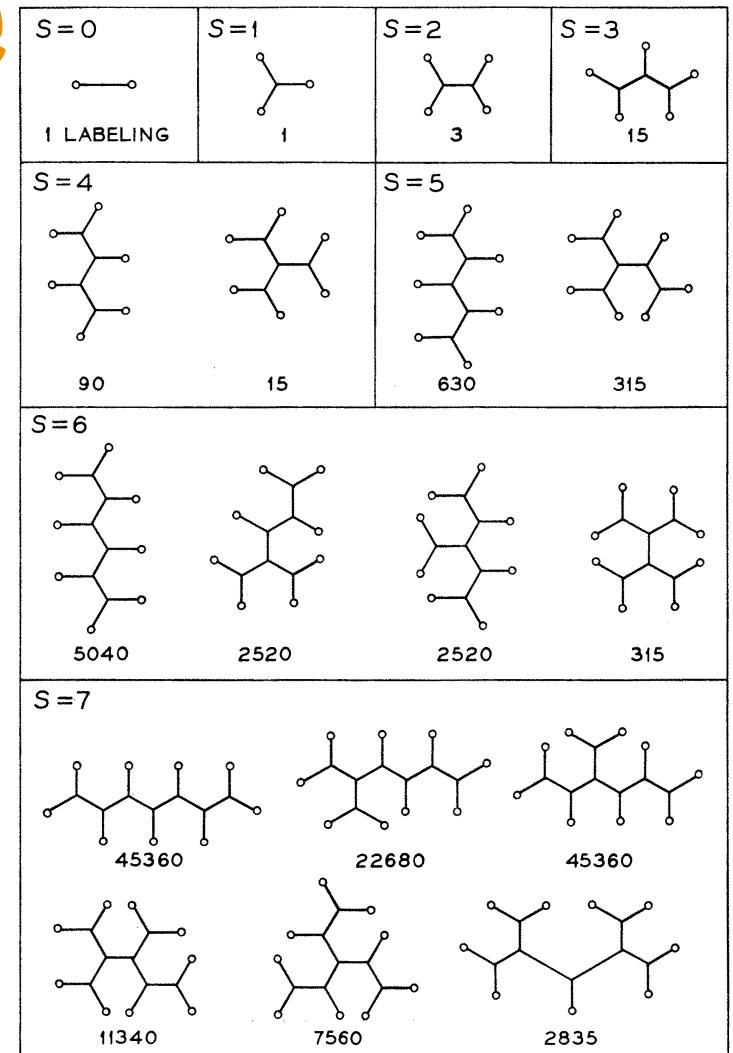
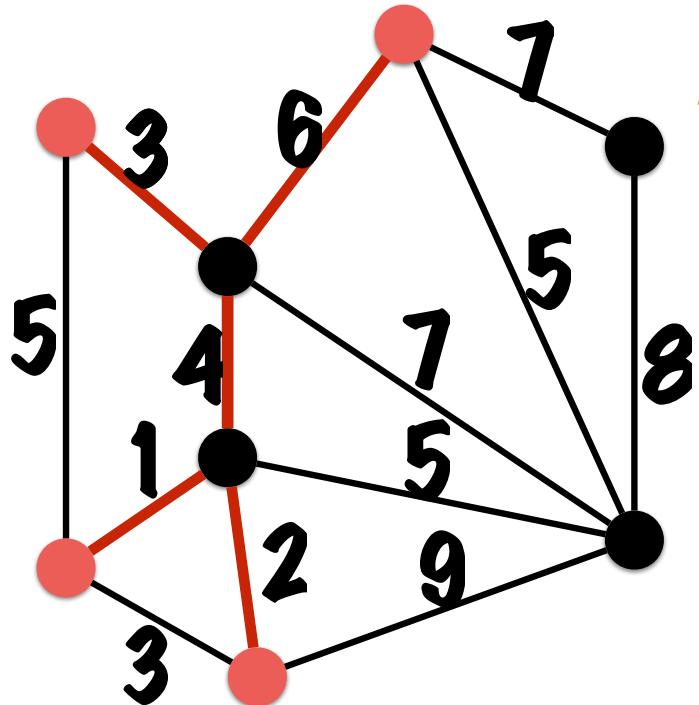
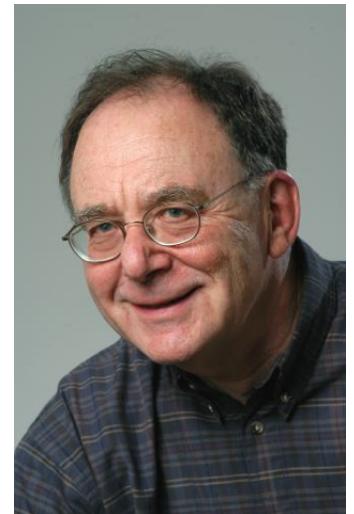


FIG. 3. Full Steiner trees with s Steiner points



Steiner tree

Richard
Karp



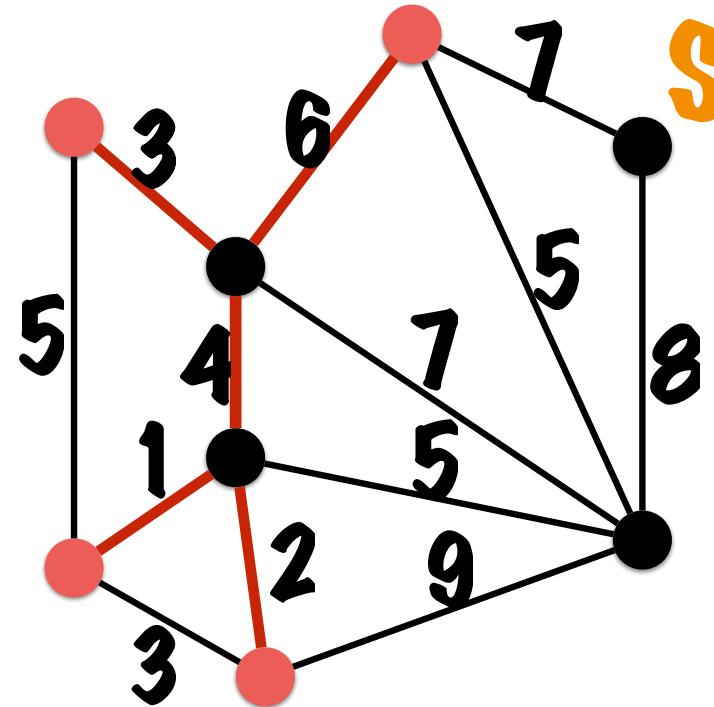
Jakob
Steiner



Marshall Bern



Paul Plassmann



Steiner tree 1.39

**Jaroslaw
Byrka**



**Fabrizio
Grandoni**



Thomas Rothvoss

Steiner forest

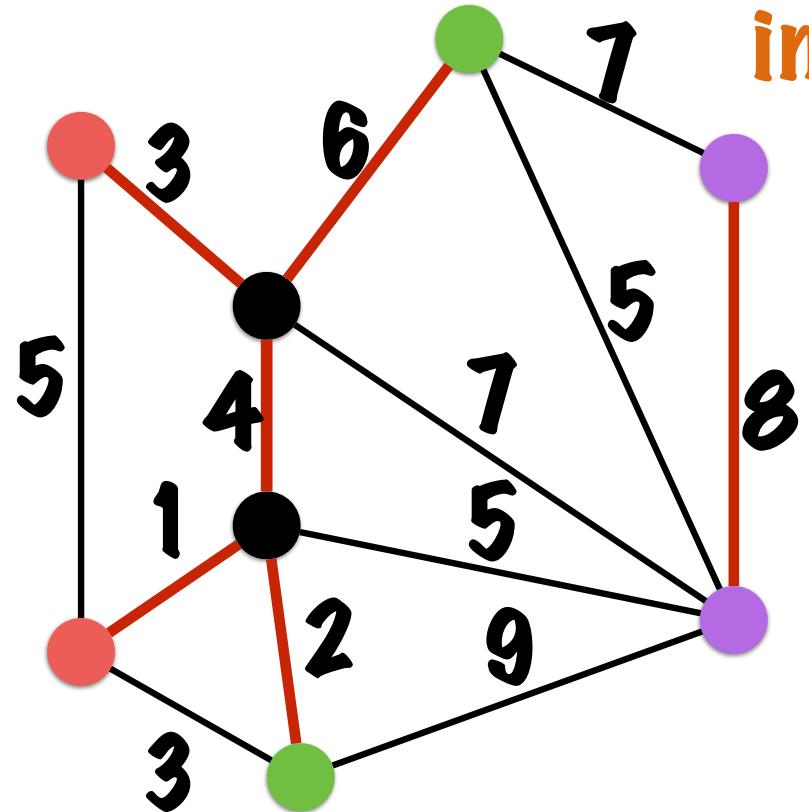


Steiner forest



Integer programming model

input: sets of terminals S_1, S_2, \dots
output: set of edges



min output cost
s.t.

for all i
for all u, v in S_i
u and v are connected

Integer programming model

output: set of edges

$$x_e \in \{0, 1\}$$

min output cost
s.t.

for all i

for all u, v in S_i

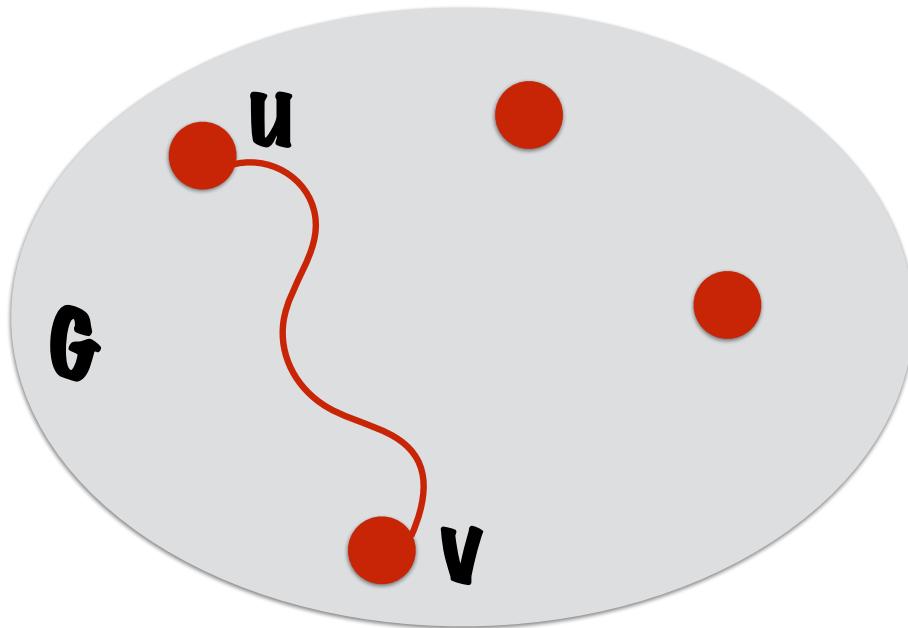
u and v are connected

$$\min \sum_e c_e x_e$$

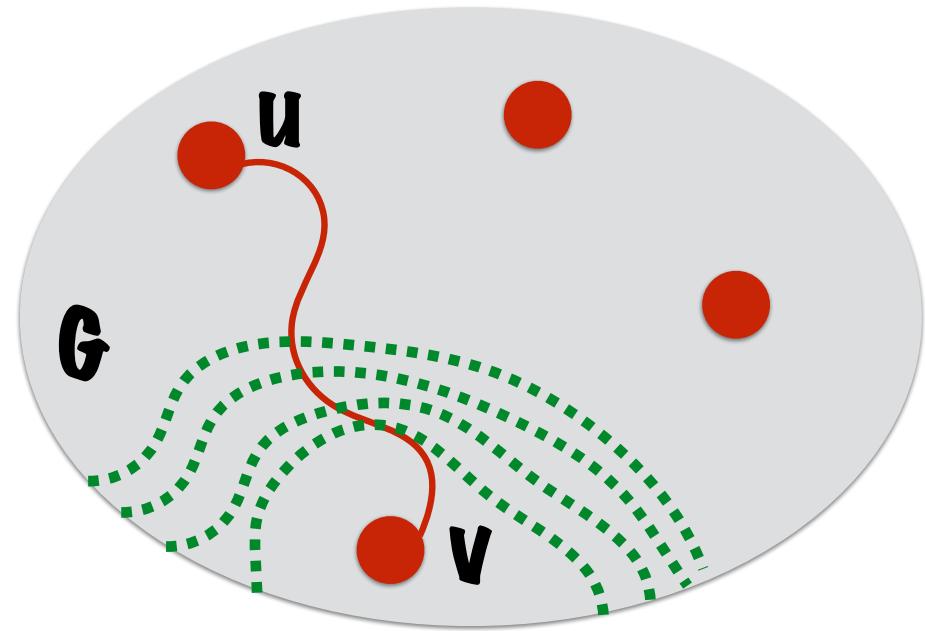
$$\forall i \forall u, v \in S_i$$

?

Integer programming model

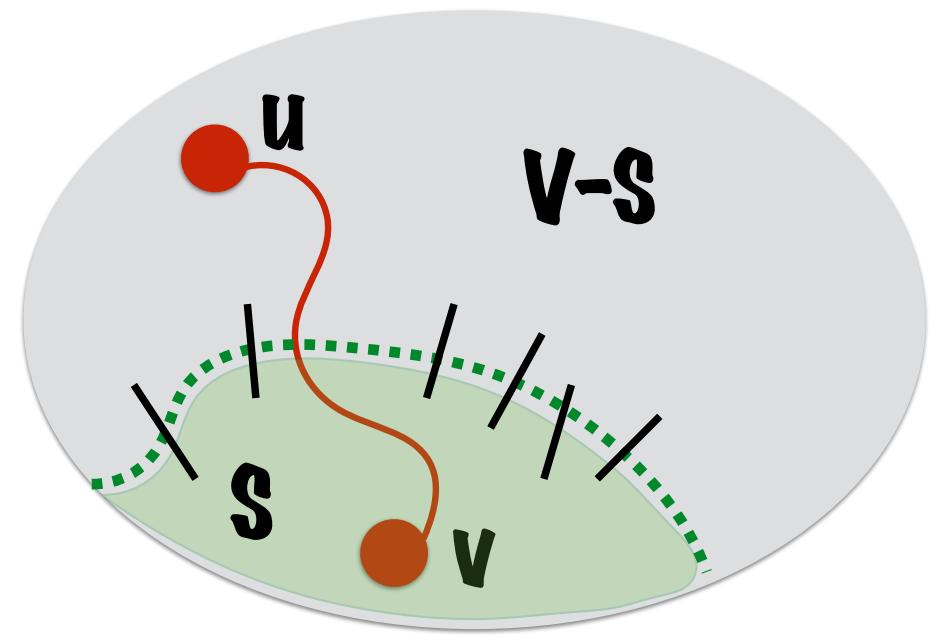
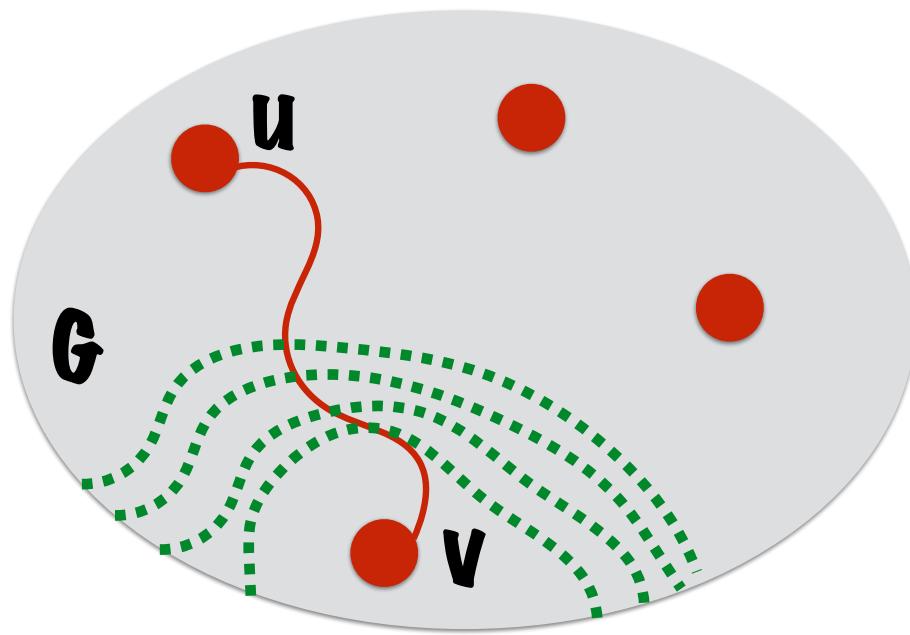


u and v are connected



every u - v cut is crossed

Integer programming model



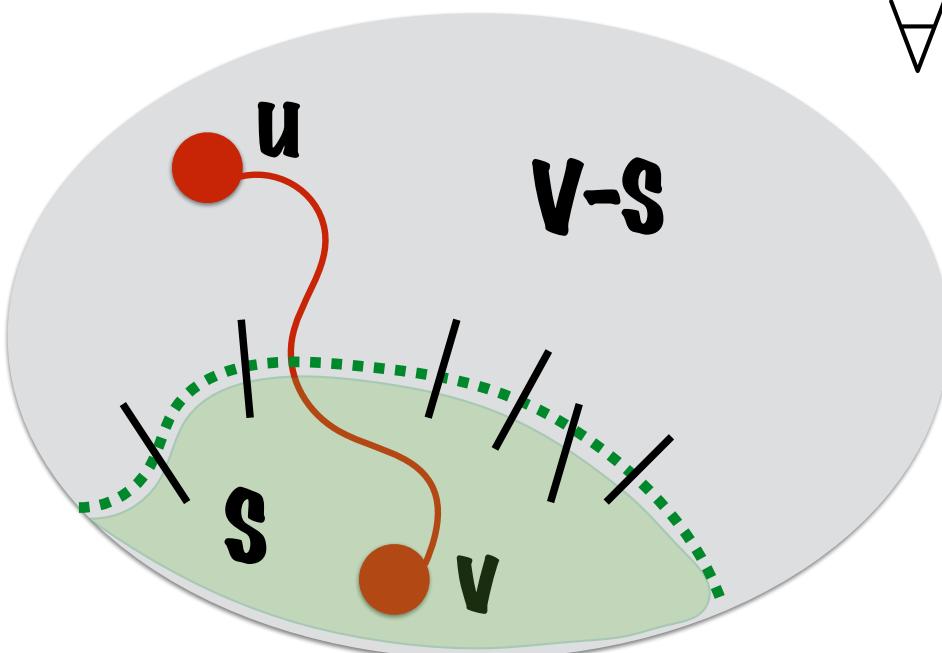
every $u-v$ cut is crossed

$$S \subseteq V : |S \cap \{u, v\}| = 1$$

$$\delta(S) = \{e \in S \times V \setminus S\}$$

for all i
for all u, v in S_i
 u and v are connected

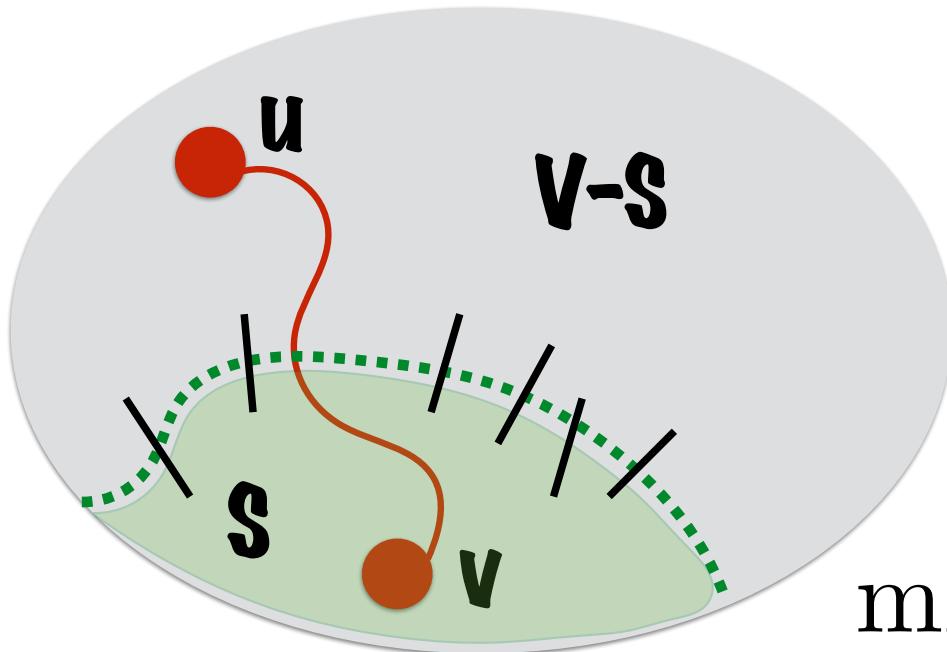
$$x_e \in \{0, 1\}$$



$$\forall i \forall u, v \in S_i \\ \forall S \text{ s.t. } |S \cap \{u, v\}| = 1$$

$$\sum_{e \in \delta(S)} x_e \geq 1$$

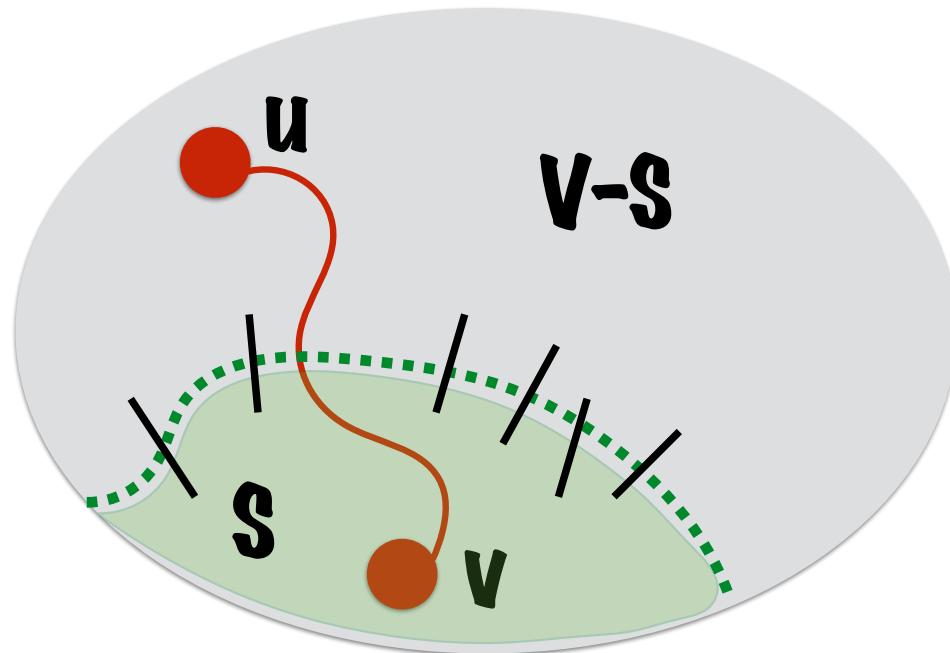
$$\mathcal{S} = \{S : \exists i \exists u, v \in S_i : |S \cap \{u, v\}| = 1\}$$



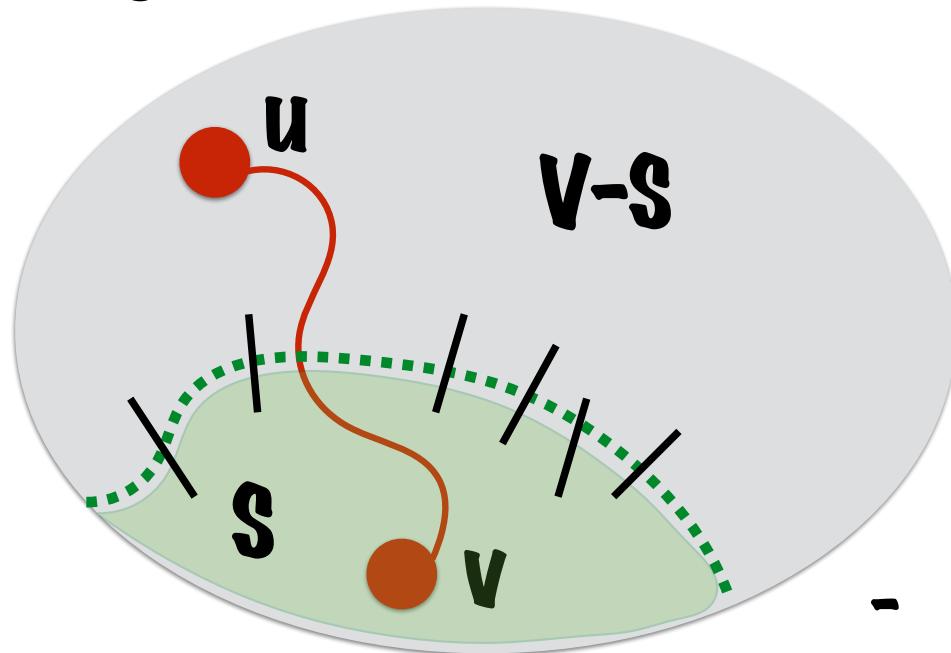
$$\begin{aligned} & \min \sum_e c_e x_e : \\ & \forall S \in \mathcal{S} \sum_{e \in \delta(S)} x_e \geq 1 \\ & x_e \in \{0, 1\} \end{aligned}$$

$$\begin{aligned} \min \sum_e c_e x_e : \\ \forall S \in \mathcal{S} \sum_{e \in \delta(S)} x_e \geq 1 \\ x_e \in \{0, 1\} \end{aligned}$$

$$\begin{aligned} \min \sum_e c_e x_e : \\ \forall S \in \mathcal{S} \sum_{e \in \delta(S)} x_e \geq 1 \\ x_e \geq 0 \end{aligned}$$



$$\begin{aligned}
 & \min \sum_e c_e x_e : \\
 & \forall S \in \mathcal{S} \sum_{e \in \delta(S)} x_e \geq 1 \\
 & x_e \geq 0
 \end{aligned}$$



**Exponential
number of constraints
but
still solvable
- does not need to be solved**

Steiner forest

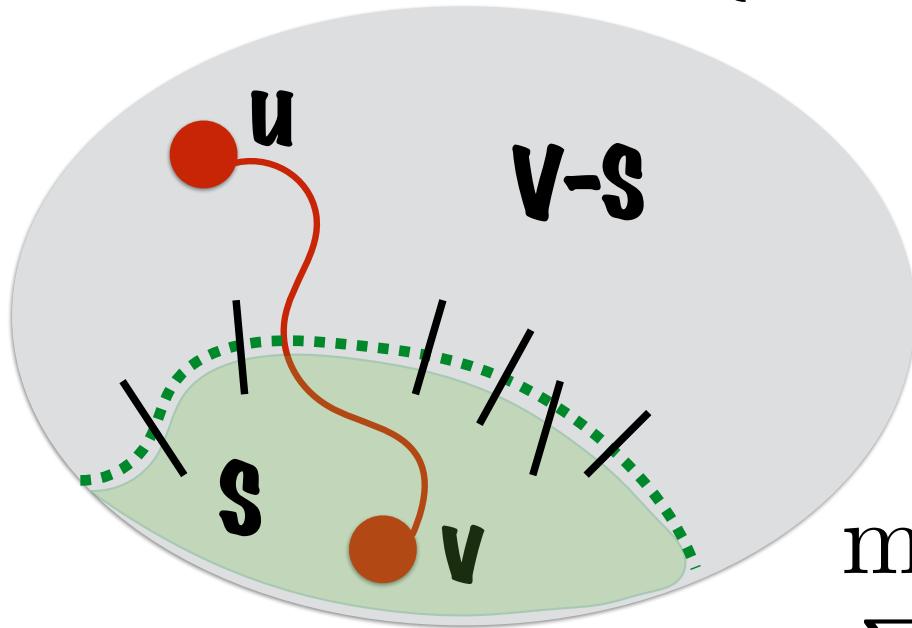


Steiner forest



Linear programming relaxation

$$\mathcal{S} = \{S : \exists i \exists u, v \in S_i : |S \cap \{u, v\}| = 1\}$$



$$\begin{aligned} & \min \sum_e c_e x_e : \\ & \sum_{e \in \delta(S)} x_e \geq 1 \quad \forall S \in \mathcal{S} \\ & x_e \geq 0 \quad \forall e \in E \end{aligned}$$

Taking the dual

$$\begin{aligned} \min \sum_e c_e x_e : \\ \sum_{e \in \delta(S)} x_e \geq 1 & \quad \forall S \in \mathcal{S} \quad [y_S] \\ x_e \geq 0 & \quad \forall e \in E \end{aligned}$$

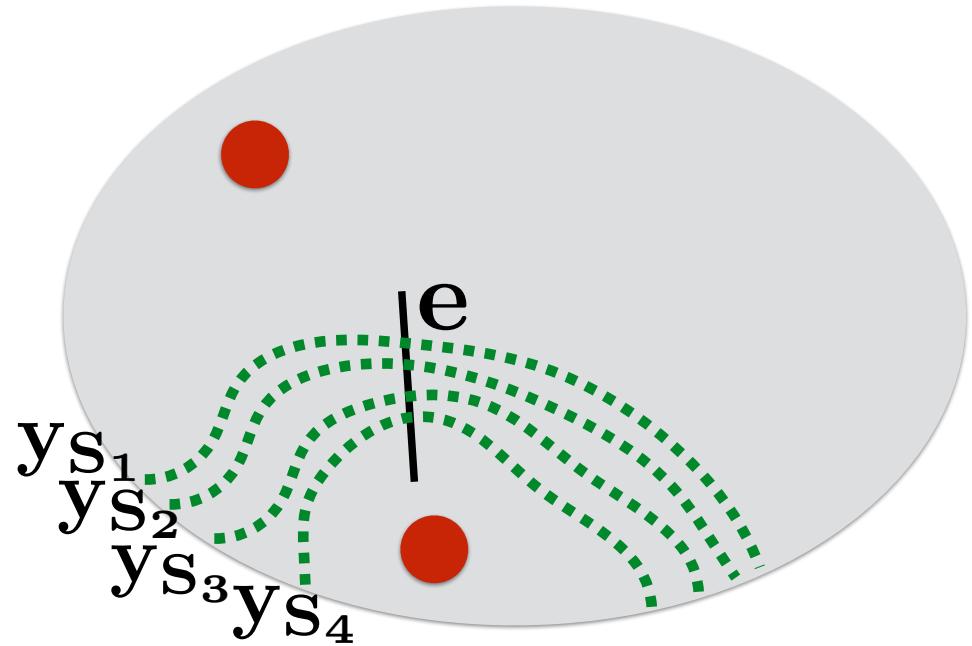
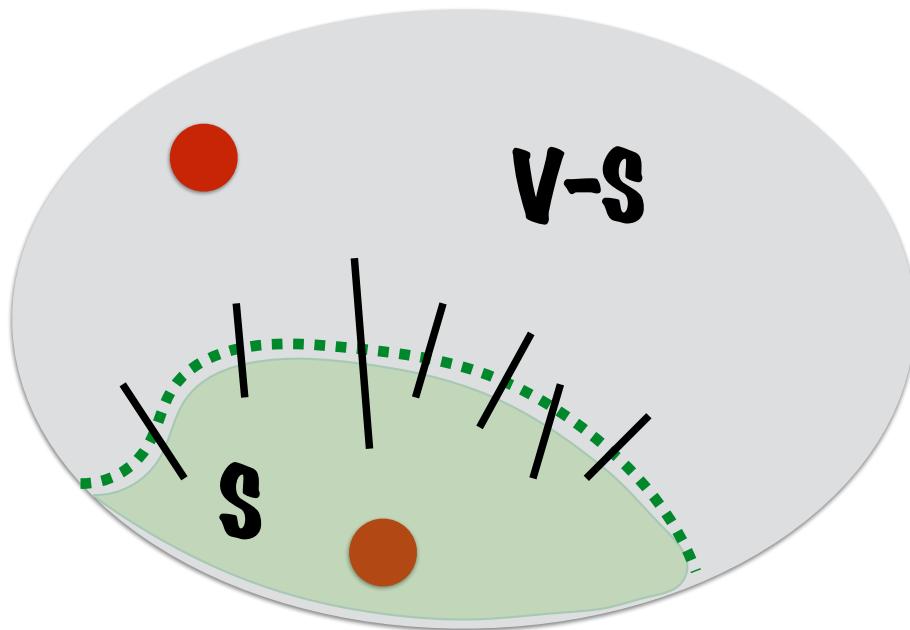
$$\begin{aligned} \max \sum_S y_S : \\ \sum_{S: e \in \delta(S)} y_S \leq c_e & \quad \forall e \in E \quad [x_e] \\ y_S \geq 0 & \quad \forall S \in \mathcal{S} \end{aligned}$$

Interpreting the dual

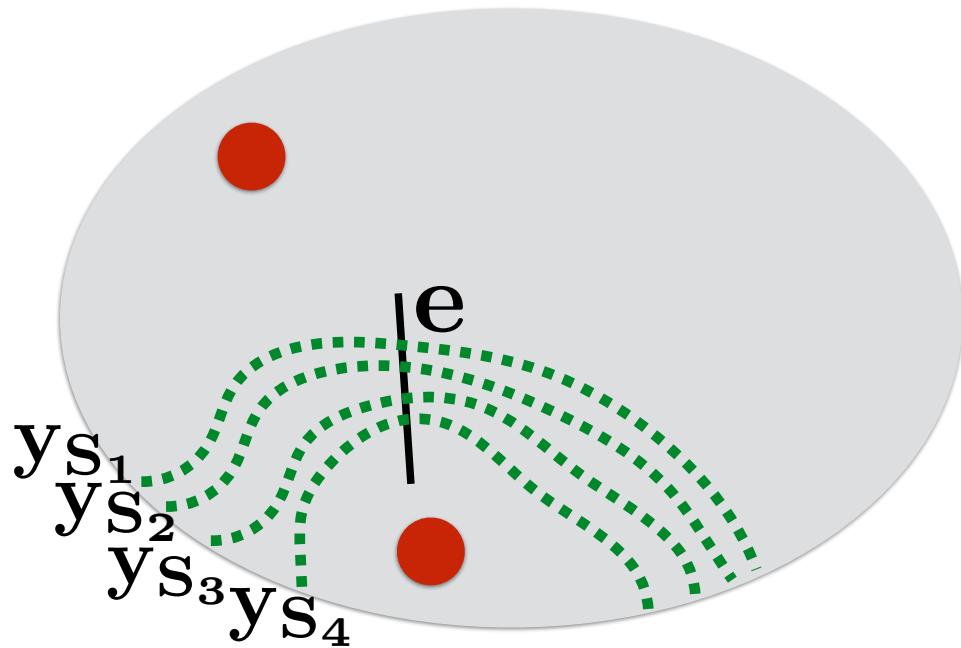
$$\max \sum_S y_S :$$

$$\sum_{S : e \in \delta(S)} y_S \leq c_e \quad \forall e \in E \quad [x_e]$$

$$y_S \geq 0 \quad \forall S \in \mathcal{S}$$



$$\begin{aligned}
 & \max \sum_S y_S : \\
 & \sum_{S: e \in \delta(S)} y_S \leq c_e \quad \forall e \in E \quad [x_e] \\
 & y_S \geq 0 \quad \forall S \in \mathcal{S}
 \end{aligned}$$



Soooo many cuts containing e!!

Q: How can we hope for feasibility?

A: most variables will be 0

Steiner forest



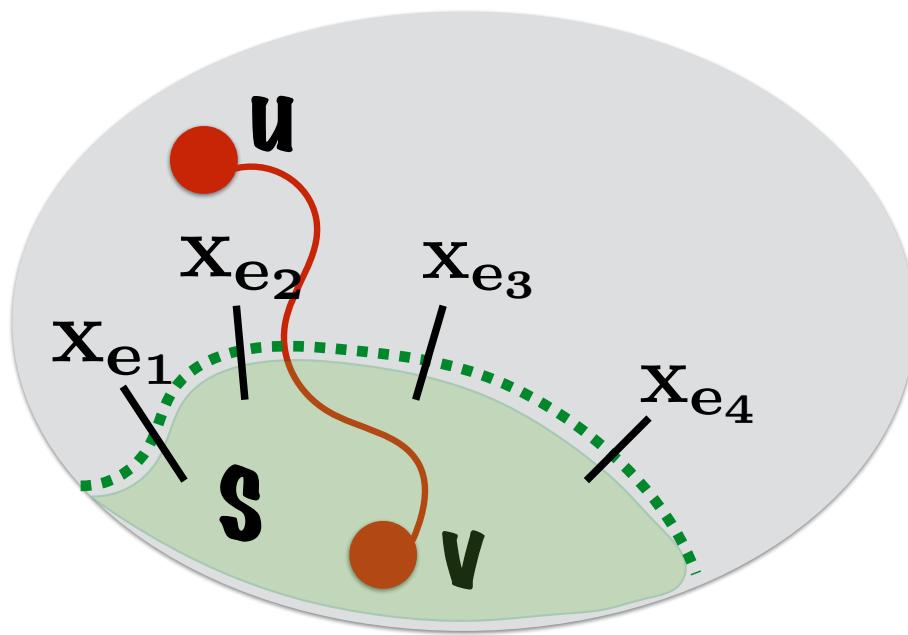
Steiner forest



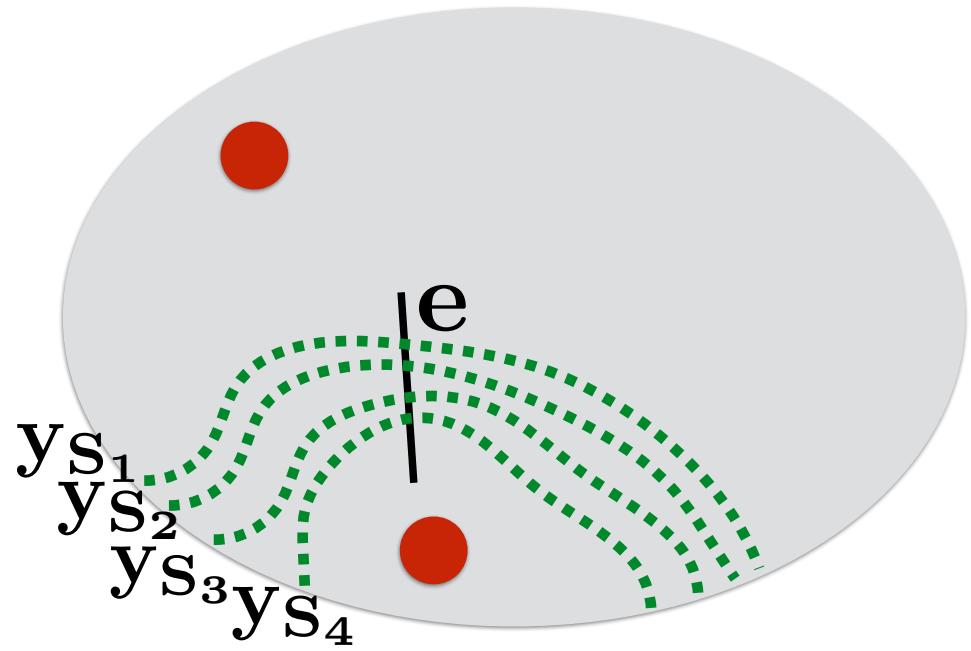
primal-dual pair

$$\begin{aligned} \min \sum_e c_e x_e : \\ \sum_{e \in \delta(S)} x_e \geq 1 & \quad \forall S \in \mathcal{S} \quad [y_S] \\ x_e \geq 0 & \quad \forall e \in E \end{aligned}$$

$$\begin{aligned} \max \sum_S y_S : \\ \sum_{S: e \in \delta(S)} y_S \leq c_e & \quad \forall e \in E \quad [x_e] \\ y_S \geq 0 & \quad \forall S \in \mathcal{S} \end{aligned}$$



$$\sum_{e \in \delta(S)} x_e \geq 1$$



$$\sum_{S: e \in \delta(S)} y_S \leq c_e$$

$$\begin{aligned} \min \sum_e c_e x_e : \\ \sum_{e \in \delta(S)} x_e \geq 1 & \quad \forall S \in \mathcal{S} \\ x_e \geq 0 & \quad \forall e \in E \end{aligned}$$

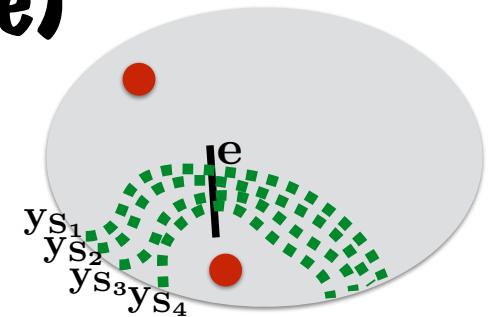
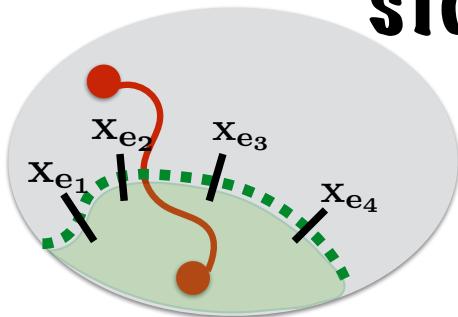
$$\begin{aligned} \max \sum_S y_S : \\ \sum_{S: e \in \delta(S)} y_S \leq c_e & \quad \forall e \in E \\ y_S \geq 0 & \quad \forall S \in \mathcal{S} \end{aligned}$$

Initialization:

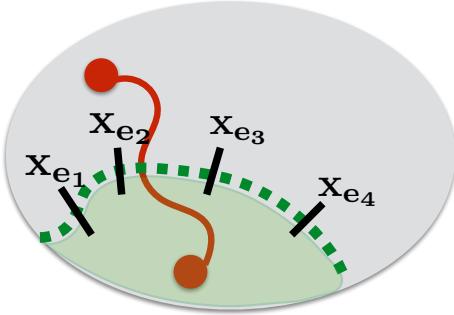
$$x \leftarrow 0, y \leftarrow 0$$

Iteration: while x not satisfiable
 raise y as much as possible
 stopped by tight constraint (e)

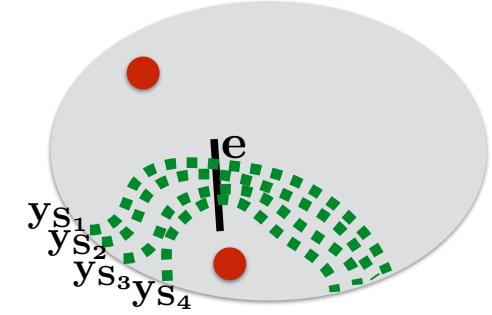
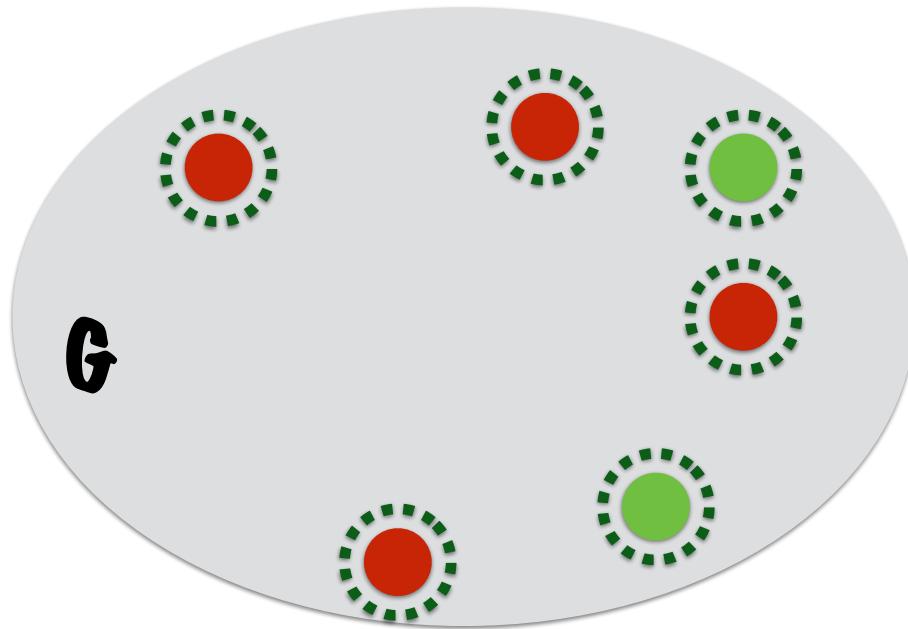
$$x_e \leftarrow 1$$



Q: How do we raise y ?



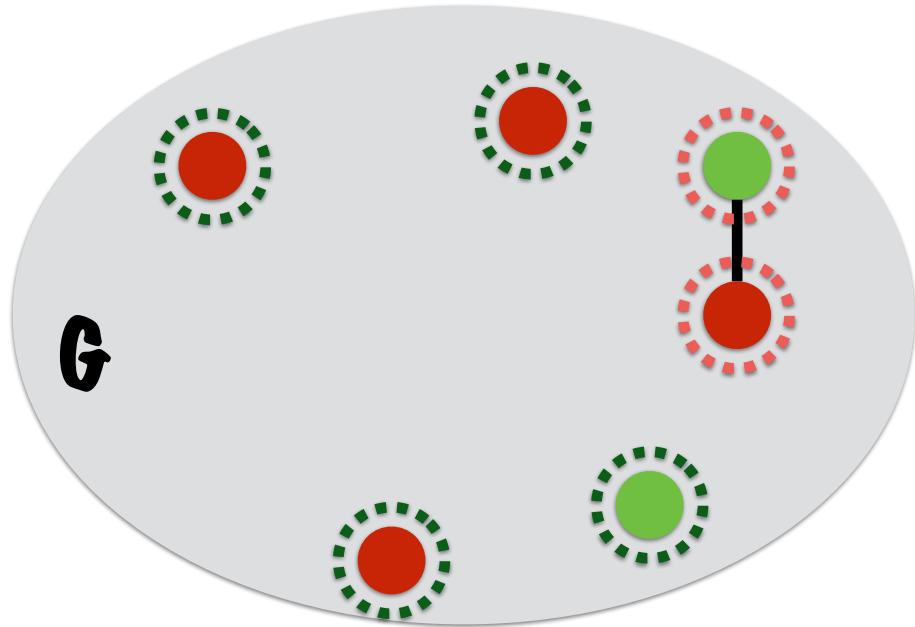
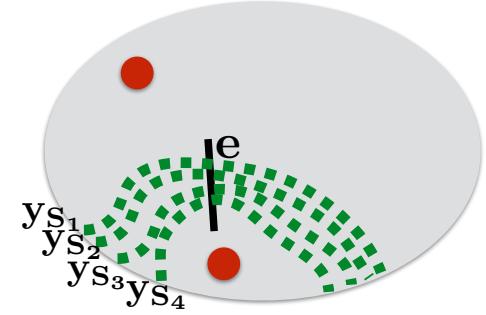
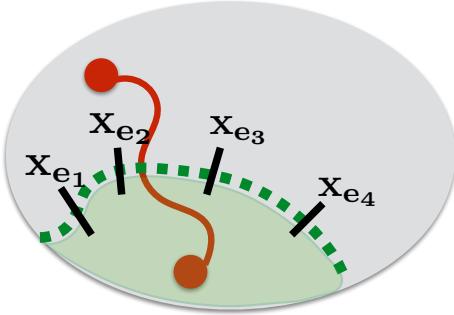
Initially...



Raise all singleton terminals simultaneously

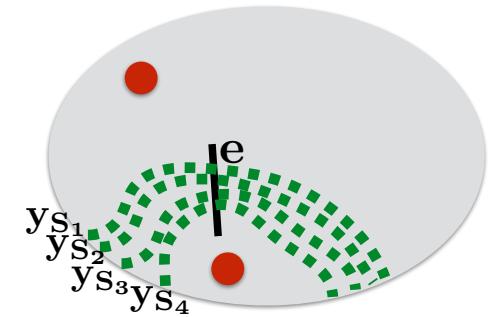
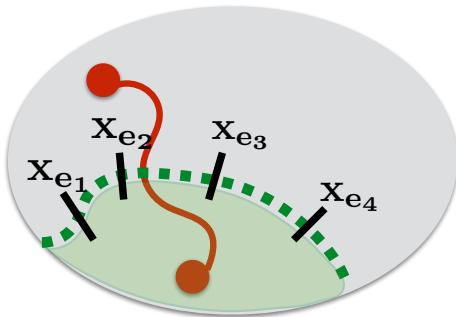
$$y_S \leftarrow y_S + \epsilon \quad \forall S = \{u\}$$

Q: How do we raise y ?

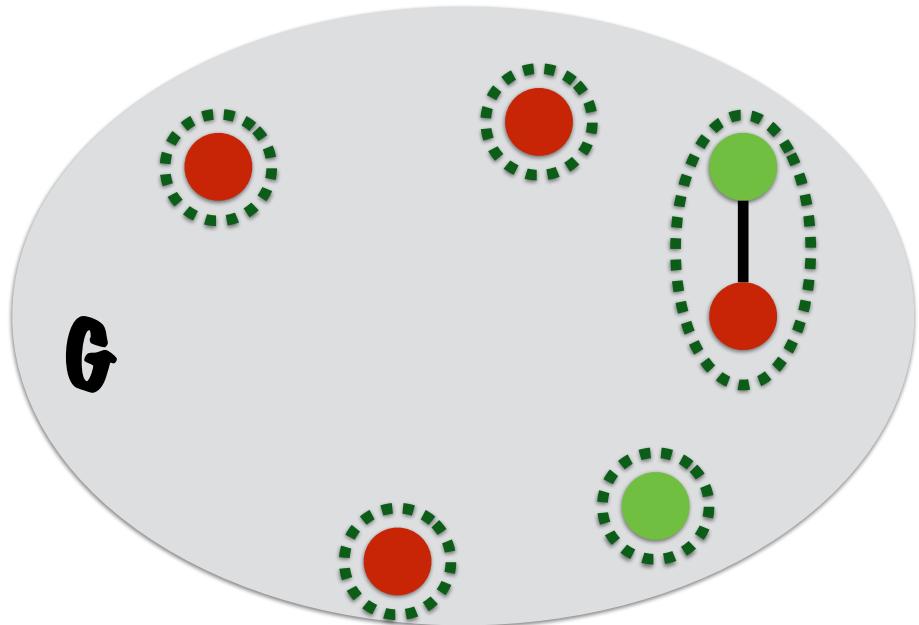


Tight
constraint...

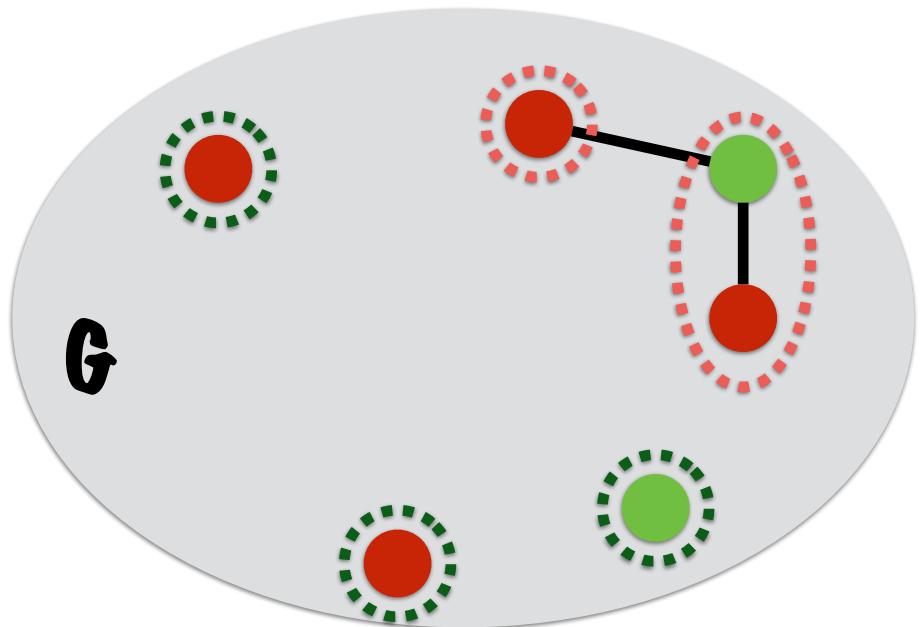
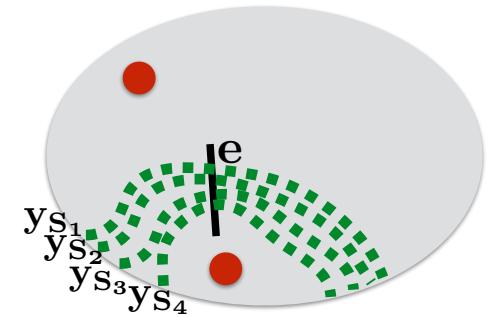
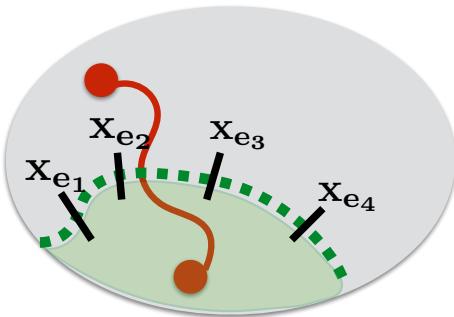
Q: How do we raise y?



Then...

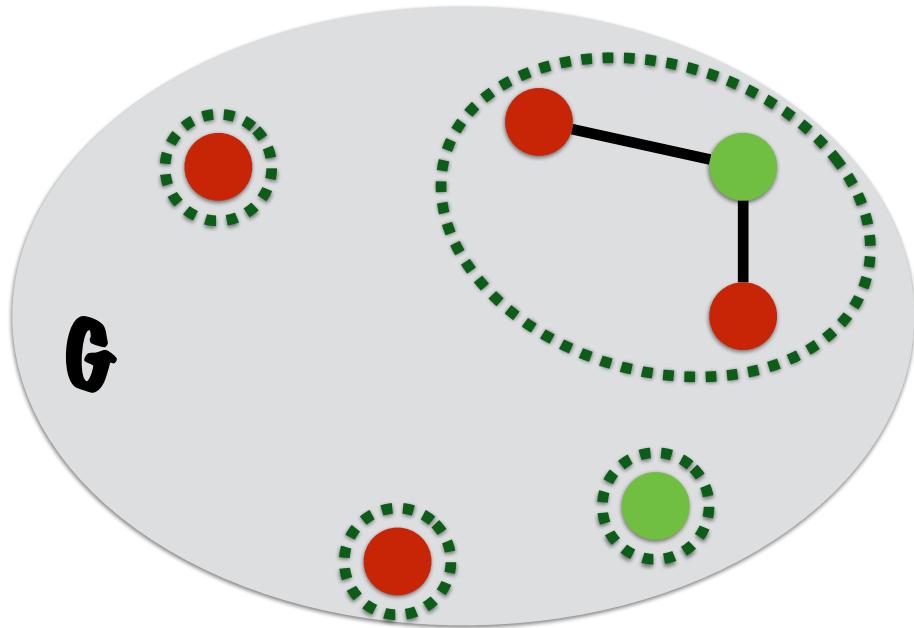
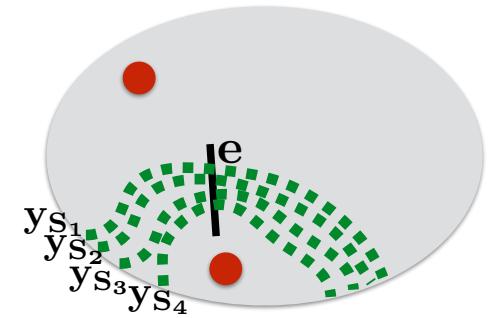
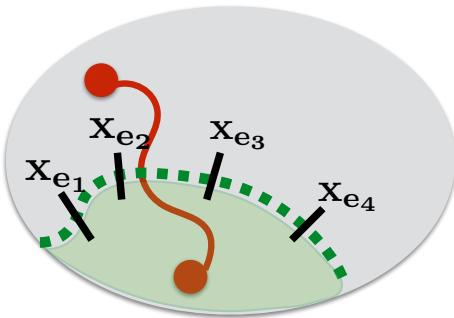


Q: How do we raise y?



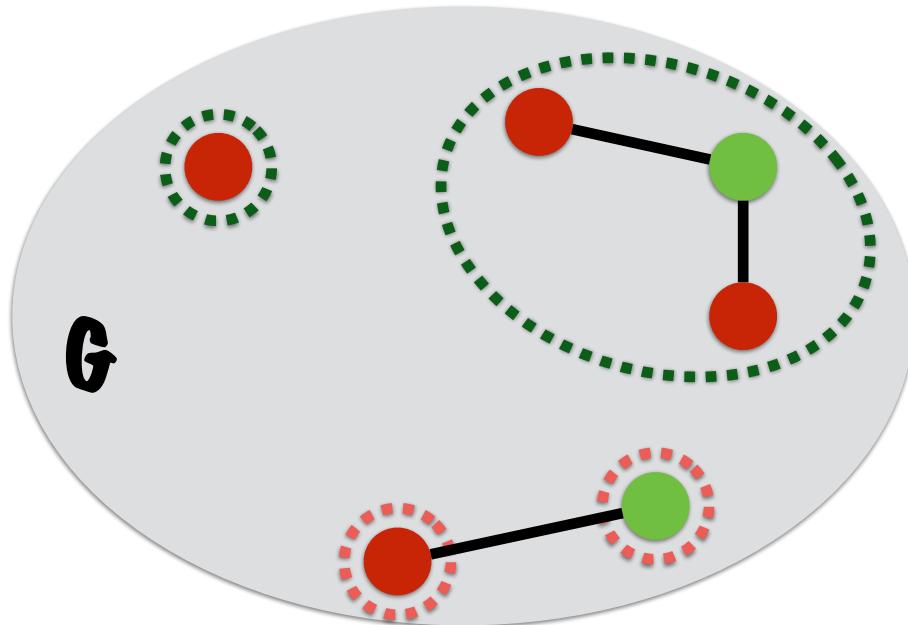
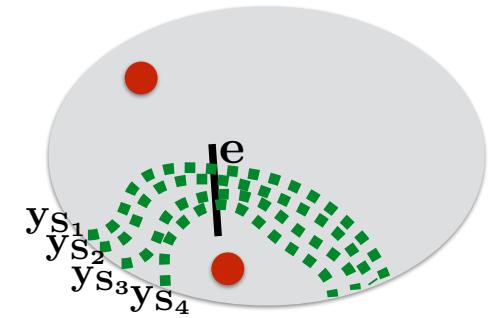
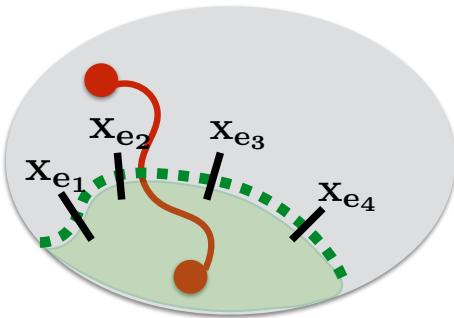
**Tight
constraint...**

Q: How do we raise y?



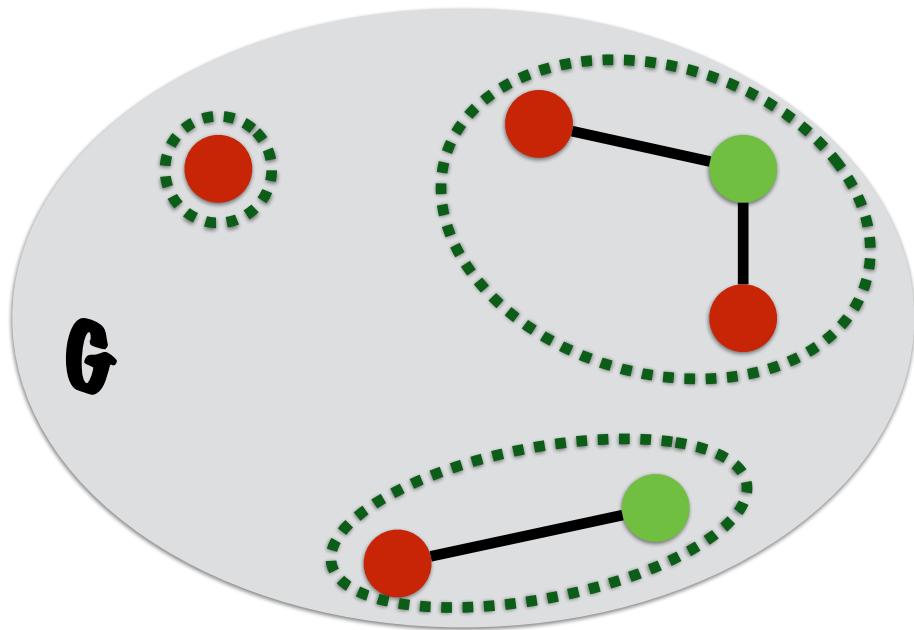
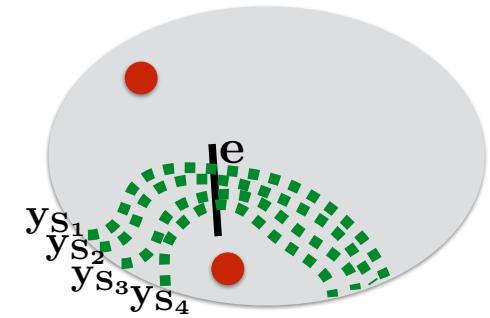
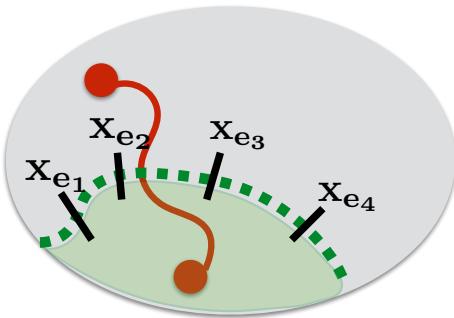
Then...

Q: How do we raise y?



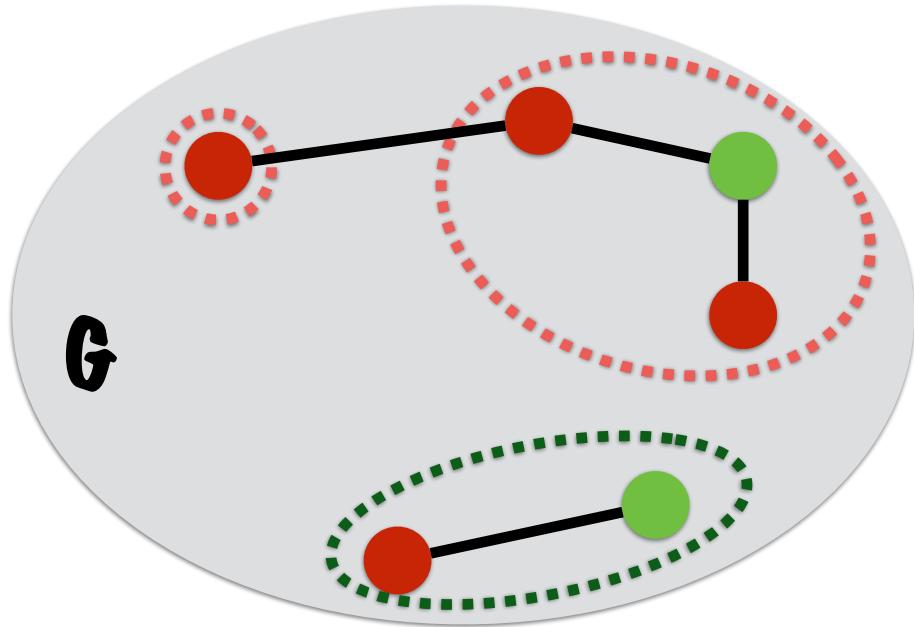
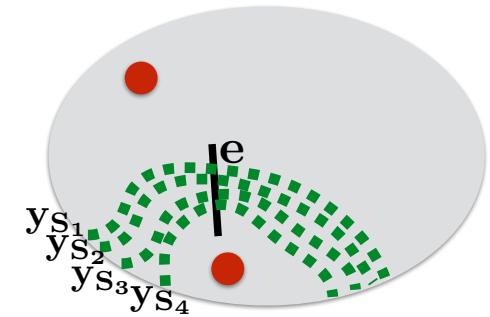
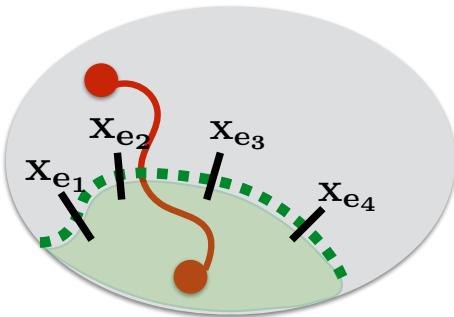
**Tight
constraint...**

Q: How do we raise y?



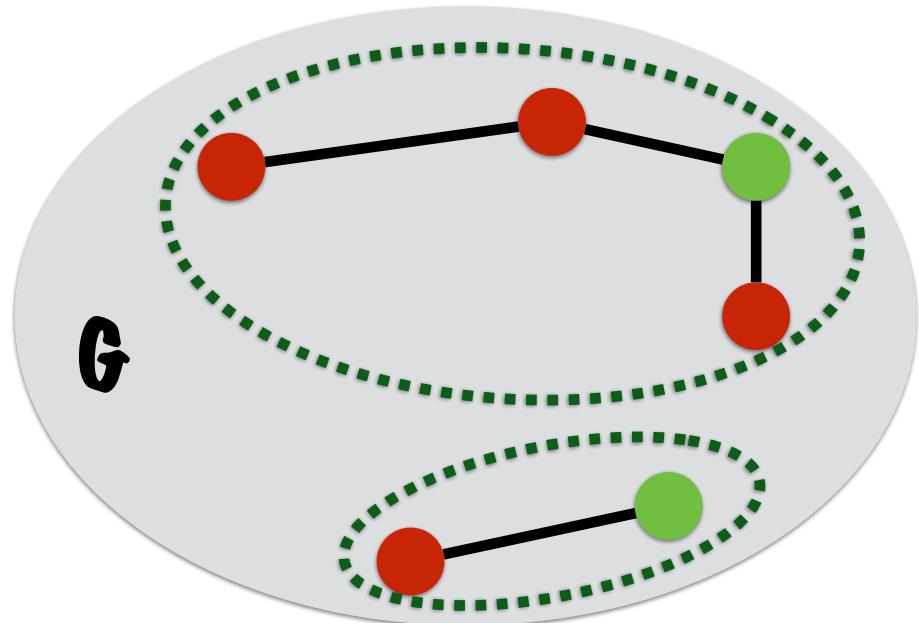
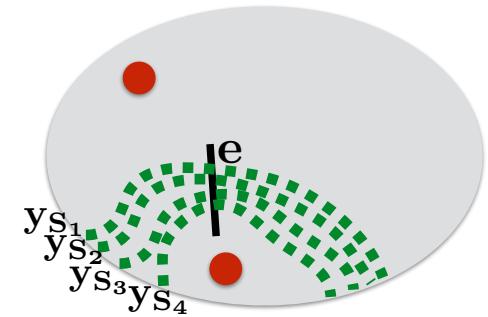
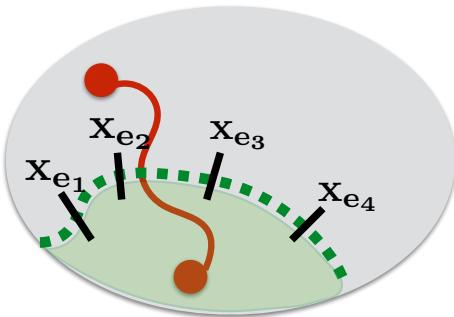
Then...

Q: How do we raise y?



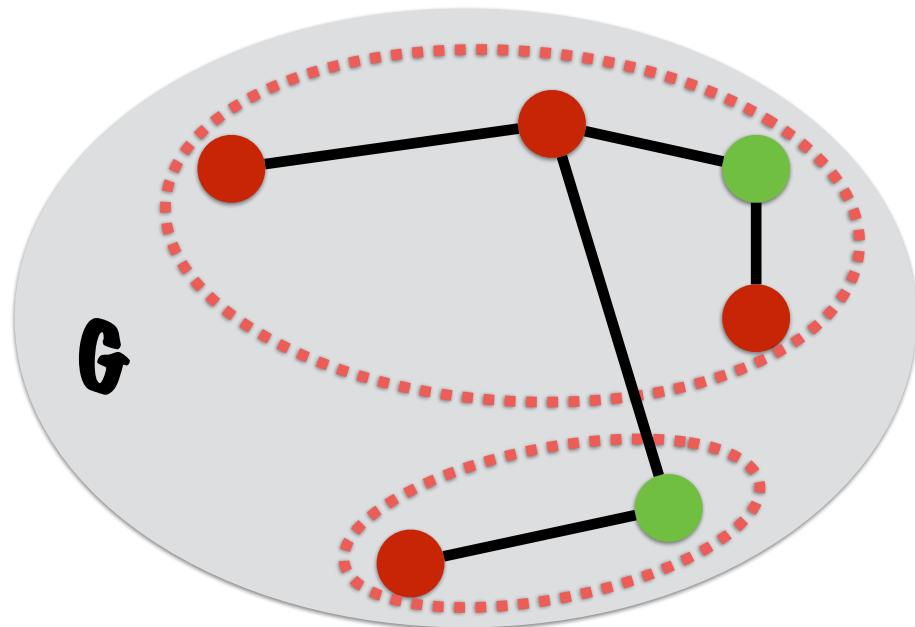
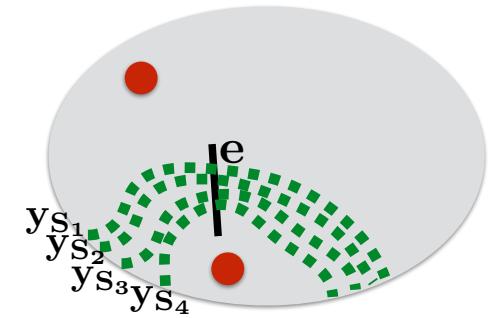
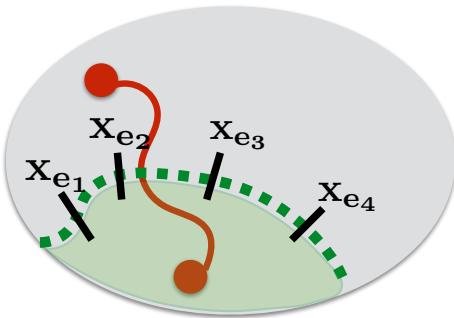
**Tight
constraint...**

Q: How do we raise y?



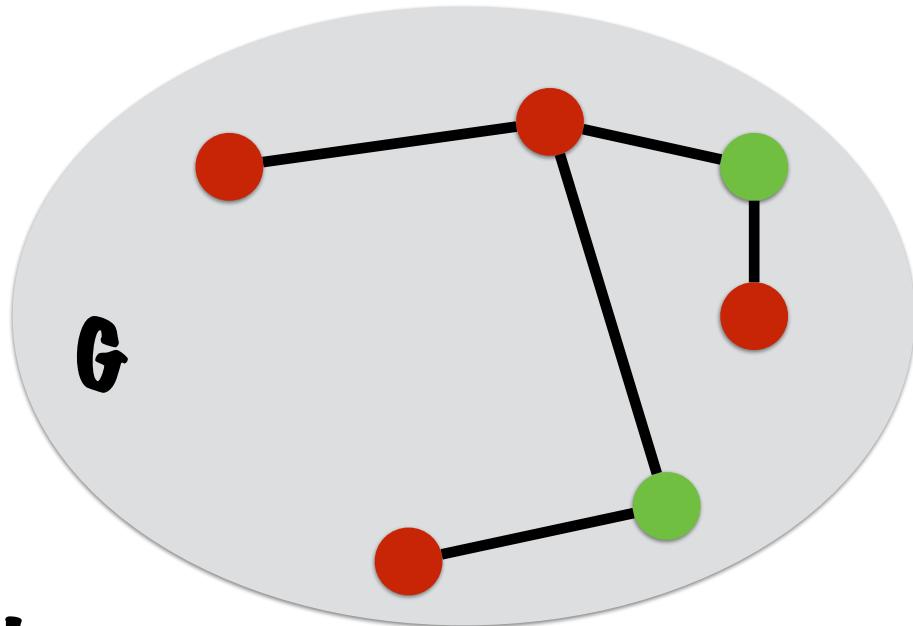
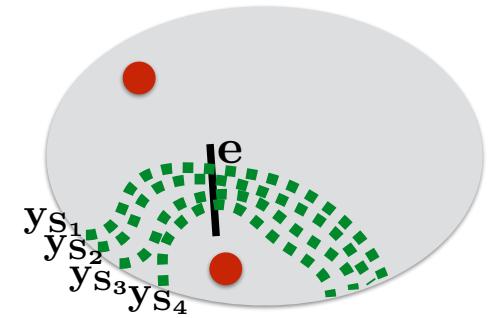
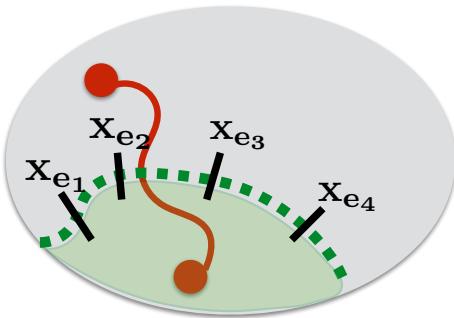
Then...

Q: How do we raise y?



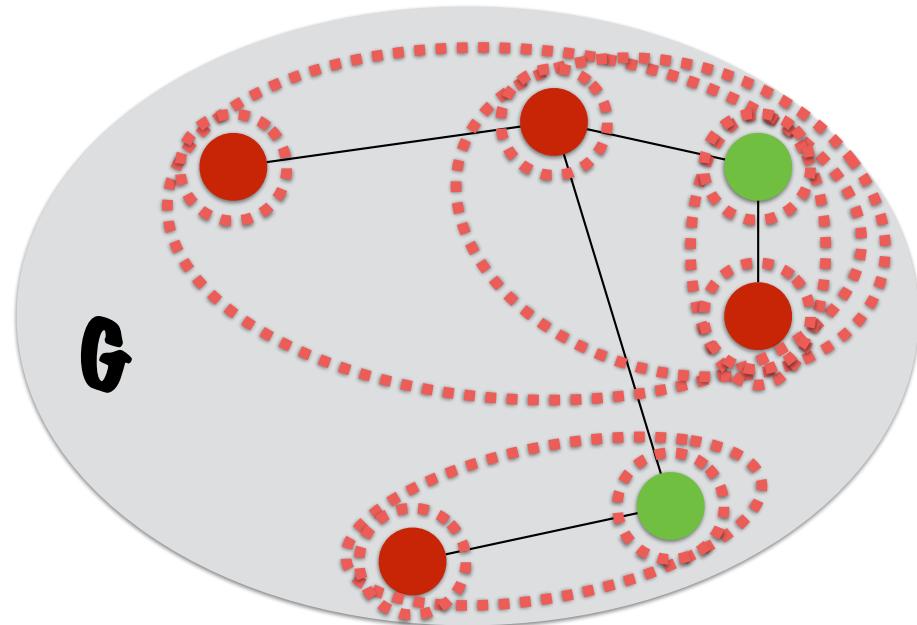
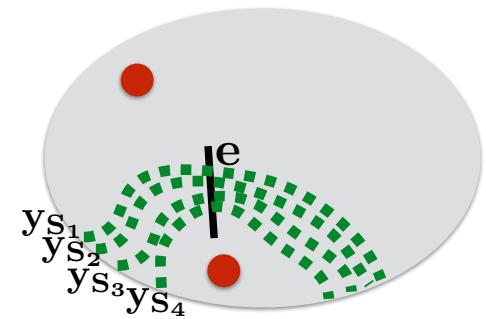
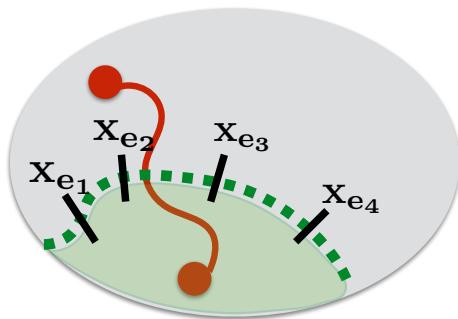
**Tight
constraint...**

Q: How do we raise y ?

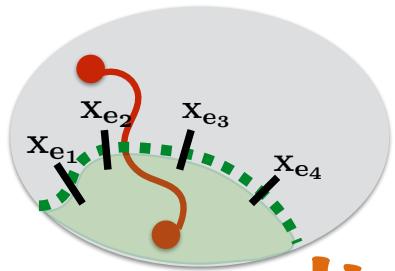


x is feasible

Q: How did we raise y ?



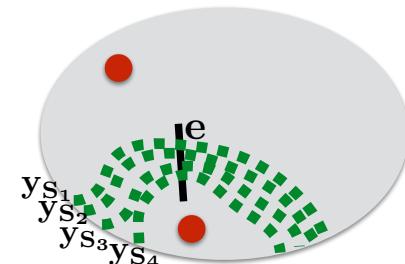
$$\begin{aligned} \min \sum_e c_e x_e : \\ \sum_{e \in \delta(S)} x_e \geq 1 & \quad \forall S \in \mathcal{S} \\ x_e \geq 0 & \quad \forall e \in E \end{aligned}$$



$$\begin{aligned} \max \sum_S y_S : \\ \sum_{S: e \in \delta(S)} y_S \leq c_e & \quad \forall e \in E \\ y_S \geq 0 & \quad \forall S \in \mathcal{S} \end{aligned}$$

Initialization:

$$x \leftarrow 0, y \leftarrow 0$$



Iteration: while x not satisfiable
in parallel, raise every **unfrozen** y_S with
minimal S

stopped by tight constraint (e)
 $x_e \leftarrow 1$

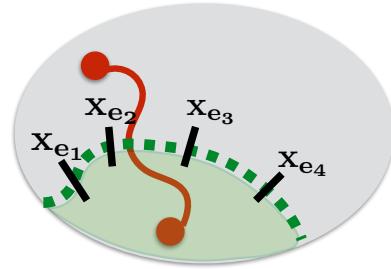
freeze y_S in tight constraints

Steiner forest



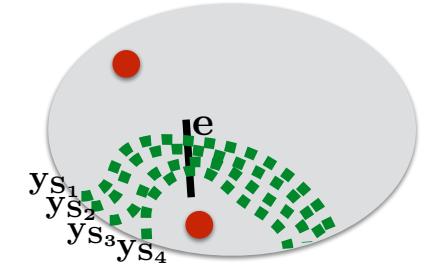
Steiner forest





Initialization:

$x \leftarrow 0, y \leftarrow 0$



Iteration: while x not satisfiable
in parallel, raise every unfrozen y_S with
minimal S

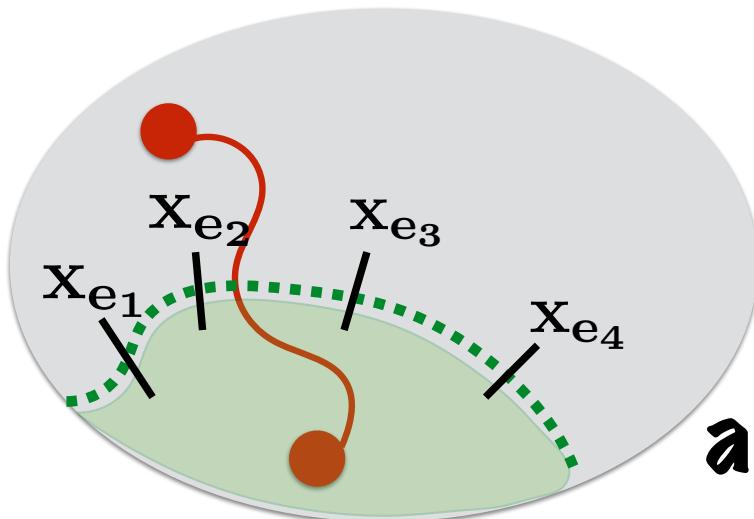
stopped by tight constraint (e)
 $X_e \leftarrow 1$

freeze y_S in tight constraints

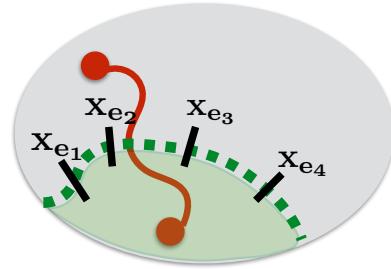
Do we ever get stuck?

**Do we ever get stuck?
Suppose we do**

cut: x not feasible?

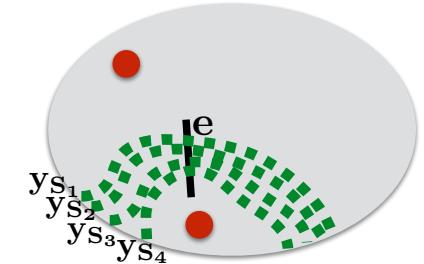


**If we cannot raise $y(S)$
it's because
some e in the cut is tight
and then we would have put $x(e)$
in solution
QED**



Initialization:

$$x \leftarrow 0, y \leftarrow 0$$



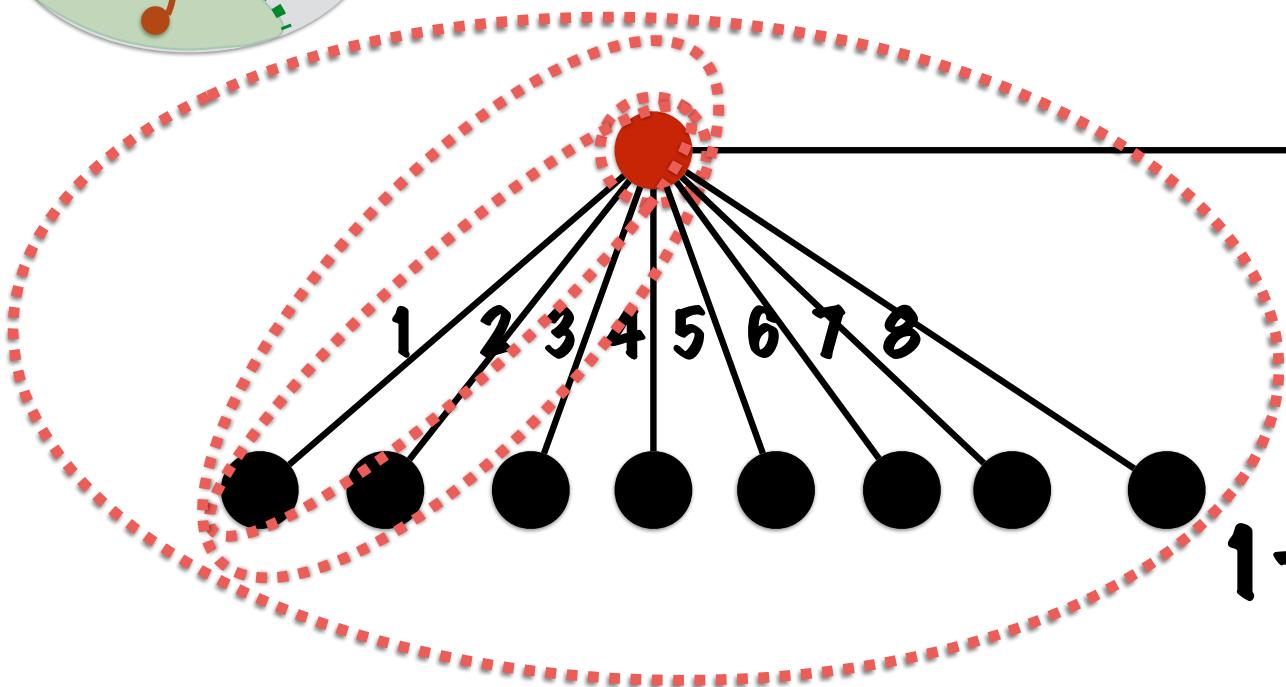
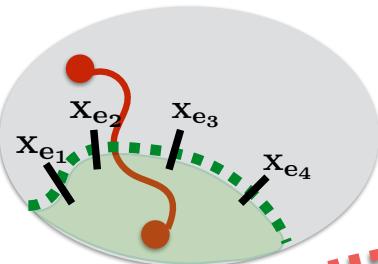
Iteration: while x not satisfiable
in parallel, raise every unfrozen y_S with
minimal S

stopped by tight constraint (e)
 $X_e \leftarrow 1$

freeze y_S in tight constraints

Fact: final x, y are feasible.

What about output cost?



18

Output cost
 $1+2+\dots+8+18$

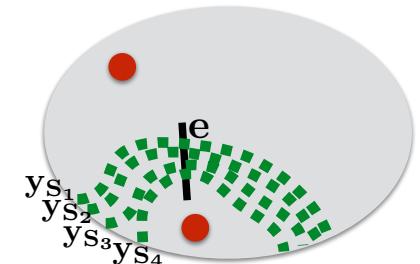
dual value

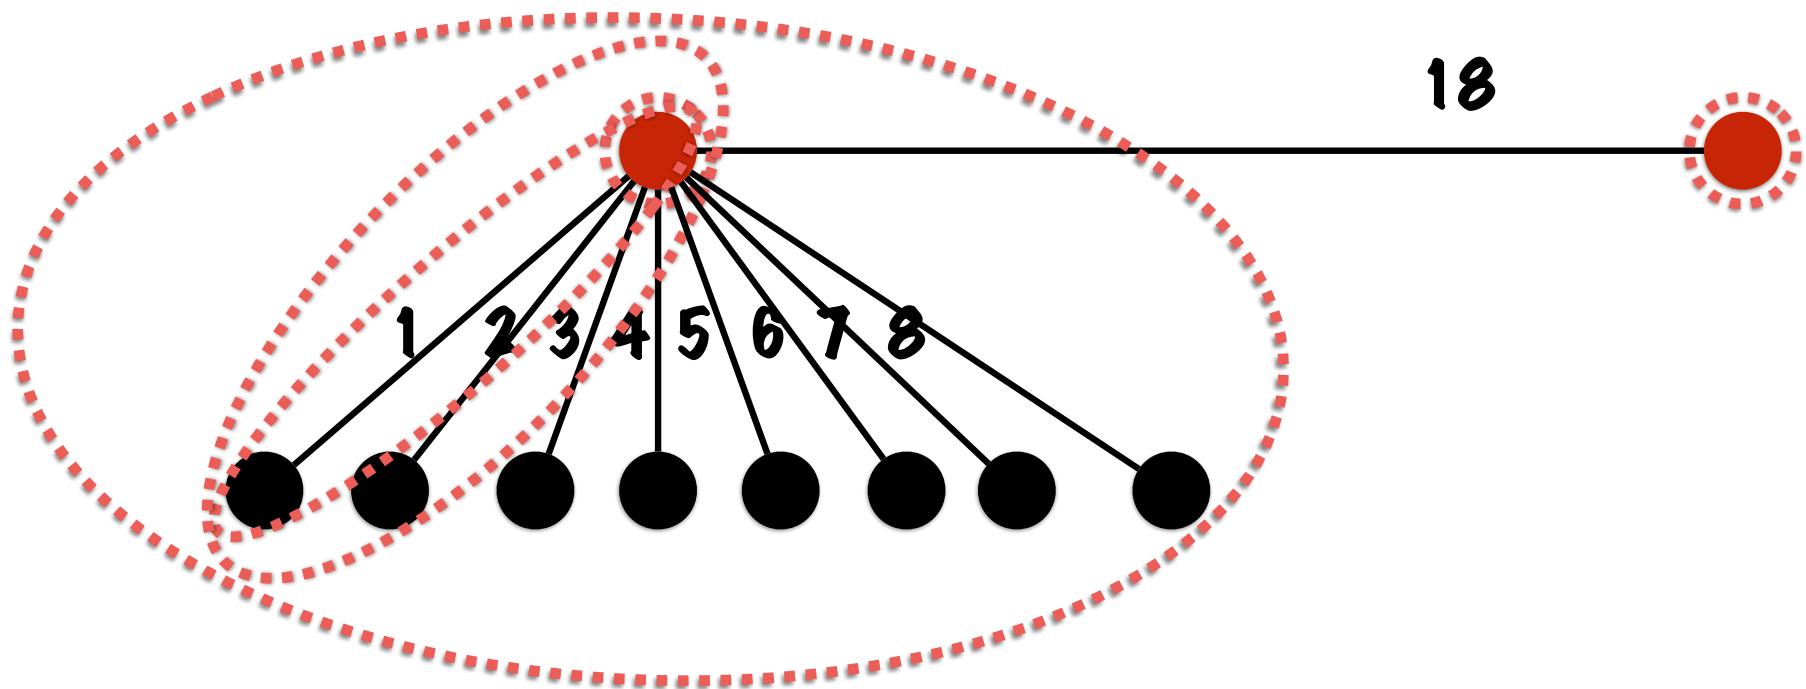
$$1+1+\dots+1+1+9=18$$

OPT

18

Bad!



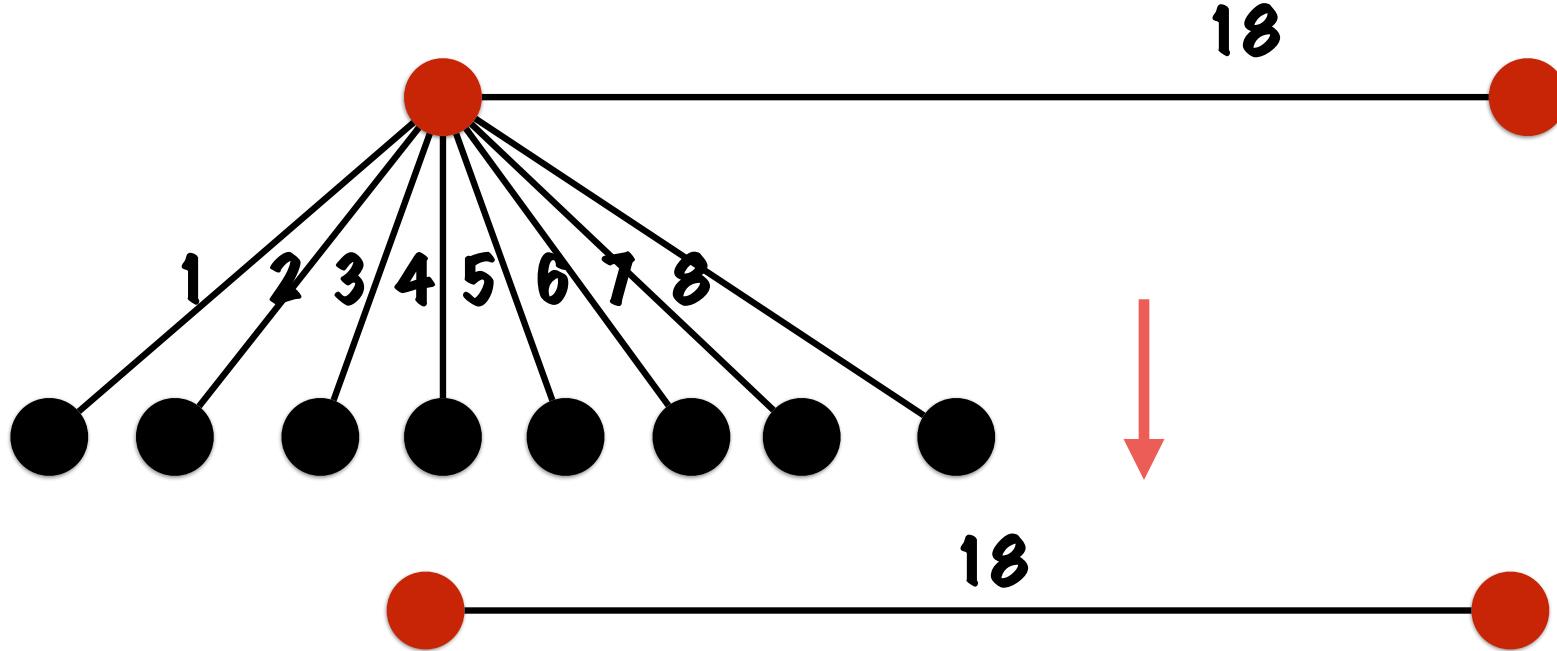


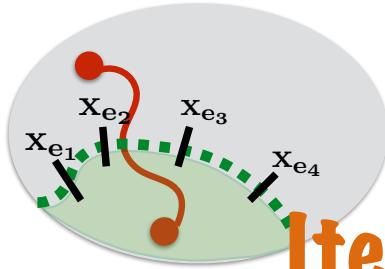
Observe: Many of the edges in the output
are useless

Idea: prune useless edges

Modified algorithm

Consider set of edges defined by x
remove unnecessary edges
Output resulting set





Initialization:

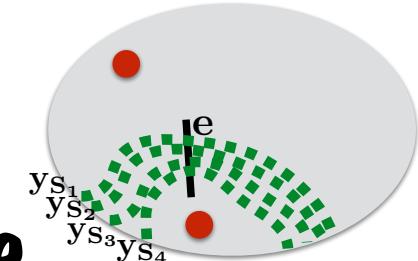
$$x \leftarrow 0, y \leftarrow 0$$

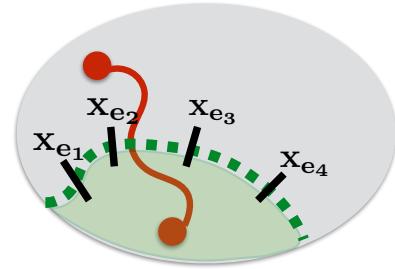
Iteration: while x not satisfiable
in parallel, raise every unfrozen y_S with
 S minimal

stopped by tight constraint (e)
 $x_e \leftarrow 1$

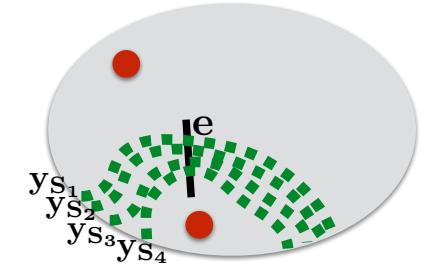
freeze y_S in tight constraints

Pruning: let $F = \{\text{edges defined by } x\}$
for each edge e of F in reverse order
remove e if unnecessary





Theorem



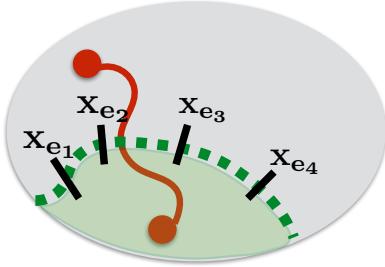
It's a 2-approximation for Steiner forest

Steiner forest



Steiner forest





Initialization:

$$x \leftarrow 0, y \leftarrow 0$$

Iteration: while x not satisfiable
in parallel, raise every unfrozen y_S with
 S minimal

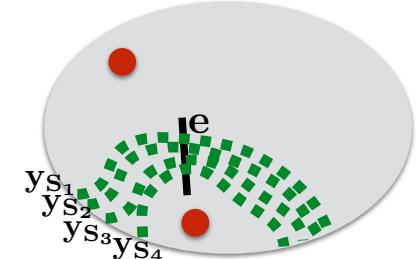
stopped by tight constraint (e)
 $x_e \leftarrow 1$

freeze y_S in tight constraints

Pruning: let $F = \{edges defined by x\}$
for each edge e of F in reverse order,
remove e if unnecessary

Theorem:

It's a 2-approximation for Steiner forest



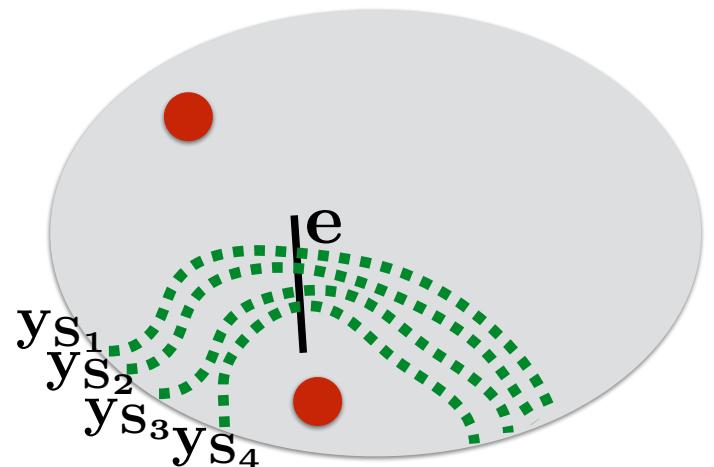
Observations:

- Final y is feasible
- Final x is a feasible forest
- Output F' is a forest
- its leaves are terminals
- slackness condition:

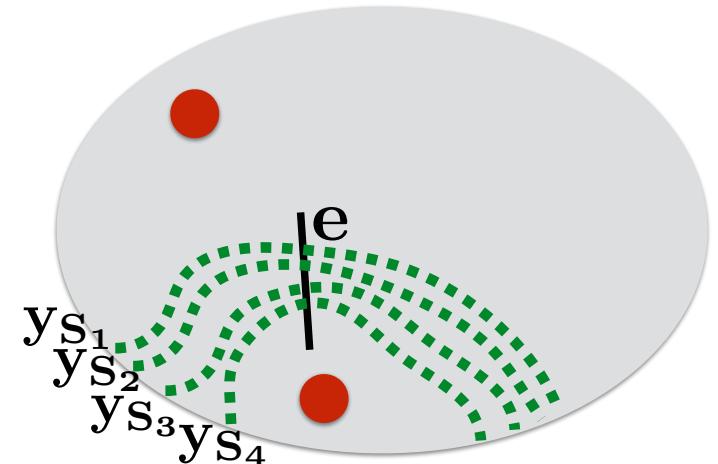
$$e \in F' \implies e \in F$$

$$\implies x_e = 1$$

$$\implies \sum_{S: e \in \delta(S)} y_S = c_e$$



**slackness condition implies
bound on cost of output
by using dual variables:**



$$\sum_{e \in F'} c_e = \sum_{e \in F'} \sum_{S: e \in \delta(S)} y_S$$

As usual, invert summations:

$$\sum_{e \in F'} \sum_{S: e \in \delta(S)} y_S = \sum_S y_S |F' \cap \delta(S)|$$

Q: How to upper bound

$$\sum_S y_S |F' \cap \delta(S)|$$

A: Incrementally Bound

$$\sum_{\text{time } t} \sum_S \text{"active" between } t \text{ and } t + dt \epsilon |F' \cap \delta(S)|$$

S “active” if its dual variable is being raised

Definition of “active” at current time

Initialization:

$$x \leftarrow 0, y \leftarrow 0$$

Iteration: while x not satisfiable
in parallel, raise every unfrozen y_s with
 S minimal

stopped by tight constraint (e)

$$x_e \leftarrow 1$$

freeze y_s in tight constraints

Pruning: let $F = \{ \text{edges defined by } x \}$
remove unnecessary edges from F

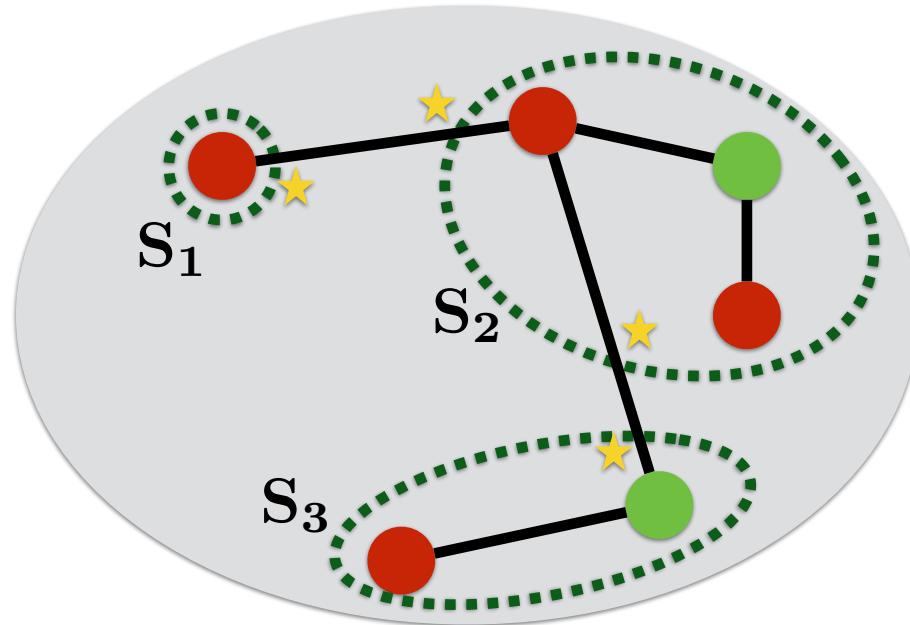
“active”:
unfrozen
with
 S minimal

Q: How to upper bound

$$\sum_{S \text{ currently active}} |F' \cap \delta(S)|$$

A: Let's try some examples

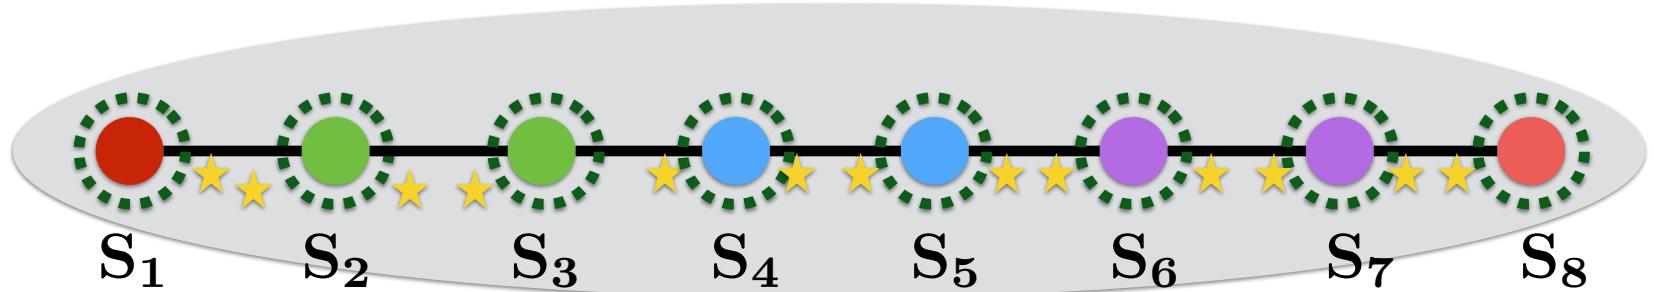
$\overline{F'}$



3 active sets: S_1, S_2, S_3

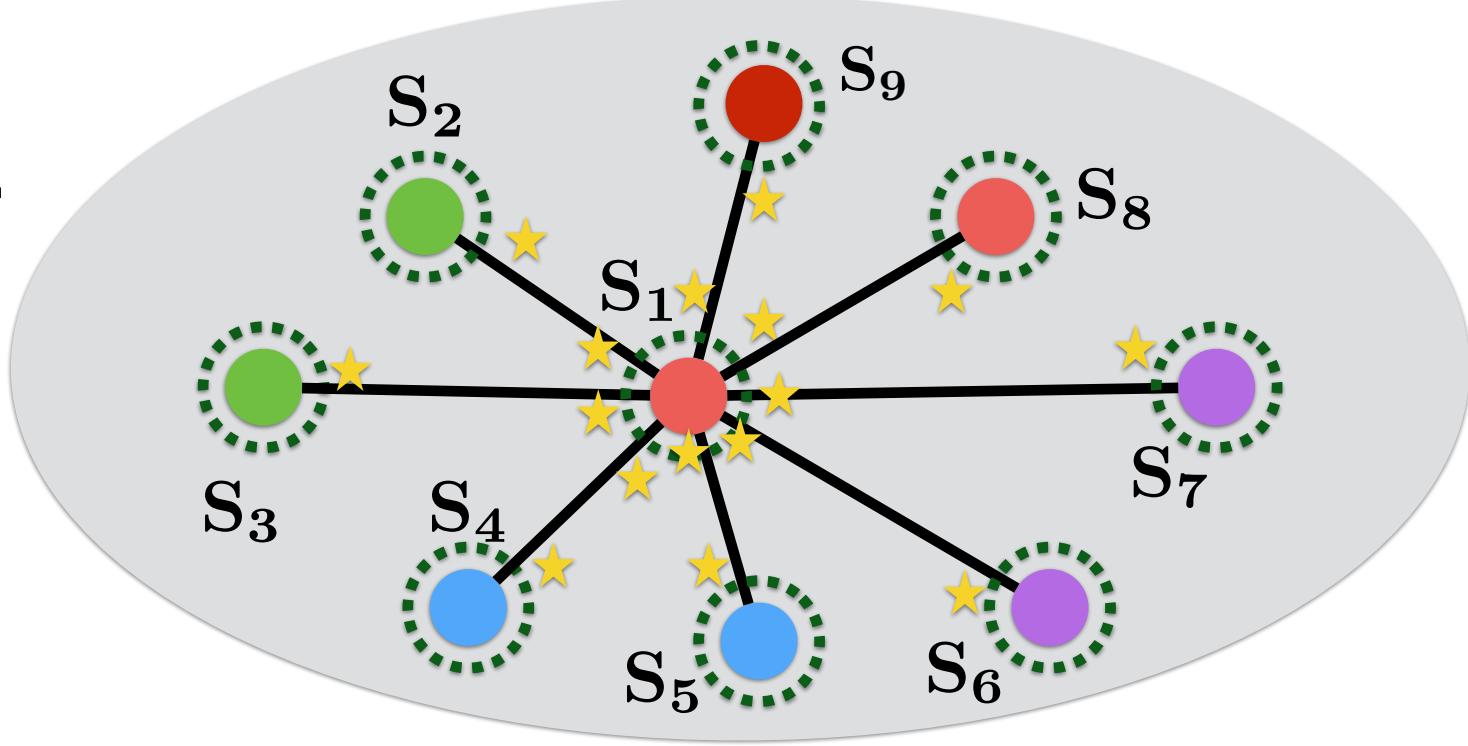
$$\sum_{S \text{ active}} |F' \cap \delta(S)| = \#(\text{stars}) = 1 + 2 + 1$$

F'

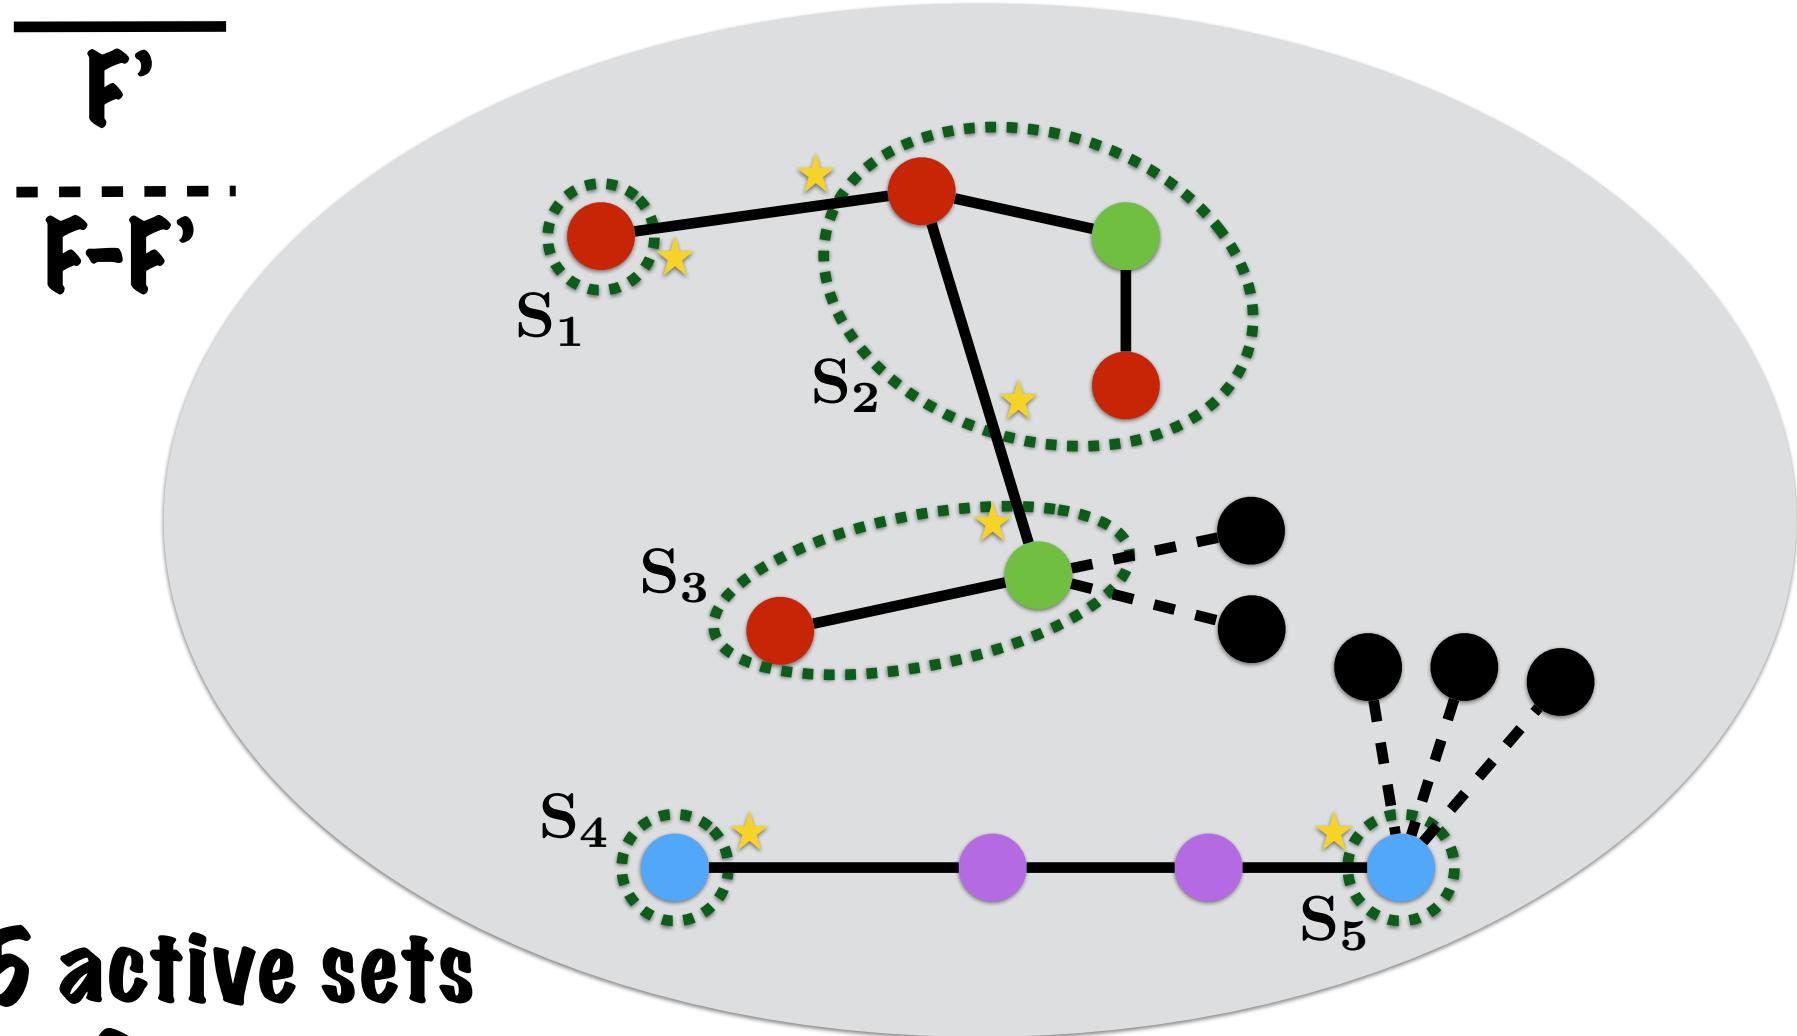


**8 active sets
14 stars**

$\overline{F'}$



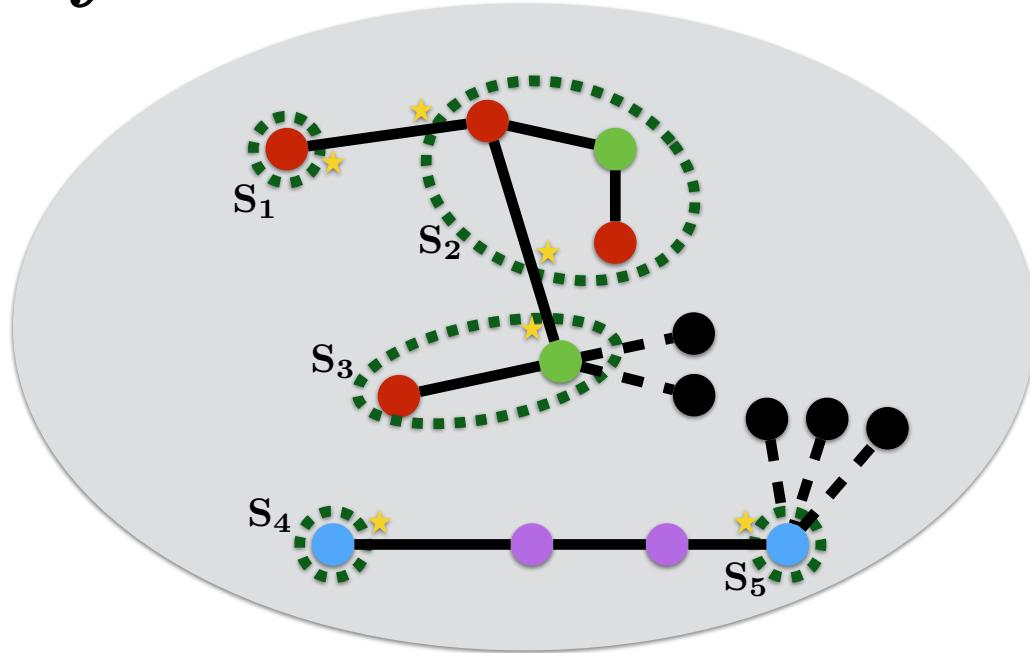
**9 active sets
16 stars**



5 active sets
6 stars

Q: How to upper bound

$$\sum_{S \text{ currently active}} |F' \cap \delta(S)|$$



Lemma:

$$\frac{\sum_{S \text{ currently active}} |F' \cap \delta(S)|}{\#(\text{currently active sets})} \leq 2$$

Analysis assuming the lemma holds

Cost(output) =

$$\begin{aligned}\sum_{e \in F'} c_e &= \sum_{e \in F'} \sum_{S: e \in \delta(S)} y_S \\&= \sum_S y_S |F' \cap \delta(S)| \\&= \sum_t \sum_{S \text{ active}} \epsilon |F' \cap \delta(S)| \\&\leq \sum_t 2\epsilon \cdot \#(S \text{ active}) \\&= 2 \sum_S y_S \\&\leq 2 \cdot \mathbf{OPT}\end{aligned}$$

Steiner forest



Steiner forest



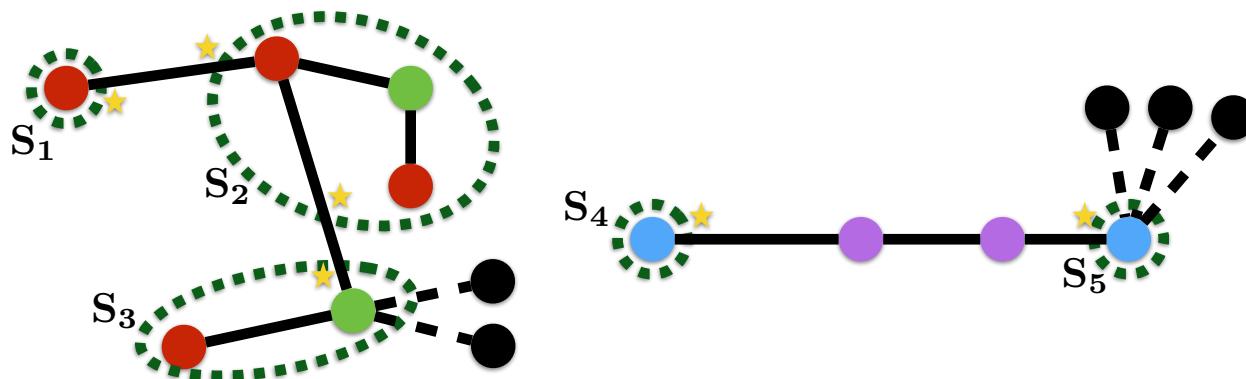
Lemma

Let F' denote the output set of edges

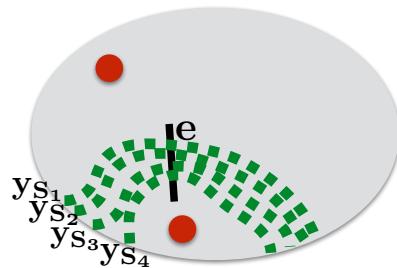
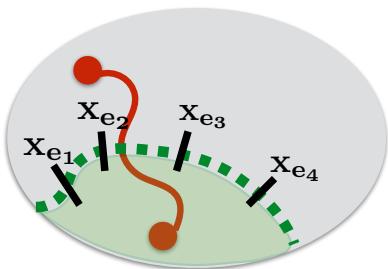
Fix a time of the execution and say a set S is active if its dual variable is being raised.

Then:

$$\frac{\sum_{S \text{ active}} |F' \cap \delta(S)|}{\#(\text{active sets})} \leq 2$$



What properties do F' and active sets have?



Initialization:

$$x \leftarrow 0, y \leftarrow 0$$

Iteration: while x not satisfiable
in parallel, raise every unfrozen y_S with
 S minimal

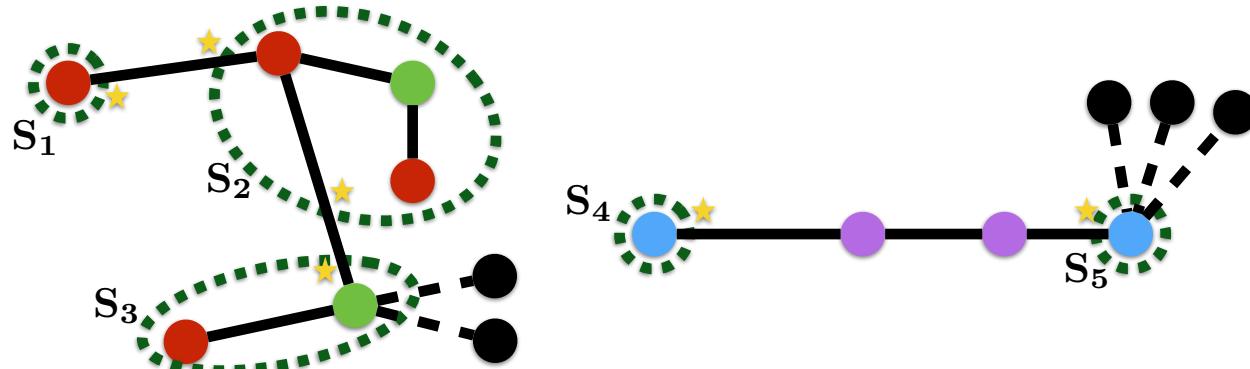
stopped by tight constraint (e)
 $x_e \leftarrow 1$

freeze y_S in tight constraints

Pruning: let $F = \{ \text{edges defined by } x \}$
for each edge e of F in reverse order,
remove e if unnecessary

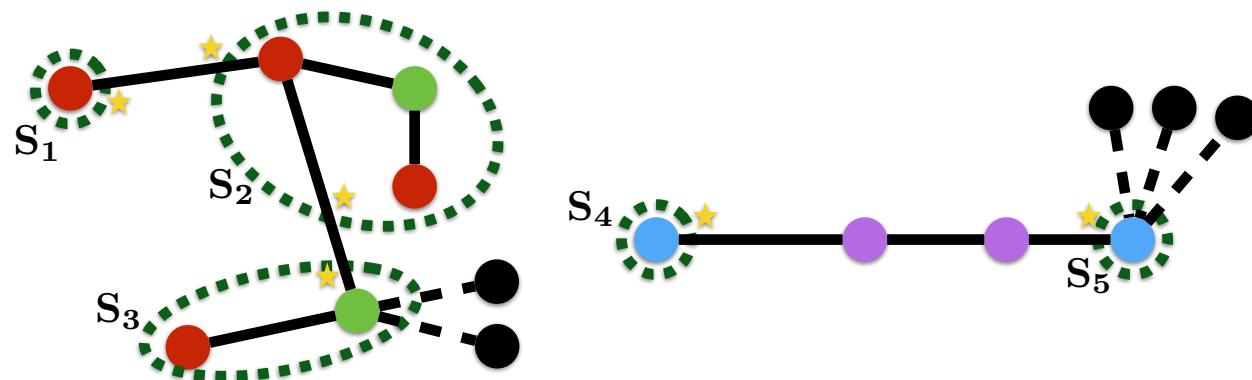
Properties of active sets

active sets are disjoint subsets of nodes containing terminals



Properties of F'

F' is a forest
and its leaves are terminals

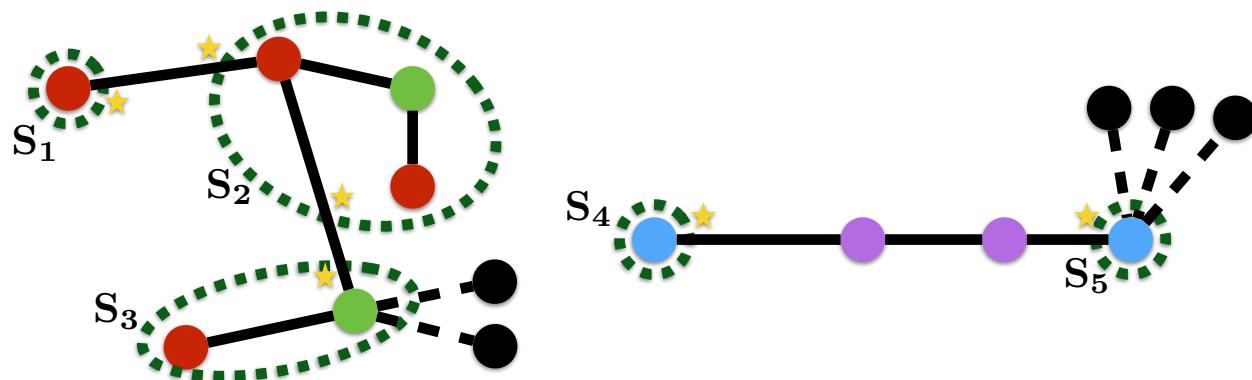


Joint properties of F' and of active sets

Consider a tree of F'

Assume it intersects some active sets

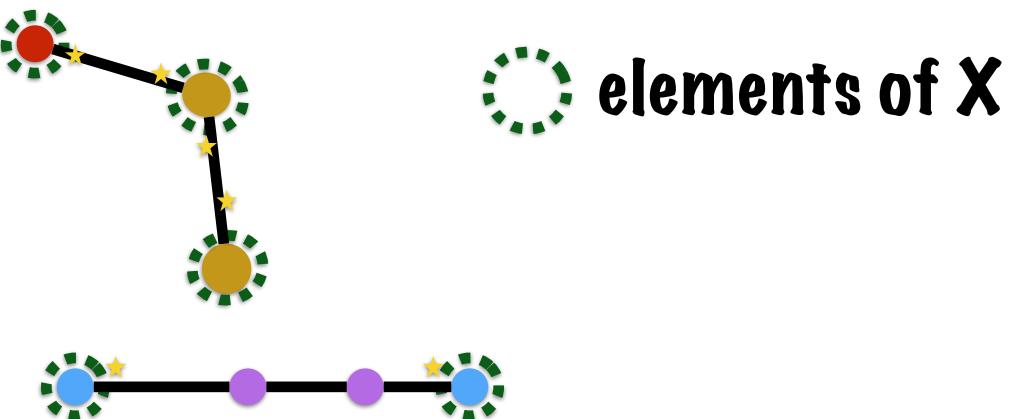
- the active sets are subtrees
- all the leaves are in active sets



Graph theory lemma

Given a tree T and
a subset X of its vertices

that includes all the leaves:
 $\sum_{v \in X} \deg(v) \leq 2 \cdot |X|$



Graph theory lemma

T tree, X subset of vertices
including all leaves:

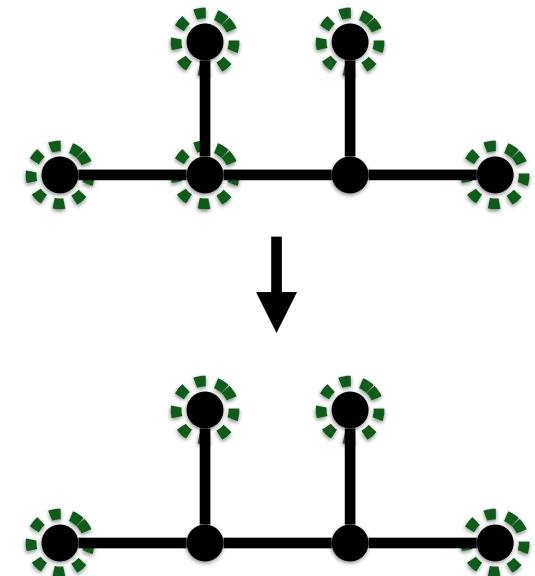
$$\sum_{v \in X} \deg(v) \leq 2 \cdot |X|$$

Proof: induction

- true for $X = V$

$$\sum_{v \in V} \deg(v) = 2|E| = 2(|V| - 1) < 2 \cdot |V|$$

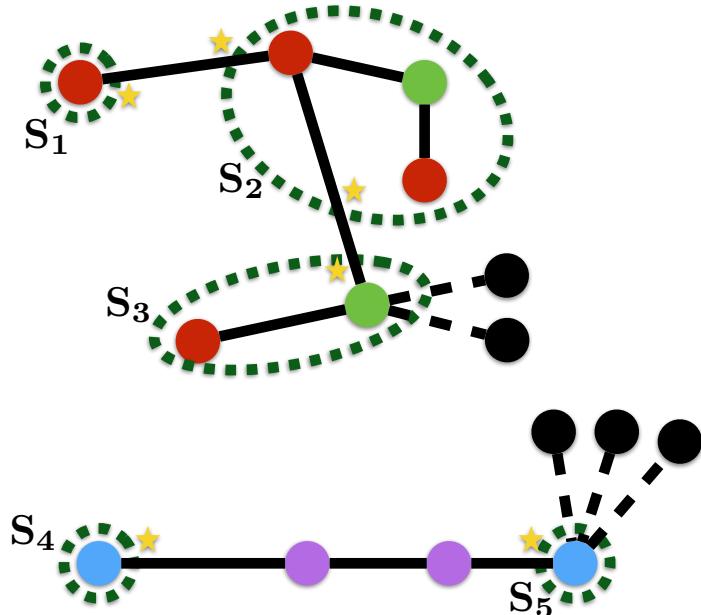
- if true for X then remains true when
removing an internal node from X , QED.



Lemma

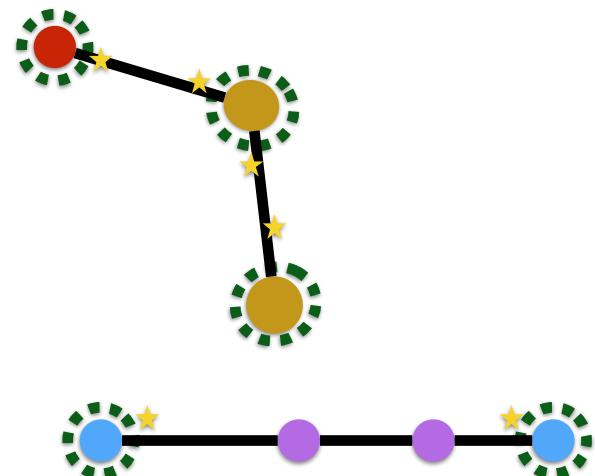
F' output forest

$$\frac{\sum_{S \text{ active}} |F' \cap \delta(S)|}{\#(\text{active sets})} \leq 2$$

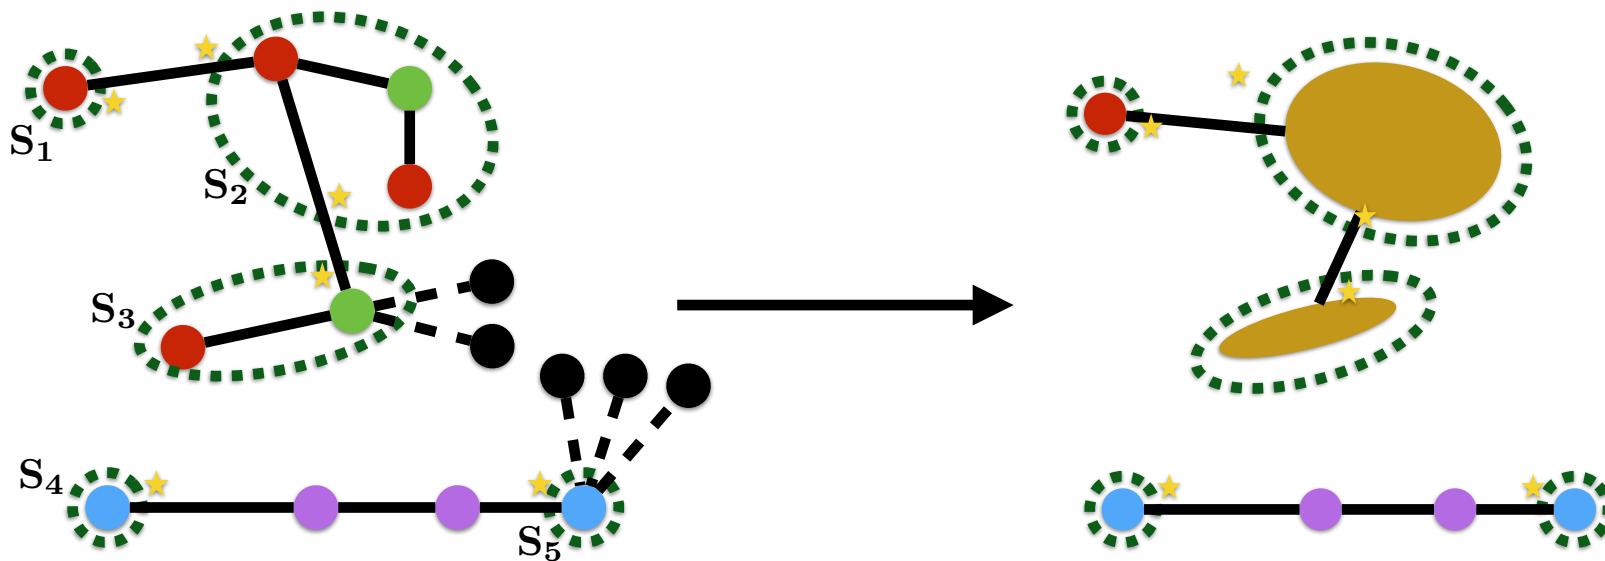


Graph theory lemma
 T tree, X subset of vertices including all leaves:

$$\frac{\sum_{v \in X} \deg(v)}{|X|} \leq 2$$



Lemma F' output forest: $\frac{\sum_{S \text{ active}} |F' \cap \delta(S)|}{\#(\text{active sets})} \leq 2$



Contracting active sets does not change

$\sum_{S \text{ active}} |F' \cap \delta(S)|$ **nor** $\#(\text{active sets})$

QED

Result:
**a 2-approximation algorithm
for the Steiner forest problem**

Comments

- extends: connectivity problems, and beyond
- primal-dual is greedy
- dual gives insight
- post-processing can be necessary
- combinatorial: LPs are guides, not computational tools

Steiner forest

