

Global Linear Models Part 2: The Perceptron Algorithm for Tagging

Michael Collins, Columbia University

Recap: Three Components of Global Linear Models

- ▶ **f** is a function that maps a structure (x, y) to a **feature vector** $\mathbf{f}(x, y) \in \mathbb{R}^d$
- ▶ **GEN** is a function that maps an input x to a set of **candidates** $\mathbf{GEN}(x)$
- ▶ \mathbf{v} is a parameter vector (also a member of \mathbb{R}^d)
- ▶ Training data is used to set the value of \mathbf{v}

Recap: Putting it all Together

- ▶ \mathcal{X} is set of sentences, \mathcal{Y} is set of possible outputs (e.g. trees)
- ▶ Need to learn a function $F : \mathcal{X} \rightarrow \mathcal{Y}$
- ▶ **GEN**, **f**, **v** define


$$F(x) = \arg \max_{y \in \mathbf{GEN}(x)} \mathbf{f}(x, y) \cdot \mathbf{v}$$

Choose the highest scoring candidate as the most plausible structure

- ▶ Given examples (x_i, y_i) , how to set **v**?

She announced a program to promote safety in trucks and vans

\Downarrow **GEN**



$\langle 1, 1, 3, 5 \rangle$

$\langle 2, 0, 0, 5 \rangle$

$\langle 1, 0, 1, 5 \rangle$

$\langle 0, 0, 3, 0 \rangle$

$\langle 0, 1, 0, 5 \rangle$

$\langle 0, 0, 1, 5 \rangle$

$\Downarrow \mathbf{f} \cdot \mathbf{v}$
13.6

$\Downarrow \mathbf{f} \cdot \mathbf{v}$
12.2


$\Downarrow \mathbf{f} \cdot \mathbf{v}$
12.1

$\Downarrow \mathbf{f} \cdot \mathbf{v}$
3.3

$\Downarrow \mathbf{f} \cdot \mathbf{v}$
9.4

$\Downarrow \mathbf{f} \cdot \mathbf{v}$
11.1

$\Downarrow \arg \max$



Recap: A Variant of the Perceptron Algorithm

Inputs: Training set (x_i, y_i) for $i = 1 \dots n$

Initialization: $\mathbf{v} = 0$

Define: $\textcolor{violet}{F}(x) = \operatorname{argmax}_{y \in \textcolor{red}{GEN}(x)} \textcolor{blue}{f}(x, y) \cdot \mathbf{v}$

Algorithm: For $t = 1 \dots T$, $i = 1 \dots n$
 $z_i = \textcolor{violet}{F}(x_i)$
 If $(z_i \neq y_i)$ $\mathbf{v} = \mathbf{v} + \textcolor{blue}{f}(x_i, y_i) - \textcolor{blue}{f}(x_i, z_i)$

Output: Parameters \mathbf{v}

Tagging Problems

TAGGING: Strings to Tagged Sequences

a b e e a f h j \Rightarrow a/C b/D e/C e/C a/D f/C h/D j/C

Example 1: Part-of-speech tagging

Profits/N soared/V at/P Boeing/N Co./N ,/, easily/ADV
topping/V forecasts/N on/P Wall/N Street/N ,/, as/P
their/POSS CEO/N Alan/N Mulally/N announced/V first/ADJ
quarter/N results/N ./.

Example 2: Named Entity Recognition

Profits/NA soared/NA at/NA Boeing/SC Co./CC ,/NA easily/NA
topping/NA forecasts/NA on/NA Wall/SL Street/CL ,/NA as/NA
their/NA CEO/NA Alan/SP Mulally/CP announced/NA first/NA
quarter/NA results/NA ./NA

Tagging using Global Linear Models

- ▶ Inputs x are sentences $w_{[1:n]} = \{w_1 \dots w_n\}$
- ▶ Define \mathcal{T} to be the set of possible tags
- ▶ **GEN** $(w_{[1:n]}) = \mathcal{T}^n$ i.e. all tag sequences of length n
- ▶ Note: The size of **GEN** is exponential in the sentence length
- ▶ How do we define **f**?

Representation: Histories

- ▶ A **history** is a 4-tuple $\langle t_{-2}, t_{-1}, w_{[1:n]}, i \rangle$
 - ▶ t_{-2}, t_{-1} are the previous two tags.
 - ▶ $w_{[1:n]}$ are the n words in the input sentence.
 - ▶ i is the index of the word being tagged
-

Hispaniola/**NNP** quickly/**RB** became/**VB** an/**DT** important/**JJ**
base/**??** from which Spain expanded its empire into the rest of the
Western Hemisphere .

- ▶ $t_{-2}, t_{-1} = \text{DT, JJ}$
- ▶ $w_{[1:n]} = \langle \text{Hispaniola, quickly, became, } \dots, \text{Hemisphere, .} \rangle$
- ▶ $i = 6$

Local Feature-Vector Representations

- ▶ Take a history/tag pair (h, t) .
 - ▶ $g_s(h, t)$ for $s = 1 \dots d$ are **local features** representing tagging decision t in context h .
-

Example: POS Tagging

- **Word/tag features**

$$g_{100}(h, t) = \begin{cases} 1 & \text{if current word } w_i \text{ is base and } t = \text{VB} \\ 0 & \text{otherwise} \end{cases}$$

$$g_{101}(h, t) = \begin{cases} 1 & \text{if current word } w_i \text{ ends in ing and } t = \text{VBG} \\ 0 & \text{otherwise} \end{cases}$$

- **Contextual Features**

$$g_{103}(h, t) = \begin{cases} 1 & \text{if } \langle t_{-2}, t_{-1}, t \rangle = \langle \text{DT}, \text{JJ}, \text{VB} \rangle \\ 0 & \text{otherwise} \end{cases}$$

A tagged sentence with n words has n history/tag pairs

Hispaniola/**NNP** quickly/**RB** became/**VB** an/**DT** important/**JJ** base/**NN**

History				Tag
t_{-2}	t_{-1}	$w_{[1:n]}$	i	t
*	*	$\langle \text{Hispaniola, quickly, ...}, \rangle$	1	NNP
*	NNP	$\langle \text{Hispaniola, quickly, ...}, \rangle$	2	RB
NNP	RB	$\langle \text{Hispaniola, quickly, ...}, \rangle$	3	VB
RB	VB	$\langle \text{Hispaniola, quickly, ...}, \rangle$	4	DT
VP	DT	$\langle \text{Hispaniola, quickly, ...}, \rangle$	5	JJ
DT	JJ	$\langle \text{Hispaniola, quickly, ...}, \rangle$	6	NN

A tagged sentence with n words has n history/tag pairs

Hispaniola/**NNP** quickly/**RB** became/**VB** an/**DT** important/**JJ** base/**NN**

History				Tag
t_{-2}	t_{-1}	$w_{[1:n]}$	i	t
*	*	$\langle \text{Hispaniola, quickly, ...} \rangle$	1	NNP
*	NNP	$\langle \text{Hispaniola, quickly, ...} \rangle$	2	RB
NNP	RB	$\langle \text{Hispaniola, quickly, ...} \rangle$	3	VB
RB	VB	$\langle \text{Hispaniola, quickly, ...} \rangle$	4	DT
VP	DT	$\langle \text{Hispaniola, quickly, ...} \rangle$	5	JJ
DT	JJ	$\langle \text{Hispaniola, quickly, ...} \rangle$	6	NN

Define global features through local features:

$$\mathbf{f}(t_{[1:n]}, w_{[1:n]}) = \sum_{i=1}^n g(h_i, t_i)$$

where t_i is the i 'th tag, h_i is the i 'th history

Global and Local Features

- Typically, local features are indicator functions, e.g.,

$$g_{101}(h, t) = \begin{cases} 1 & \text{if current word } w_i \text{ ends in ing and } t = \text{VBG} \\ 0 & \text{otherwise} \end{cases}$$

- and global features are then counts,

$f_{101}(w_{[1:n]}, t_{[1:n]}) =$ Number of times a word ending in ing is tagged as VBG in $(w_{[1:n]}, t_{[1:n]})$

Putting it all Together

- ▶ **GEN**($w_{[1:n]}$) is the set of all tagged sequences of length n
- ▶ **GEN**, **f**, **v** define

$$\begin{aligned} F(w_{[1:n]}) &= \arg \max_{t_{[1:n]} \in \text{GEN}(w_{[1:n]})} \mathbf{v} \cdot \mathbf{f}(w_{[1:n]}, t_{[1:n]}) \\ &= \arg \max_{t_{[1:n]} \in \text{GEN}(w_{[1:n]})} \mathbf{v} \cdot \sum_{i=1}^n g(h_i, t_i) \\ &= \arg \max_{t_{[1:n]} \in \text{GEN}(w_{[1:n]})} \sum_{i=1}^n \mathbf{v} \cdot g(h_i, t_i) \end{aligned}$$

Dynamic programming can be used to find the argmax!

A Variant of the Perceptron Algorithm

Inputs: Training set (x_i, y_i) for $i = 1 \dots n$

Initialization: $\mathbf{v} = 0$

Define: $\textcolor{violet}{F}(x) = \operatorname{argmax}_{y \in \textcolor{red}{GEN}(x)} \textcolor{blue}{f}(x, y) \cdot \mathbf{v}$

Algorithm: For $t = 1 \dots T$, $i = 1 \dots n$
 $z_i = \textcolor{violet}{F}(x_i)$
 If $(z_i \neq y_i)$ $\mathbf{v} = \mathbf{v} + \textcolor{blue}{f}(x_i, y_i) - \textcolor{blue}{f}(x_i, z_i)$

Output: Parameters \mathbf{v}

Training a Tagger Using the Perceptron Algorithm

Inputs: Training set $(w_{[1:n_i]}^i, t_{[1:n_i]}^i)$ for $i = 1 \dots n$.

Initialization: $\mathbf{v} = 0$

Algorithm: For $t = 1 \dots T, i = 1 \dots n$

$$z_{[1:n_i]} = \arg \max_{u_{[1:n_i]} \in \mathcal{T}^{n_i}} \mathbf{v} \cdot \mathbf{f}(w_{[1:n_i]}^i, u_{[1:n_i]})$$

$z_{[1:n_i]}$ can be computed with the dynamic programming (Viterbi) algorithm

If $z_{[1:n_i]} \neq t_{[1:n_i]}^i$ then

$$\mathbf{v} = \mathbf{v} + \mathbf{f}(w_{[1:n_i]}^i, t_{[1:n_i]}^i) - \mathbf{f}(w_{[1:n_i]}^i, z_{[1:n_i]})$$

Output: Parameter vector \mathbf{v} .

An Example

Say the correct tags for *i*'th sentence are

the/**DT** man/**NN** bit/**VBD** the/**DT** dog/**NN**

Under current parameters, output is

the/**DT** man/**NN** bit/**NN** the/**DT** dog/**NN**

Assume also that features track: (1) all bigrams; (2) word/tag pairs

Parameters incremented:

$\langle \text{NN}, \text{VBD} \rangle, \langle \text{VBD}, \text{DT} \rangle, \langle \text{VBD} \rightarrow \text{bit} \rangle$

Parameters decremented:

$\langle \text{NN}, \text{NN} \rangle, \langle \text{NN}, \text{DT} \rangle, \langle \text{NN} \rightarrow \text{bit} \rangle$

Experiments

- ▶ Wall Street Journal part-of-speech tagging data

Perceptron = 2.89% error, Log-linear tagger = 3.28% error

- ▶ [Ramshaw and Marcus, 1995] NP chunking data

Perceptron = 93.63% accuracy, Log-linear tagger = 93.29% accuracy