

▼ Predicting Bounding Boxes

Welcome to Course 3, Week 1 Programming Assignment!

In this week's assignment, you'll build a model to predict bounding boxes around images.

- You will use transfer learning on any of the pre-trained models available in Keras.
- You'll be using the [Caltech Birds - 2010](#) dataset.

How to submit your work

Notice that there is not a "submit assignment" button in this notebook.

To check your work and get graded on your work, you'll train the model, save it and then upload the model to Coursera for grading.

- [Initial steps](#)
 - [0.1 Set up your Colab](#)
 - [0.2 Set up the data location](#)
 - [0.3 Choose the GPU Runtime](#)
 - [0.4 Mount your drive](#)
 - [0.5 Imports](#)
- [1. Visualization Utilities](#)
 - [1.1 Bounding Boxes Utilities](#)
 - [1.2 Data and Predictions Utilities](#)
- [2. Preprocessing and Loading the Dataset](#)
 - [2.1 Preprocessing Utilities](#)
 - [2.2 Visualize the prepared Data](#)
 - [2.3 Loading the Dataset](#)
- [3. Define the Network](#)
 - [Exercise 1](#)
 - [Exercise 2](#)
 - [Exercise 3](#)
 - [Exercise 4](#)
 - [Exercise 5](#)
- [4. Training the Model](#)
 - [Prepare to train the model](#)
 - [Exercise 6](#)
 - [Fit the model to the data](#)
 - [Exercise 7](#)
- [5. Validate the Model](#)
 - [5.1 Loss](#)
 - [5.2 Save your Model](#)

- [5.3 Plot the Loss Function](#)
- [5.4 Evaluate performance using IoU](#)
- [6. Visualize Predictions](#)
- [7. Upload your model for grading](#)

0. Initial steps

0.1 Set up your Colab

- As you cannot save the changes you make to this colab, you have to make a copy of this notebook in your own drive and run that.
- You can do so by going to File -> Save a copy in Drive.
- Close this colab and open the copy which you have made in your own drive. Then continue to the next step to set up the data location.

Set up the data location

A copy of the dataset that you'll be using is stored in a publicly viewable Google Drive folder. You'll want to add a shortcut to it to your own Google Drive.

- Go to this google drive folder named [TF3 C3 W1 Data](#)
- Next to the folder name "TF3 C3 W1 Data" (at the top of the page beside "Shared with me"), hover your mouse over the triangle to reveal the drop down menu.
- Use the drop down menu to select "Add shortcut to Drive" A pop-up menu will open up.
- In the pop-up menu, "My Drive" is selected by default. Click the ADD SHORTCUT button. This should add a shortcut to the folder TF3 C3 W1 Data within your own google drive at the location `content/drive`.
- To verify, go to the left-side menu and click on "My Drive". Scroll through your files to look for the shortcut TF3 C3 W1 Data.

Please make sure this happens, as you'll be reading the data for this notebook from this folder.

0.3 Choose the GPU Runtime

- Make sure your runtime is **GPU** (*not* CPU or TPU). And if it is an option, make sure you are using *Python 3*. You can select these settings by going to Runtime -> Change runtime type -> Select the above mentioned settings and then press SAVE

▼ 0.4 Mount your drive

Please run the next code cell and follow these steps to mount your Google Drive so that it can be accessed by this Colab.

- Run the code cell below. A web link will appear below the cell.

- Please click on the web link, which will open a new tab in your browser, which asks you to choose your google account.
- Choose your google account to login.
- The page will display "Google Drive File Stream wants to access your Google Account". Please click "Allow".
- The page will now show a code (a line of text). Please copy the code and return to this Colab.
- Paste the code the textbox that is labeled "Enter your authorization code:" and hit <Enter>
- The text will now say "Mounted at [/content/drive/](#)"
- Please look at the files explorer of this Colab (left side) and verify that you can navigate to `drive/MyDrive/TF3 C3 W1 Data/caltech_birds2010/0.1.1` . If the folder is not there, please redo the steps above and make sure that you're able to add the shortcut to the hosted dataset.

```
from google.colab import drive
drive.mount('/content/drive/', force_remount=True)
```

```
Mounted at /content/drive/
```

▼ 0.5 Imports

```
import os, re, time, json
import PIL.Image, PIL.ImageFont, PIL.ImageDraw
import numpy as np
import tensorflow as tf
from matplotlib import pyplot as plt
import tensorflow_datasets as tfds
import cv2
```

Store the path to the data.

- Remember to follow the steps to set up the data location (above) so that you'll have a shortcut to the data in your Google Drive.

```
data_dir = "/content/drive/My Drive/TF3 C3 W1 Data/"
```

▼ 1. Visualization Utilities

▼ 1.1 Bounding Boxes Utilities

We have provided you with some functions which you will use to draw bounding boxes around the birds in the image .

- `draw_bounding_box_on_image`: Draws a single bounding box on an image.
- `draw_bounding_boxes_on_image`: Draws multiple bounding boxes on an image.
- `draw_bounding_boxes_on_image_array`: Draws multiple bounding boxes on an array of images.

```
def draw_bounding_boxes_on_image(image, boxes, labels, scores):
    """Draws bounding boxes on an image.
    Args:
        image: A PIL image.
        boxes: A list of bounding boxes in the format (x1, y1, x2, y2).
        labels: A list of labels for the bounding boxes.
        scores: A list of scores for the bounding boxes.
    """
```

```
def draw_bounding_box_on_image(image, ymin, xmin, ymax, xmax, color=(255, 0, 0), thickness=5):
```

```
    """
```

Adds a bounding box to an image.

Bounding box coordinates can be specified in either absolute (pixel) or normalized coordinates by setting the `use_normalized_coordinates` argument.

Args:

image: a PIL.Image object.

ymin: ymin of bounding box.

xmin: xmin of bounding box.

ymax: ymax of bounding box.

xmax: xmax of bounding box.

color: color to draw bounding box. Default is red.

thickness: line thickness. Default value is 4.

```
    """
```

```
image_width = image.shape[1]
```

```
image_height = image.shape[0]
```

```
cv2.rectangle(image, (int(xmin), int(ymin)), (int(xmax), int(ymax)), color, thickness)
```

```
def draw_bounding_boxes_on_image(image, boxes, color=[], thickness=5):
```

```
    """
```

Draws bounding boxes on image.

Args:

image: a PIL.Image object.

boxes: a 2 dimensional numpy array of [N, 4]: (ymin, xmin, ymax, xmax).

The coordinates are in normalized format between [0, 1].

color: color to draw bounding box. Default is red.

thickness: line thickness. Default value is 4.

Raises:

ValueError: if boxes is not a [N, 4] array

```
    """
```

```
boxes_shape = boxes.shape
```

```
if not boxes_shape:
```

```
    return
```

```
if len(boxes_shape) != 2 or boxes_shape[1] != 4:
```

```
    raise ValueError('Input must be of size [N, 4]')
```

```
for i in range(boxes_shape[0]):
```

```
    draw_bounding_box_on_image(image, boxes[i, 1], boxes[i, 0], boxes[i, 3],
                               boxes[i, 2], color[i], thickness)
```

```
def draw_bounding_boxes_on_image_array(image, boxes, color=[], thickness=5):
```

```
    """
```

Draws bounding boxes on image (numpy array).

Args:

image: a numpy array object.

boxes: a 2 dimensional numpy array of [N, 4]: (ymin, xmin, ymax, xmax).

The coordinates are in normalized format between [0, 1].

color: color to draw bounding box. Default is red.

thickness: line thickness. Default value is 4.

display_str_list_list: a list of strings for each bounding box.

Raises:

ValueError: if boxes is not a [N, 4] array

```

ValueError: If boxes is not a [N, 4] array

"""

draw_bounding_boxes_on_image(image, boxes, color, thickness)

return image

```

▼ 1.2 Data and Predictions Utilities

We've given you some helper functions and code that are used to visualize the data and the model's predictions.

- `display_digits_with_boxes`: This displays a row of "digit" images along with the model's predictions for each image.
- `plot_metrics`: This plots a given metric (like loss) as it changes over multiple epochs of training.

```

# Matplotlib config
plt.rc('image', cmap='gray')
plt.rc('grid', linewidth=0)
plt.rc('xtick', top=False, bottom=False, labelsizе='large')
plt.rc('ytick', left=False, right=False, labelsizе='large')
plt.rc('axes', facecolor='F8F8F8', titlesize="large", edgecolor='white')
plt.rc('text', color='a8151a')
plt.rc('figure', facecolor='F0F0F0')# Matplotlib fonts
MATPLOTLIB_FONT_DIR = os.path.join(os.path.dirname(plt.__file__), "mpl-data/fonts/ttf")

# utility to display a row of digits with their predictions
def display_digits_with_boxes(images, pred_bboxes, bboxes, iou, title, bboxes_normalized=False):

    n = len(images)

    fig = plt.figure(figsize=(20, 4))
    plt.title(title)
    plt.yticks([])
    plt.xticks([])

    for i in range(n):
        ax = fig.add_subplot(1, 10, i+1)
        bboxes_to_plot = []
        if (len(pred_bboxes) > i):
            bbox = pred_bboxes[i]
            bbox = [bbox[0] * images[i].shape[1], bbox[1] * images[i].shape[0], bbox[2] * images[i].shape[1],
                    bbox[3] * images[i].shape[0]]
            bboxes_to_plot.append(bbox)

        if (len(bboxes) > i):
            bbox = bboxes[i]
            if bboxes_normalized == True:
                bbox = [bbox[0] * images[i].shape[1],bbox[1] * images[i].shape[0], bbox[2] * images[i].shape[1],
                        bbox[3] * images[i].shape[0]]
            bboxes_to_plot.append(bbox)

    img_to_draw = draw_bounding_boxes_on_image_array(image=images[i], boxes=np.asarray(bboxes_to_plot))
    plt.xticks([])
    plt.yticks([])

    plt.imshow(img_to_draw)

```

```

if len(iou) > i :
    color = "black"
    if (iou[i][0] < iou_threshold):
        color = "red"
    ax.text(0.2, -0.3, "iou: %s" %(iou[i][0]), color=color, transform=ax.transAxes)

```

```

# utility to display training and validation curves
def plot_metrics(metric_name, title, ylim=5):
    plt.title(title)
    plt.ylim(0,ylim)
    plt.plot(history.history[metric_name],color='blue',label=metric_name)
    plt.plot(history.history['val_' + metric_name],color='green',label='val_' + metric_name)

```

▼ 2. Preprocess and Load the Dataset

▼ 2.1 Preprocessing Utilities

We have given you some helper functions to pre-process the image data.

▼ read_image_tfds

- Resizes image to (224, 224)
- Normalizes image
- Translates and normalizes bounding boxes

```

def read_image_tfds(image, bbox):
    image = tf.cast(image, tf.float32)
    shape = tf.shape(image)

    factor_x = tf.cast(shape[1], tf.float32)
    factor_y = tf.cast(shape[0], tf.float32)

    image = tf.image.resize(image, (224, 224,))

    image = image/127.5
    image -= 1

    bbox_list = [bbox[0] / factor_x ,
                  bbox[1] / factor_y,
                  bbox[2] / factor_x ,
                  bbox[3] / factor_y]

    return image, bbox_list

```

▼ read_image_with_shape

This is very similar to `read_image_tfds` except it also keeps a copy of the original image (before pre-processing) and returns this as well.

- Makes a copy of the original image.
- Resizes image to (224, 224)

- Normalizes image
- Translates and normalizes bounding boxes

```
def read_image_with_shape(image, bbox):
    original_image = image

    image, bbox_list = read_image_tfds(image, bbox)

    return original_image, image, bbox_list
```

▼ read_image_tfds_with_original_bbox

- This function reads image from data
- It also denormalizes the bounding boxes (it undoes the bounding box normalization that is performed by the previous two helper functions.)

```
def read_image_tfds_with_original_bbox(data):
    image = data["image"]
    bbox = data["bbox"]

    shape = tf.shape(image)
    factor_x = tf.cast(shape[1], tf.float32)
    factor_y = tf.cast(shape[0], tf.float32)

    bbox_list = [bbox[1] * factor_x ,
                  bbox[0] * factor_y,
                  bbox[3] * factor_x,
                  bbox[2] * factor_y]
    return image, bbox_list
```

▼ dataset_to_numpy_util

This function converts a dataset into numpy arrays of images and boxes.

- This will be used when visualizing the images and their bounding boxes

```
def dataset_to_numpy_util(dataset, batch_size=0, N=0):

    # eager execution: loop through datasets normally
    take_dataset = dataset.shuffle(1024)

    if batch_size > 0:
        take_dataset = take_dataset.batch(batch_size)

    if N > 0:
        take_dataset = take_dataset.take(N)

    if tf.executing_eagerly():
        ds_images, ds_bboxes = [], []
        for images, bboxes in take_dataset:
            ds_images.append(images.numpy())
            ds_bboxes.append(bboxes.numpy())
```

```
return (np.array(ds_images), np.array(ds_bboxes))
```

▼ dataset_to_numpy_with_original_bboxes_util

- This function converts a `dataset` into numpy arrays of
 - original images
 - resized and normalized images
 - bounding boxes
- This will be used for plotting the original images with true and predicted bounding boxes.

```
def dataset_to_numpy_with_original_bboxes_util(dataset, batch_size=0, N=0):  
  
    normalized_dataset = dataset.map(read_image_with_shape)  
    if batch_size > 0:  
        normalized_dataset = normalized_dataset.batch(batch_size)  
  
    if N > 0:  
        normalized_dataset = normalized_dataset.take(N)  
  
    if tf.executing_eagerly():  
        ds_original_images, ds_images, ds_bboxes = [], [], []  
  
    for original_images, images, bboxes in normalized_dataset:  
        ds_images.append(images.numpy())  
        ds_bboxes.append(bboxes.numpy())  
        ds_original_images.append(original_images.numpy())  
  
    return np.array(ds_original_images), np.array(ds_images), np.array(ds_bboxes)
```

▼ 2.2 Visualize the images and their bounding box labels

Now you'll take a random sample of images from the training and validation sets and visualize them by plotting the corresponding bounding boxes.

Visualize the **training** images and their bounding box labels

```
def get_visualization_training_dataset():  
    dataset, info = tfds.load("caltech_birds2010", split="train", with_info=True, data_dir=data_dir,  
    print(info)  
    visualization_training_dataset = dataset.map(read_image_tfds_with_original_bbox,  
                                                num_parallel_calls=16)  
    return visualization_training_dataset  
  
visualization_training_dataset = get_visualization_training_dataset()  
  
(visualization_training_images, visualization_training_bboxes) = dataset_to_numpy_util(visualization_training_dataset)  
display_digits_with_boxes(np.array(visualization_training_images), np.array([]), np.array(visualization_training_bboxes))
```



```

tfds.core.DatasetInfo(
    name='caltech_birds2010',
    version=0.1.1,
    description='Caltech-UCSD Birds 200 (CUB-200) is an image dataset with photos
of 200 bird species (mostly North American). The total number of
categories of birds is 200 and there are 6033 images in the 2010
dataset and 11,788 images in the 2011 dataset.
Annotations include bounding boxes, segmentation labels.',
    homepage='http://www.vision.caltech.edu/visipedia/CUB-200.html',
    features=FeaturesDict({
        'bbox': BBoxFeature(shape=(4,), dtype=tf.float32),
        'image': Image(shape=(None, None, 3), dtype=tf.uint8),
        'image/filename': Text(shape=(), dtype=tf.string),
        'label': ClassLabel(shape=(), dtype=tf.int64, num_classes=200),
        'label_name': Text(shape=(), dtype=tf.string),
        'segmentation_mask': Image(shape=(None, None, 1), dtype=tf.uint8),
    }),
    total_num_examples=6033,
    splits={
        'test': 3033,
        'train': 3000,
    },
    supervised_keys=('image', 'label'),
    citation="""@techreport{WelinderEtal2010,
Author = {P. Welinder and S. Branson and T. Mita and C. Wah and F. Schroff and S. Belongie
Institution = {California Institute of Technology},
Number = {CNS-TR-2010-001},
Title = {{Caltech-UCSD Birds 200}},
Year = {2010}
}""",
    redistribution_info=,
)

```

/usr/local/lib/python3.7/dist-packages/ipykernel_launcher.py:18: VisibleDeprecationWarning: Cr

training images and their bboxes



Visualize the **validation** images and their bounding boxes

```

def get_visualization_validation_dataset():
    dataset = tfds.load("caltech_birds2010", split="test", data_dir=data_dir, download=False)
    visualization_validation_dataset = dataset.map(read_image_tfds_with_original_bbox, num_parallel=
    return visualization_validation_dataset

```

```


visualization_validation_dataset = get_visualization_validation_dataset()

```

```

(visualization_validation_images, visualization_validation_bboxes) = dataset_to_numpy_util(visualiza
display_digits_with_boxes(np.array(visualization_validation_images), np.array([]), np.array(visualiz

```



▼ 2.3 Load and prepare the datasets for the model

These next two functions read and prepare the datasets that you'll feed to the model.

- They use `read_image_tfds` to resize, and normalize each image and its bounding box label.
- They performs shuffling and batching.
- You'll use these functions to create `training_dataset` and `validation_dataset`, which you will give to the model that you're about to build.

```
BATCH_SIZE = 64
```

```
def get_training_dataset(dataset):
    dataset = dataset.map(read_image_tfds, num_parallel_calls=16)
    dataset = dataset.shuffle(512, reshuffle_each_iteration=True)
    dataset = dataset.repeat()
    dataset = dataset.batch(BATCH_SIZE)
    dataset = dataset.prefetch(-1)
    return dataset

def get_validation_dataset(dataset):
    dataset = dataset.map(read_image_tfds, num_parallel_calls=16)
    dataset = dataset.batch(BATCH_SIZE)
    dataset = dataset.repeat()
    return dataset

training_dataset = get_training_dataset(visualization_training_dataset)
validation_dataset = get_validation_dataset(visualization_validation_dataset)
```

▼ 3. Define the Network

Bounding box prediction is treated as a "regression" task, in that you want the model to output numerical values.

- You will be performing transfer learning with **MobileNet V2**. The model architecture is available in TensorFlow Keras.
- You'll also use pretrained `'imagenet'` weights as a starting point for further training. These weights are also readily available
- You will choose to retrain all layers of **MobileNet V2** along with the final classification layers.

Note: For the following exercises, please use the TensorFlow Keras Functional API (as opposed to the Sequential API).

▼ Exercise 1

Please build a feature extractor using MobileNetV2.

- First, create an instance of the mobilenet version 2 model
 - Please check out the documentation for [MobileNetV2](#)
 - Set the following parameters:
 - `input_shape`: (height, width, channel): input images have height and width of 224 by 224, and have red, green and blue channels.
 - `include_top`: you do not want to keep the "top" fully connected layer, since you will customize your model for the current task.
 - `weights`: Use the pre-trained 'imagenet' weights.
- Next, make the feature extractor for your specific inputs by passing the `inputs` into your mobilenet model.
 - For example, if you created a model object called `some_model` and have inputs stored in `x`, you'd invoke the model and pass in your inputs like this: `some_model(x)` to get the feature extractor for your given inputs `x`.

Note: please use `mobilenet_v2` and not `mobile_net` or `mobile_net_v3`

```
def feature_extractor(inputs):  
    ### YOUR CODE HERE ###  
  
    # Create a mobilenet version 2 model object  
    mobilenet_model = tf.keras.applications.MobileNetV2(input_shape=(224,224,3), include_top=False,  
  
    # pass the inputs into this model object to get a feature extractor for these inputs  
    feature_extractor = mobilenet_model(inputs)  
  
    ### END CODE HERE ###  
  
    # return the feature_extractor  
    return feature_extractor
```

▼ Exercise 2

Next, you'll define the dense layers to be used by your model.

You'll be using the following layers

- [GlobalAveragePooling2D](#): pools the features .
- [Flatten](#): flattens the pooled layer.
- [Dense](#): Add two dense layers:

- A dense layer with 1024 neurons and a relu activation.
- A dense layer following that with 512 neurons and a relu activation.

Note: Remember, please build the model using the Functional API syntax (as opposed to the Sequential API).

```
def dense_layers(features):
    ### YOUR CODE HERE ###

    # global average pooling 2D layer.
    x = tf.keras.layers.GlobalAveragePooling2D()(features)

    # flatten layer
    x = tf.keras.layers.Flatten()(x)

    # 1024 Dense layer, with relu
    x = tf.keras.layers.Dense(1024, activation="relu")(x)

    # 512 Dense layer, with relu
    x = tf.keras.layers.Dense(512, activation="relu")(x)

    ### END CODE HERE ###

    return x
```

▼ Exercise 3

Now you'll define a layer that outputs the bounding box predictions.

- You'll use a [Dense](#) layer.
- Remember that you have *4 units* in the output layer, corresponding to (xmin, ymin, xmax, ymax).
- The prediction layer follows the previous dense layer, which is passed into this function as the variable `x`.
- For grading purposes, please set the `name` parameter of this Dense layer to be `'bounding_box'`

```
def bounding_box_regression(x):
    ### YOUR CODE HERE ###

    # Dense layer named `bounding_box`
    bounding_box_regression_output = tf.keras.layers.Dense(units=4, name="bounding_box")(x)

    ### END CODE HERE ###

    return bounding_box_regression_output
```

▼ Exercise 4

Now, you'll use those functions that you have just defined above to construct the model.

- `feature_extractor(inputs)`
- `dense_layers(features)`
- `bounding_box_regression(x)`

Then you'll define the model object using [Model](#). Set the two parameters:

- inputs
- outputs

```
def final_model(inputs):
    ### YOUR CODE HERE ###

    # features
    feature_cnn = feature_extractor(inputs)

    # dense layers
    last_dense_layer = dense_layers(feature_cnn)

    # bounding box
    bounding_box_output = bounding_box_regression(last_dense_layer)

    # define the TensorFlow Keras model using the inputs and outputs to your model
    model = tf.keras.Model(inputs=inputs, outputs=[bounding_box_output], name='model')

    ### END CODE HERE ###

    return model
```

▼ Exercise 5

Define the input layer, define the model, and then compile the model.

- inputs: define an [Input](#) layer
 - Set the `shape` parameter. Check your definition of `feature_extractor` to see the expected dimensions of the input image.
- model: use the `final_model` function that you just defined to create the model.
- compile the model: Check the [Model](#) documentation for how to compile the model.
 - Set the `optimizer` parameter to Stochastic Gradient Descent using [SGD](#)
 - When using SGD, set the `momentum` to 0.9 and keep the default learning rate.
 - Set the loss function of SGD to mean squared error (see the SGD documentation for an example of how to choose mean squared error loss).

```
def define_and_compile_model():

    ### YOUR CODE HERE ###

    # define the input layer
    inputs = tf.keras.layers.Input(shape=(224, 224, 3))

    # create the model
    model = final_model(inputs)

    # compile your model
    model.compile(optimizer = tf.keras.optimizers.SGD(momentum=0.9), loss = 'mean_squared_error')
```

END CODE HERE

return model

Run the cell below to define your model and print the model summary.

```
# define your model
model = define_and_compile_model()
# print model layers
model.summary()
```

Downloading data from https://storage.googleapis.com/tensorflow/keras-applications/mobilenet_v2_1.00_224 [=====] - 0s 0us/step
Model: "model"

Layer (type)	Output Shape	Param #
input_1 (InputLayer)	[(None, 224, 224, 3)]	0
mobilenetv2_1.00_224 (Function)	(None, 7, 7, 1280)	2257984
global_average_pooling2d (GlobalAveragePooling2D)	(None, 1280)	0
flatten (Flatten)	(None, 1280)	0
dense (Dense)	(None, 1024)	1311744
dense_1 (Dense)	(None, 512)	524800
bounding_box (Dense)	(None, 4)	2052
Total params: 4,096,580		
Trainable params: 4,062,468		
Non-trainable params: 34,112		

Your expected model summary:

Layer (type)	Output Shape	Param #
input_1 (InputLayer)	[(None, 224, 224, 3)]	0
mobilenetv2_1.00_224 (Function)	(None, 7, 7, 1280)	2257984
global_average_pooling2d (GlobalAveragePooling2D)	(None, 1280)	0
flatten (Flatten)	(None, 1280)	0
dense (Dense)	(None, 1024)	1311744
dense_1 (Dense)	(None, 512)	524800
bounding_box (Dense)	(None, 4)	2052
Total params: 4,096,580		
Trainable params: 4,062,468		
Non-trainable params: 34,112		

Train the Model

4.1 Prepare to Train the Model

You'll fit the model here, but first you'll set some of the parameters that go into fitting the model.

- EPOCHS: You'll train the model for 50 epochs
- BATCH_SIZE: Set the `BATCH_SIZE` to an appropriate value. You can look at the ungraded labs from this week for some examples.
- `length_of_training_dataset`: this is the number of training examples. You can find this value by getting the length of `visualization_training_dataset`.
 - Note: You won't be able to get the length of the object `training_dataset`. (You'll get an error message).
- `length_of_validation_dataset`: this is the number of validation examples. You can find this value by getting the length of `visualization_validation_dataset`.
 - Note: You won't be able to get the length of the object `validation_dataset`.
- `steps_per_epoch`: This is the number of steps it will take to process all of the training data.
 - If the number of training examples is not evenly divisible by the batch size, there will be one last batch that is not the full batch size.

- Try to calculate the number steps it would take to train all the full batches plus one more batch containing the remaining training examples. There are a couples ways you can calculate this.
 - You can use regular division `/` and import `math` to use `math.ceil()` [Python math module docs](#)
 - Alternatively, you can use `//` for integer division, `%` to check for a remainder after integer division, and an `if` statement.
- `validation_steps`: This is the number of steps it will take to process all of the validation data. You can

▼ Exercise 6

```
# You'll train 50 epochs
EPOCHS = 50

### START CODE HERE ###

# Choose a batch size
BATCH_SIZE = 128

# Get the length of the training set
length_of_training_dataset = len(visualization_training_dataset)

# Get the length of the validation set
length_of_validation_dataset = len(visualization_validation_dataset)

# Get the steps per epoch (may be a few lines of code)
steps_per_epoch = length_of_training_dataset//BATCH_SIZE
if length_of_training_dataset % BATCH_SIZE > 0:
    steps_per_epoch += 1

# get the validation steps (per epoch) (may be a few lines of code)
validation_steps = length_of_validation_dataset//BATCH_SIZE
if length_of_validation_dataset % BATCH_SIZE > 0:
    validation_steps += 1

### END CODE HERE
```

4.2 Fit the model to the data

Check out the parameters that you can set to fit the [Model](#). Please set the following parameters.

- `x`: this can be a tuple of both the features and labels, as is the case here when using a `tf.Data` dataset.
 - Please use the variable returned from `get_training_dataset()`.
 - Note, don't set the `y` parameter when the `x` is already set to both the features and labels.
- `steps_per_epoch`: the number of steps to train in order to train on all examples in the training dataset.
- `validation_data`: this is a tuple of both the features and labels of the validation set.
 - Please use the variable returned from `get_validation_dataset()`
- `validation_steps`: teh number of steps to go through the validation set, batch by batch.
- `epochs`: the number of epochs.

If all goes well your model's training will start.

Exercise 7

YOUR CODE HERE

```
# Fit the model, setting the parameters noted in the instructions above.
history = model.fit(training_dataset,
                    steps_per_epoch=steps_per_epoch, validation_data=validation_dataset, validation_steps=validation_steps)
```

END CODE HERE

```
24/24 [-----] - 17s 714ms/step - loss: 0.0040 - val_loss: 0.0540
Epoch 22/50
24/24 [=====] - 17s 715ms/step - loss: 0.0044 - val_loss: 0.0540
Epoch 23/50
24/24 [=====] - 17s 707ms/step - loss: 0.0042 - val_loss: 0.0511
Epoch 24/50
24/24 [=====] - 17s 710ms/step - loss: 0.0040 - val_loss: 0.0501

Epoch 25/50
24/24 [=====] - 17s 694ms/step - loss: 0.0039 - val_loss: 0.0455
Epoch 26/50
24/24 [=====] - 17s 702ms/step - loss: 0.0038 - val_loss: 0.0460
Epoch 27/50
24/24 [=====] - 17s 725ms/step - loss: 0.0037 - val_loss: 0.0430
Epoch 28/50
24/24 [=====] - 17s 726ms/step - loss: 0.0039 - val_loss: 0.0430
Epoch 29/50
24/24 [=====] - 17s 701ms/step - loss: 0.0035 - val_loss: 0.0378
Epoch 30/50
24/24 [=====] - 17s 710ms/step - loss: 0.0035 - val_loss: 0.0365
Epoch 31/50
24/24 [=====] - 17s 709ms/step - loss: 0.0037 - val_loss: 0.0383
Epoch 32/50
24/24 [=====] - 17s 721ms/step - loss: 0.0034 - val_loss: 0.0375
Epoch 33/50
24/24 [=====] - 17s 706ms/step - loss: 0.0034 - val_loss: 0.0349
Epoch 34/50
24/24 [=====] - 17s 724ms/step - loss: 0.0032 - val_loss: 0.0329
Epoch 35/50
24/24 [=====] - 17s 719ms/step - loss: 0.0034 - val_loss: 0.0310
Epoch 36/50
24/24 [=====] - 17s 716ms/step - loss: 0.0030 - val_loss: 0.0327
Epoch 37/50
24/24 [=====] - 17s 709ms/step - loss: 0.0032 - val_loss: 0.0310
Epoch 38/50
24/24 [=====] - 17s 720ms/step - loss: 0.0029 - val_loss: 0.0328
Epoch 39/50
24/24 [=====] - 17s 709ms/step - loss: 0.0030 - val_loss: 0.0284
Epoch 40/50
24/24 [=====] - 17s 721ms/step - loss: 0.0027 - val_loss: 0.0303
Epoch 41/50
24/24 [=====] - 17s 707ms/step - loss: 0.0030 - val_loss: 0.0261
Epoch 42/50
24/24 [=====] - 17s 712ms/step - loss: 0.0031 - val_loss: 0.0304
Epoch 43/50
24/24 [=====] - 17s 715ms/step - loss: 0.0029 - val_loss: 0.0274
Epoch 44/50
24/24 [=====] - 17s 726ms/step - loss: 0.0027 - val_loss: 0.0286
Epoch 45/50
24/24 [=====] - 17s 704ms/step - loss: 0.0032 - val_loss: 0.0269
```

```
Epoch 46/50
24/24 [=====] - 18s 746ms/step - loss: 0.0025 - val_loss: 0.0267
Epoch 47/50
24/24 [=====] - 17s 726ms/step - loss: 0.0025 - val_loss: 0.0274
Epoch 48/50
24/24 [=====] - 17s 727ms/step - loss: 0.0027 - val_loss: 0.0262
Epoch 49/50
24/24 [=====] - 17s 712ms/step - loss: 0.0026 - val_loss: 0.0250
Epoch 50/50
24/24 [=====] - 17s 708ms/step - loss: 0.0027 - val_loss: 0.0254
```

▼ 5. Validate the Model

5.1 Loss

You can now evaluate your trained model's performance by checking its loss value on the validation set.

```
loss = model.evaluate(validation_dataset, steps=validation_steps)
print("Loss: ", loss)
```

```
24/24 [=====] - 6s 247ms/step - loss: 0.0254
Loss: 0.025356346741318703
```

▼ 5.2 Save your Model for Grading

When you have trained your model and are satisfied with your validation loss, please you save your model so that you can upload it to the Coursera classroom for grading.

```
# Please save your model
model.save("birds.h5")
```


```
# And download it using this shortcut or from the "Files" panel to the left
from google.colab import files
```

```
files.download("birds.h5")
```

▼ 5.3 Plot Loss Function

You can also plot the loss metrics.

```
plot_metrics("loss", "Bounding Box Loss", ylim=0.2)
```



5.4 Evaluate performance using IoU

You can see how well your model predicts bounding boxes on the validation set by calculating the Intersection-over-union (IoU) score for each image.

- You'll find the IoU calculation implemented for you.
- Predict on the validation set of images.
- Apply the `intersection_over_union` on these predicted bounding boxes.

```
def intersection_over_union(pred_box, true_box):

    xmin_pred, ymin_pred, xmax_pred, ymax_pred = np.split(pred_box, 4, axis = 1)
    xmin_true, ymin_true, xmax_true, ymax_true = np.split(true_box, 4, axis = 1)

    #Calculate coordinates of overlap area between boxes
    xmin_overlap= np.maximum(xmin_pred, xmin_true)
    xmax_overlap = np.minimum(xmax_pred, xmax_true)
    ymin_overlap = np.maximum(ymin_pred, ymin_true)
    ymax_overlap = np.minimum(ymax_pred, ymax_true)

    #Calculates area of true and predicted boxes
    pred_box_area = (xmax_pred - xmin_pred) * (ymax_pred - ymin_pred)
    true_box_area = (xmax_true - xmin_true) * (ymax_true - ymin_true)

    #Calculates overlap area and union area.
    overlap_area = np.maximum((xmax_overlap - xmin_overlap),0) * np.maximum((ymax_overlap - ymin_overlap),0)
    union_area = (pred_box_area + true_box_area) - overlap_area

    # Defines a smoothing factor to prevent division by 0
    smoothing_factor = 1e-10

    #Updates iou score
    iou = (overlap_area + smoothing_factor) / (union_area + smoothing_factor)

    return iou

#Makes predictions
original_images, normalized_images, normalized_bboxes = dataset_to_numpy_with_original_bboxes_util(v
predicted_bboxes = model.predict(normalized_images, batch_size=32)

#Calculates IOU and reports true positives and false positives based on IOU threshold
iou = intersection_over_union(predicted_bboxes, normalized_bboxes)
iou_threshold = 0.5

print("Number of predictions where iou > threshold(0.5): %s" % (iou_threshold, (iou >= iou_threshold).sum()))
print("Number of predictions where iou < threshold(0.5): %s" % (iou_threshold, (iou < iou_threshold).sum()))

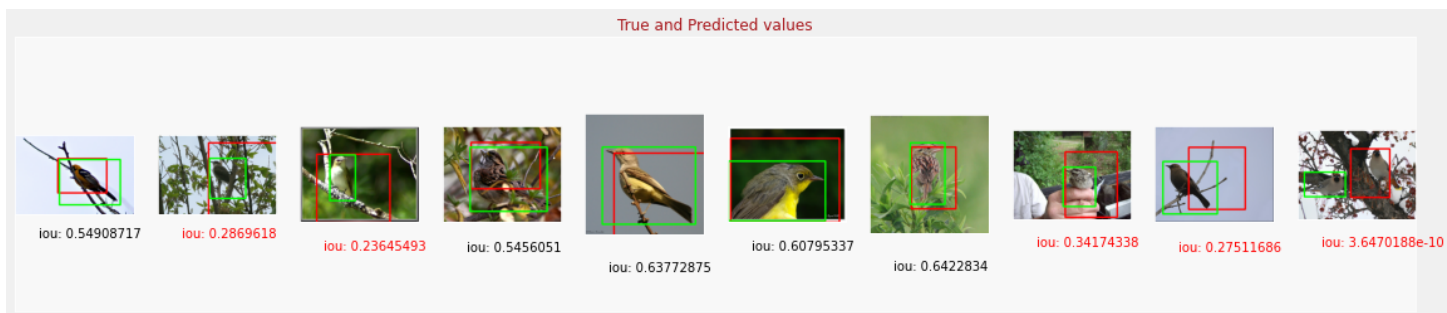
/usr/local/lib/python3.7/dist-packages/ipykernel_launcher.py:18: VisibleDeprecationWarning: Cr
Number of predictions where iou > threshold(0.5): 205
Number of predictions where iou < threshold(0.5): 295
```

6. Visualize Predictions

Lastly, you'll plot the predicted and ground truth bounding boxes for a random set of images and visually see how well you did!

```
n = 10
indexes = np.random.choice(len(predicted_bboxes), size=n)

iou_to_draw = iou[indexes]
norm_to_draw = original_images[indexes]
display_digits_with_boxes(original_images[indexes], predicted_bboxes[indexes], normalized_bboxes[indexes])
```



7 Upload your model for grading

Please return to the Coursera classroom and find the section that allows you to upload your 'birds.h5' model for grading. Good luck!

