

Theoretical Computer Science Stack Exchange is a question and answer site for theoretical computer scientists and researchers in related fields. It's 100% free, no registration required.

Take the 2-minute tour ×

Testing/Identifying a Topological Sorting

You're given a set of n Directed Acyclic Graphs G_1, G_2, \dots, G_n over the same set of m vertices V . You're also given a permutation of the set of vertices (v_1, v_2, \dots, v_m) . What is the best algorithm that could identify the graphs among G_1, G_2, \dots, G_n that have (v_1, v_2, \dots, v_m) as a topological sort? Could someone test whether (v_1, v_2, \dots, v_m) is a topological sort of a DAG G over V in sub-linear time?

ds.algorithms graph-algorithms directed-acyclic-graph

edited Jan 20 '11 at 4:57

 Hsien-Chih Chang 張顯之
5,715 2 31 69

asked Jan 20 '11 at 4:09

Steve

- 6 Are you able to build a data structure based on the set of graphs before being presented with the ordering of the vertices? You need to look at all n graphs and all $m - 1$ edges in the ordering, so unless you're allowed to preprocess the graphs somehow, it doesn't seem like you could beat linear time. – [mjqxxx](#) Jan 20 '11 at 5:29

Hsien-Chih Chang, what would be a good pre-processing technique that can allow a better solution? Some type of hashing? I guess you can beat linear time if can approximate the solution (probabilistic algorithm). – [user2471](#) Jan 21 '11 at 0:37

@user2471: As I said in your previous answer, this post is written by @Steve, not me ;) – [Hsien-Chih Chang](#) 張顯之 Jan 21 '11 at 1:00

Sorry Hsien-Chih Chang, my question was meant for everyone :) – [user2471](#) Jan 21 '11 at 2:04

@user2471, no need to apologize! Hope someone who is familiar to this question will post a nice answer :D – [Hsien-Chih Chang](#) 張顯之 Jan 21 '11 at 5:28

2 Answers

This can be done in nearly linear time.

Let the permutation be $\pi = (v_1, \dots, v_m)$, and let $k = k(\pi)$ be the number of steps needed to check a single edge (u, v) against π . It is then enough to check that each of the M_i edges of G_i is compatible with π , which can be done in $O(kM_i)$ time, or $O(k \sum M_i)$ overall.

By preprocessing π one can reduce k down to two lookups in an array containing m entries each of $\log m$ size, and a comparison between two $(\log m)$ -bit entries in the array; the array element $a[w]$ contains the index of w in π considered as an ordered list. This means that $k = O(\log m)$ yielding $O((\log m) \sum M_i)$ time overall for the upper bound.

As @mjqxxx points out, every edge of every graph may be relevant. This creates a lower bound of $\Omega(K \sum M_i)$ steps, where K is the least amount of work that needs to be done for every graph edge; it is possible that some approaches can amortize the cost so that $K = o(\log m)$. This is still going to be $\Omega(\sum M_i)$ at best, so there is not much of a gap left.

edited Apr 23 '11 at 15:52

answered Apr 20 '11 at 23:21

 András Salamon
11.8k 3 40 117

What algorithm can be used to make $k = \log m$. Is it suffix trees? – [vincent mathew](#) Mar 30 at 11:03

Trivial Method:

```
GS = {G1,G2...G3}
for v in (v1,v2,...vm)
    remove all graphs from GS where indegree(v) != 0
    remove v and attached edges from remaining GS
```

It is not as fast as you want. But it solves one problem that "there can be multiple valid topological orderings of DAG". And finding them all is not a good idea.

answered Jan 20 '11 at 5:46

