

[Personal](#) [Open source](#) [Business](#) [Explore](#)[Pricing](#) [Blog](#) [Support](#)[This repository](#)[Sign in](#)[Sign up](#)[contravariant](#) / **pandas\_dplyr**[Watch](#)

3

[★ Star](#)

50

[Fork](#)

14

[Code](#)[Issues](#) 0[Pull requests](#) 0[Pulse](#)[Graphs](#)Branch: **master** ▾**pandas\_dplyr** / pandas\_dplyr.md[Find file](#)[Copy path](#)**contravariant** Update pandas\_dplyr.md

1951da5 Jul 20, 2015

2 contributors



591 lines (567 sloc) 18 KB

[Raw](#)[Blame](#)[History](#)

# Super-short super-basic Data Munging in R and Python

*Giuseppe Paleologo* [paleologo@gmail.com](mailto:paleologo@gmail.com)

2015-04-01

In October 2012 I wrote a document to help translate between R and python pandas idioms for data manipulation. Back then, It was reasonable to cover a good fraction of the operations in the two languages. Today that is no longer the case. Pandas has tripled in size since I first covered its capabilities and is now nearing 200K lines of code. Meanwhile, R has seen intense

development of `data.table` and the release of `dplyr`, `tidyr` and `magrittr`, steps toward a "grammar of data". Pandas and `dplyr` take different approaches to data manipulation. Pandas takes the "maximalist" approach: it works with 1, 2, and 3-dimensional arrays; and with hierarchical indices, one can represent tree-like data structures (e.g., nested lists); moreover, it allows one to align and operate on numerical arrays. Finally, it has plotting capabilities as well. `dplyr` takes the "minimalist" approach. It works with 2-D tables only, with a small set of verbs that are easily composable with each other (as well as with base R functions). As a result, the manipulations are easy to write and read. Plotting is delegated to other libraries; so is numerical algebra; so is time series analysis. To learn well the data manipulation capabilities alone of these two languages requires a significant time investment. To reflect these changes, I chose to focus on manipulation of tabular data alone, therefore using a subset of pandas for Python, and of `dplyr` + `tidyr` in R. The choice of the latter is motivated by its balance of conceptual elegance, performance and fast rate of adoption. Pandas' documentation has a section comparing R to pandas, but the idioms presented here are more modern, regular and concise.

I make no claims of completeness, but hope that readers will find this useful to go back and forth between the languages.

## Preliminaries

---

### Python

```
# python preliminaries
import numpy as np
randn = np.random.randn
from pandas import *
```

### R

```
# R preliminaries
library(dplyr)
library(tidyr)
library(data.table) # used only for fread()
```

## Reading a dataframe from file or URL

---

The functions `read_csv` and `fread` have similar performance. They attempt to infer from data field separator and field type.

### Python

```
DF = read_csv(filepath)
```

### R

```
DF <- fread(filepath)
```

Both pandas and R offer functions to access DMBS (though SQLAlchemy for Pandas, specific packages for R) and a variety of formats (e.g., Json).

## Generating a dataframe from existing data

---

### Python

```
# From dictionary
x = {'city': ['Rome', 'New York', 'Moscow'],
     'continent': ['Europe', 'America', 'Europe'],
     'inhabitants': [8.4e6, 2.8e6, 11.5e6]}
DF = DataFrame(x)

# from recarray
m = np.zeros((2,), dtype=[('A', 'i4'), ('B', 'f4'), ('C', 'a10')])
```

```
m[:] = [(1,2., 'Hello'), (2,3., "World")]
DataFrame(m)

# from ndarray
m = np.arange(12.).reshape((3, 4))
DataFrame(m, columns=['a', 'b', 'c', 'd'])
```

## R

```
# native
DF <- data.frame(city=c('New York', 'Rome', 'Moscow'),
                  continent = c('Europe', 'America', 'Europe'),
                  inhabitants=c(8.4e6, 2.8e6, 11.5e6))

# From list
x <- list(city=c('New York', 'Rome', 'Moscow'),
          continent = c('Europe', 'America', 'Europe'),
          inhabitants=c(8.4e6, 2.8e6, 11.5e6))
DF <- as.data.frame(x)

# from matrix
m <- matrix(1:12, ncol=2)
colnames(m) <- letters[1:3]
as.data.frame(m)
```

# Getting row names

---

## Python

```
DF.index
```

## R

```
row.names(DF)
DF %>% row.names      # magrittr version
```

## Filtering rows

---

### Python

```
DF.ix[DF.inhabitants > 5e6]
DF.ix[[x in ['Europe', 'Asia'] for x in DF.continent]]
```

### R

```
filter(DF, inhabitants > 5e6)
filter(DF, continent %in% c('Europe', 'Asia'))
```

## Filtering columns

---

### Python

```
DF[['city', 'continent']]
```

### R

```
select(DF, city, continent)
DF[, c('city', 'continent')]
```

## Deleting columns

---

### Python

```
del DF['city']
```

### R

```
DF[, 'city'] <- NULL  
select(DF, -city)
```

## Jointly filtering rows and columns

---

### Python

```
DF.ix[[True, False, True], ['city', 'continent']]
```

### R

```
# base R; dplyr combines filter and select  
DF[c(TRUE, FALSE, TRUE), c('city', 'continent')]
```

## Setting Fields to NA

---

### Python

```
DF.ix[[True, False, True], ['city', 'continent']] = np.nan
```

## R

```
DF[c(TRUE, FALSE, TRUE), c('city', 'continent')] = NA
```

# Checking for (non)missing values

---

## Python

```
DF.isnull()      # missing  
isnull(DF)       # alternative  
DF.notnull()     # non missing  
notnull(DF)      # alternative
```

## R

```
is.na(DF)        # missing  
!is.na(DF)       # non missing
```

# dropping missing values

---

## Python

```
DF.dropna()
```

## R

```
na.omit(s)
```

# Replacing missing values

---

## Python

```
DF.fillna(-1)
```

## R

```
DF[is.na(DF)] <- -1
```

# Head and Tail

---

## Python

```
DF.head(2) # default is 5 rows  
DF.tail(2)
```

## R

```
head(DF, 2) # default is 6 rows  
tail(DF, 2)
```



# Updating/adding columns

---

## Python

```
DF['birthdate'] = [-621, 1625, 1147]
```

## R

```
DF[, 'birthdate'] <- c(-621, 1625, 1147)  
DF %<>% mutate(birthdate = c(-621, 1625, 1147))
```

# Renaming columns

---

## Python

```
DF.columns = [x.upper() for x in list(DF.columns)]
```

## R

```
names(DF) <- toupper(names(DF))      # base R  
DF %<>% set_names(toupper(names(.)))
```

# Concatenating rows

---

## Python

```
DF_list = [DF, DF]  
concat(DF_list)
```

## R

```
rbind_all(DF, DF)  
rbind_list(list(DF, DF))
```

N.B.: Pandas allows data with different column names and types to be concatenated. Moreover, pandas allows the creation of a hierarchical index for the combined data frame. See the section "Introduction to Pandas indices".

# Remove duplicated rows

---

## Python

```
s = DataFrame({'a' : [0., .0, 1., 2.]})  
s.duplicated()  
s.drop_duplicates()
```

## R

```
s = data.frame(a = c(0., .0, 1., 2.))  
distinct(s)  # in base R, unique(s)
```

# Joining

---

## Python

```
merge(DF1, DF2, how='left', on=['colname1', 'colname2'])
merge(DF1, DF2, how='right', on=['colname1', 'colname2'])
merge(DF1, DF2, how='inner', on=['colname1', 'colname2'])
merge(DF1, DF2, how='outer', on=['colname1', 'colname2'])
```

## R

```
inner_join(DF1, DF2, by = c('colname1', 'colname2'))
left_join (DF1, DF2, by = c('colname1', 'colname2'))
right_join(DF1, DF2, by = c('colname1', 'colname2'))
full_join (DF1, DF2, by = c('colname1', 'colname2'))
```

dplyr has also a `semi_join` and `anti_join`.

## Sorting

---

### Python

```
DF.sort('city', ascending=False)
DF.sort('city', ascending=True)
DF.sort(['city', 'continent'], ascending=True)
```

### R

```
DF %>% arrange(desc(city))
DF %>% arrange(city)
DF %>% arrange(city, continent)
```

# Summarizing

---

## Python

```
DF.describe()
```

## R

```
summary(DF)
```

# Getting dimensions

---

## Python

```
DF.shape
```

## R

```
dim(DF)
```

# Converting to arrays for numerical computation

---

Pandas and R/dplyr differ substantially on this account. In R, it is strongly inadvisable to perform binary operations on data

frames. The user should convert the data to a suitable n-way array and then perform operations. Conversely, Pandas is fully able to correctly perform operations on the underlying numpy object, with the added benefit of automatic alignment. mplyr is a package that does alignment on n-way arrays in R.

## Python

```
DF[['inhabitants', 'birthdate']].as_matrix()

# an example of how Pandas takes care of automatic alignment
m1 = np.arange(12.).reshape((3, 4))
df1 = DataFrame(m1, columns=list('abcd'))
m2 = (np.arange(20)+5).reshape((4, 5))
df2 = DataFrame(m2, columns=list('abcde'))
df1 + df2    # takes union of indices and columns, NaN applied

as.matrix(DF[, c('inhabitants', 'birthdate')])
```

## Splitting, applying, combining

---

### Python

```
# converts the index to column fields
DF_grouped = DF.groupby('city', as_index=False)
DF_grouped.agg({'C' : np.sum, 'D' : lambda x: np.std(x, ddof=1)})
DF_grouped['C'].agg(np.sum)
grouped['C'].agg({'result1' : np.sum, 'result2' : np.mean})
# sugared expression
DF_grouped.std()
```

# Converting a dataframe from wide to long format

---

## Python

```
cheese = DataFrame({'first' : ['John', 'Mary'],  
                    'last'  : ['Doe', 'Bo'],  
                    'height' : [5.5, 6.0],  
                    'weight' : [130, 150]})  
pd.melt(cheese, id_vars=['first', 'last'])
```

## R

```
cheese <- data.frame(first = c('John', 'Mary'),  
                     last  = c('Doe', 'Bo'),  
                     height = c(5.5, 6.0),  
                     weight = c(130, 150))  
cheese %>% gather(feature, value, -first, -last)
```

# Converting a dataframe from long to wide format

---

## Python

```
df = DataFrame({'Animal': ['Animal1', 'Animal2', 'Animal3', 'Animal2',  
                           'Animal1', 'Animal2', 'Animal3'],  
               'FeedType': ['A', 'B', 'A', 'A', 'B', 'B', 'A'],  
               'Amount': [10, 7, 4, 2, 5, 6, 2], })  
df.pivot_table(values='Amount', index='Animal', columns='FeedType', aggfunc='sum')
```

## R

```
df <- data.frame(  
  Animal = c('Animal1', 'Animal2', 'Animal3', 'Animal2', 'Animal1',  
             'Animal2', 'Animal3'),  
  FeedType = c('A', 'B', 'A', 'A', 'B', 'B', 'A'),  
  Amount = c(10, 7, 4, 2, 5, 6, 2)  
)  
# with tidyr  
df %>% spread(Animal, FeedType)  
# using base R  
with(df, tapply(Amount, list(Animal, FeedType), sum, na.rm = TRUE))
```

## Casting to an Array

---

### Python

```
df = pd.DataFrame({x': np.random.uniform(1., 168., 12),  
                  'y': np.random.uniform(7., 334., 12),  
                  'z': np.random.uniform(1.7, 20.7, 12),  
                  'month': [5,6,7]*4, 'week': [1,2]*6})  
mdf = pd.melt(df, id_vars=['month', 'week'])
```

### R

```
pd.pivot_table(mdf,  
               values='value',  
               index=['variable', 'week'],  
               columns=['month'], aggfunc=np.mean)  
df <- data.frame(x = runif(12, 1, 168),  
                y = runif(12, 7, 334),  
                z = runif(12, 1.7, 20.7),  
                month = rep(c(5,6,7),4),
```

```
      week = rep(c(1,2), 6)
    )
    mdf <- melt(df, id=c("month", "week"))
    acast(mdf, week ~ month ~ variable, mean)
```

## Applying functions to an entire data frame or column

---

### Python

```
np.exp(df[['x']])
```

### R

```
exp(df$x)
```

## Applying elementwise functions

---

```
s = Series(["little", "red", "fox"])
s.map(len)      # notice that map(len, s) returns a list
```

### R

```
s <- c("little", "red", "fox")
sapply(s, nchar)  # nchar(s) would have worked here
```

Notice the difference: pandas aligns by name, taking the union of indices. Base R does not take names into account. In R,



the typical process would be to join the vectors in a data frame and then operating on the columns.

## A self-contained example: baby names

I close with a concrete example: the "baby names" data set made famous by Martin Wattenberg and Fernanda Viégas. The python code below is taken verbatim from Wes McKinney's book "Pandas for Data Analysis". The R code is a translation of the same analysis in R. A few distinguishing features stand out. First, python used method extensively whereas R uses functions. It is possible to chain methods but readability is not greatly enhanced. In R, all the operations can be performed by a single chain. Second, R delegates all the plotting functions to ggplot2, which uses a syntax similar to dplyr (but with a "+"; ggvis, the successor to ggplot, uses the familiar %>%"). Third, all R analysis uses "long data frame". This is an instance of "tidy data". Python uses wide format data frames for plotting (not unlike R's the ones `matplot()` would require. Lastly, the same few functions show up in R over and over: `group_by()`, `summarize()`, `is`, `reindex`, `arrange()`, `mutate()`, `filter()`. Pandas has a larger vocabulary, with `groupby`, `apply`, `sortindex`, `searchsorted`, `unstack`, `map`, `pivot_table` methods; many of which take further arguments; but it is at the same time slightly less verbose.

## Read Files

### Python

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt

years = range(1880, 2011)
pieces = []
columns = ['name', 'sex', 'births']
for year in years:
    path = 'names/yob%d.txt' % year
    #print(path)
    frame = pd.read_csv(path, names=columns)
```

```
frame['year'] = year
pieces.append(frame)
names = pd.concat(pieces, ignore_index=True)
```

## R

```
library(dplyr)
library(ggplot2)
library(data.table)
library(stringr)
library(magrittr)

years <- 1880:2011
path <- sprintf('names/yob%d.txt', years)
columns <- c('name', 'sex', 'births')

reader <- function(yr){
  x<- fread(sprintf('names/yob%d.txt', yr), data.table=FALSE) %>%
    set_names(columns)
  x$year <- yr
  x
}
babynames <- lapply(years, reader) %>% bind_rows
```

## Plots births by sex and year

### Python

```
total_births = names.pivot_table('births', rows='year', cols='sex', aggfunc=sum)
total_births.plot(title='Total births by sex and year')
```

## R

```
babynames %>%  
  group_by(year, sex) %>%  
  summarize(births=sum(births)) %>%  
  ggplot(aes(x=year, y=births, color=sex)) + geom_line()
```

## Fraction of total names, within sex and year

### Python

```
def add_prop(group):  
    # Integer division floors  
    births = group.births.astype(float)  
    group['prop'] = births / births.sum()  
    return group  
  
names = names.groupby(['year', 'sex']).apply(add_prop)  
names.head()
```

## R

```
babynames %<>%  
  group_by(year, sex) %>%  
  mutate(prop=births/sum(births))
```

## Top 1000 names by year, sex

### Python

```
def get_top1000(group):  
    return group.sort_index(by='births', ascending=False)[:1000]  
grouped = names.groupby(['year', 'sex'])  
top1000 = grouped.apply(get_top1000)  
top1000.index = np.arange(len(top1000))  
top1000.head()
```

## R

```
top1000 <- babynames %>%  
  group_by(year, sex) %>%  
  arrange(desc(births)) %>%  
  filter(min_rank(desc(births)) <= 1000)  
top1000
```

# Plots the number of babies named John, Harry, Mary, Marilyn over time

## Python

```
total_births = top1000.pivot_table('births', rows='year',  
                                    cols='name', aggfunc=sum)  
subset = total_births[['John', 'Harry', 'Mary', 'Marilyn']]  
subset.plot(subplots=True, figsize=(12, 10),  
            grid=False, title="Number of births per year")
```

## R

```
babynames %>%  
  filter(name %in% c('John', 'Harry', 'Mary', 'Marilyn')) %>%  
  group_by(year, name) %>%
```

```
summarize(births=sum(births)) %>%
ggplot(aes(x=year, y=births)) + geom_line() + facet_grid(name ~ .)
```

## Plots the proportion of the top 1000 names as a percentage of total

### Python

```
table = top1000.pivot_table('prop', rows='year', cols='sex', aggfunc=sum)

table.plot(title='Sum of table1000.prop by year and sex',
           yticks=np.linspace(0, 1.2, 13),
           xticks=range(1880, 2020, 10))
```

### R

```
top1000 %>%
  group_by(year, sex) %>%
  summarize(prop=sum(prop)) %>%
  ggplot(aes(x=year, y=prop, color=sex)) +
    geom_line() +
    ggtitle('Sum of table1000.prop by year and sex') +
    scale_y_continuous(limits=c(0, 1))
```

## How many boy names comprise 50% of the total in 2010?

### Python

```
boys = top1000[top1000.sex == 'M']
df = boys[boys.year == 2010]
prop_cumsum = df.sort_index(by='prop', ascending=False).prop.cumsum()
```

```
prop_cumsum.values.searchsorted(0.5)
```

## R

```
top1000 %>%
  filter(year == 2010, sex == 'M') %>%
  arrange(desc(births)) %>%
  mutate(totprop = cumsum(prop)) %>%
  filter(totprop <= .50) %>%
  nrow
```

# Plots number of most popular names used by 50% of boys and girls over time

## Python

```
def get_quantile_count(group, q=0.5):
    group = group.sort_index(by='prop', ascending=False)
    return group.prop.cumsum().values.searchsorted(q) + 1
diversity = top1000.groupby(['year', 'sex']).apply(get_quantile_count)
diversity = diversity.unstack('sex')
diversity.plot(title="Number of popular names in top 50%")
```

## R

```
get_quantile_count <- function(x, qtle=0.5)
  x %>% sort(decreasing=TRUE) %>% cumsum %>% {.<= qtle} %>% sum

top1000 %>%
  group_by(year, sex) %>%
```

```
summarize(No.names = get_quantile_count(prop)) %>%
ggplot(aes(x=year, y=No.names, color=sex)) +
  geom_line() +
  ggtitle('Number of popular names in top 50%')
```

## Plot the distribution of names by last letter for three time snapshots

### Python

```
get_last_letter = lambda x: x[-1]
last_letters = names.name.map(get_last_letter)
last_letters.name = 'last_letter'
table = names.pivot_table('births', rows=last_letters, cols=['sex', 'year'], aggfunc=sum)

subtable = table.reindex(columns=[1910, 1960, 2010], level='year')
letter_prop = subtable / subtable.sum().astype(float)

fig, axes = plt.subplots(2, 1, figsize=(10, 8))
letter_prop['M'].plot(kind='bar', rot=0, ax=axes[0], title='Male', legend=True)
letter_prop['F'].plot(kind='bar', rot=0, ax=axes[1], title='Female', legend=False)
```

### R

```
letter_count <- babynames %>%
  mutate(last_letter = str_sub(name, start=-1L, end=-1L )) %>%
  group_by(year, sex, last_letter) %>%
  summarise(count=sum(births))

letter_prop <- letter_count %>%
  filter(year %in% c(1910, 1960, 2010)) %>%
  group_by(year, sex) %>%
  mutate(letter_prop=count/sum(count))
```

```
letter_prop %>%  
  ggplot(aes(x=last_letter, y=letter_prop, fill=as.factor(year))) +  
    geom_bar(stat='identity', position=position_dodge()) +  
    facet_grid(sex ~ .)
```

## Plots the proportion of boy names ending in 'd', 'n', and 'y' over time

### Python

```
letter_prop = table / table.sum().astype(float)  
dny_ts = letter_prop.ix[['d', 'n', 'y'], 'M'].T  
  
dny_ts.plot()  
# table.ix(last_letter=='d')
```

### R

```
letter_count %>%  
  filter(sex=='M') %>%  
  group_by(year) %>%  
  mutate(letter_count=sum(count), letter_prop=count/sum(count)) %>%  
  filter(last_letter %in% c('d', 'n', 'y')) %>%  
  ggplot(aes(x=year, y=letter_prop, color=last_letter)) +  
    geom_line()
```



