




R news and tutorials contributed by (573) R bloggers

- [Home](#)
- [About](#)
- [RSS](#)
- [add your blog!](#)
- [R jobs](#) 
- [Contact us](#)

## Welcome!

Follow @rbloggers  20.2K

Here you will find daily **news and tutorials about R**, contributed by over 573 bloggers. There are many ways to **follow us** - [By e-mail:](#)



[On Facebook:](#)



**If you are an R blogger yourself** you are invited to [add your own R content feed to this site](#) (Non-English R bloggers should add themselves- [here](#))

## [Jobs for R-users](#)

- [Senior Analyst for Health Insurance Intelligence with R & Shiny \(@Zurich\)](#)
- [Scientist / Sr. Scientist](#)
- [Sr. Manager of Content Analytics at Amazon.com \(@Seattle\)](#)
- [Data Scientist @Burlington](#)
- [Associate Director of Admissions Strategy and Analysis \(@Maine\)](#)

Search & Hit Enter

## Popular

## Searches

- [heatmap](#)
- [web scraping](#)
- [maps](#)
- [hadoop](#)
- [shiny](#)
- [twitterR](#)
- [ggplot2](#)
- [boxplot](#)
- [trading](#)
- [animation](#)
- [excel](#)
- [Time series](#)
- [finance](#)
- [ggplot](#)
- [latex](#)
- [pca](#)
- [googlevis](#)
- [eclipse](#)
- [quantmod](#)
- [rstudio](#)
- [market research](#)
- [tutorial](#)
- [how to import image file to R](#)
- [remdr](#)
- [knitr](#)
- [rattle](#)
- [gis](#)
- [coplot](#)
- [sweave](#)
- [ecdf](#)

## Recent Posts

- [The Workflow of Infinite Shame, and other stories from the R Summit](#)
- [What is a good Sharpe ratio?](#)
- [Open Analytics @ UseR! 2015](#)
- [When did same-sex marriage laws become effective by state?](#)
- [Bio7 2.2 for Windows, Linux and MacOSX released](#)
- [Rperform: Obtaining quantitative metrics across R package versions](#)
- [A bit about Win-Vector LLC](#)
- [R at Microsoft](#)
- [Introduction to Monte Carlo methods with R and Bayesian Essentials with R](#)
- [An Attempt to Understand Boosting Algorithm\(s\)](#)
- [remote update, paper & poster](#)
- [Geomorph update 2.1.6 Now Available!](#)
- [KDD Cup 2015: The story of how I built hundreds of predictive models....And got so close, yet so far away from 1st](#)

- [place!](#)
- [Working with the RStudio CRAN logs](#)
- [Learn to Map Census Data in R](#)

## Other sites

- [Jobs for R-users](#)
- [Statistics of Israel](#)
- [SAS blogs](#)

# Text Mining the Complete Works of William Shakespeare

September 5, 2013

By [andrew](#)

(This article was first published on [Exegetic Analytics » R](#), and kindly contributed to [R-bloggers](#))

I am starting a new project that will require some serious text mining. So, in the interests of bringing myself up to speed on the tm package, I thought I would apply it to the Complete Works of William Shakespeare and just see what falls out.

The first order of business was getting my hands on all that text. Fortunately it is available from a number of sources. I chose to use [Project Gutenberg](#).

```
> TEXTFILE = "data/pg100.txt"
> if (!file.exists(TEXTFILE)) {
+   download.file("http://www.gutenberg.org/cache/
+ }
> shakespeare = readLines(TEXTFILE)
> length(shakespeare)
[1] 124787
```

That's quite a solid chunk of data: 124787 lines. Let's take a closer look.

```
> head(shakespeare)
[1] "The Project Gutenberg EBook of The Complete Wor
[2] "William Shakespeare"
[3] ""
[4] "This eBook is for the use of anyone anywhere at
[5] "almost no restrictions whatsoever. You may cop
[6] "re-use it under the terms of the Project Gutenb
> tail(shakespeare)
[1] "http://www.gutenberg.org/2/4/6/8/24689" ""
[3] "An alternative method of locating eBooks:" "htt
[5] "" ""***
```

There seems to be some header and footer text. We will want to get rid of that! Using a text editor I checked to see how many lines were occupied with metadata and then removed them before concatenating all of the lines into a single long, long, long string.

```
> shakespeare = shakespeare[-(1:173)]
> shakespeare = shakespeare[-(124195:length(shakespe
>
> shakespeare = paste(shakespeare, collapse = " ")
> nchar(shakespeare)
[1] 5436541
```

While I had the text open in the editor I noticed that sections in the document were separated by the following text:

```
<<THIS ELECTRONIC VERSION OF THE COMPLETE WORKS OF W
SHAKESPEARE IS COPYRIGHT 1990-1993 BY WORLD LIBRARY,
PROVIDED BY PROJECT GUTENBERG ETEXT OF ILLINOIS BENE
WITH PERMISSION. ELECTRONIC AND MACHINE READABLE CO
DISTRIBUTED SO LONG AS SUCH COPIES (1) ARE FOR YOUR
PERSONAL USE ONLY, AND (2) ARE NOT DISTRIBUTED OR US
COMMERCIAL. PROHIBITED COMMERCIAL DISTRIBUTION IN
SERVICE THAT CHARGES FOR DOWNLOAD TIME OR FOR MEMBER
```

Obviously that is going to taint the analysis. But it also serves as a convenient marker to divide that long, long, long string into separate documents.

```
> shakespeare = strsplit(shakespeare, "<<[^>]*>>")[[
> length(shakespeare)
[1] 218
```

This left me with a list of 218 documents. On further inspection, some of them appeared to be a little on the short side (in my limited experience, the bard is not known for brevity). As it turns out, the short documents were the [dramatis personae](#) for his plays. I removed them as well.

```
> (dramatis.personae <- grep("Dramatis Personae", sh
[1] 2 8 11 17 23 28 33 43 49 55 62 68
[28] 164 170 176 182 188 194 200 206 212
> length(shakespeare)
[1] 218
> shakespeare = shakespeare[-dramatis.personae]
> length(shakespeare)
[1] 182
```

Down to 182 documents, each of which is a complete work.

The next task was to convert these documents into a [corpus](#).

```
> library(tm)
>
> doc.vec <- VectorSource(shakespeare)
> doc.corpus <- Corpus(doc.vec)
> summary(doc.corpus)
A corpus with 182 text documents
```

The metadata consists of 2 tag-value pairs and a data frame. Available tags are:  
create\_date creator  
Available variables in the data frame are:  
MetaID

There is a lot of information in those documents which is not particularly useful for text mining. So before proceeding any further, we will clean things up a bit. First we convert all of the text to lowercase and then remove punctuation, numbers and common English stopwords. Possibly the list of English stop words is not entirely appropriate for Shakespearean English, but it is a reasonable starting point.

```
> doc.corpus <- tm_map(doc.corpus, tolower)
> doc.corpus <- tm_map(doc.corpus, removePunctuation)
> doc.corpus <- tm_map(doc.corpus, removeNumbers)
> doc.corpus <- tm_map(doc.corpus, removeWords, stop
```

Next we perform [stemming](#), which removes affixes from words (so, for example, “run”, “runs” and “running” all become “run”).

```
> library(SnowballC)
>
> doc.corpus <- tm_map(doc.corpus, stemDocument)
```

All of these transformations have resulted in a lot of whitespace, which is then removed.

```
> doc.corpus <- tm_map(doc.corpus, stripWhitespace)
```

If we have a look at what’s left, we find that it’s just the lowercase, stripped down version of the text (which I have truncated here).

```
> inspect(doc.corpus[8])
A corpus with 1 text document
```

The metadata consists of 2 tag-value pairs and a data frame. Available tags are:  
create\_date creator  
Available variables in the data frame are:  
MetaID

```
[[1]]
act ii scene messina pompey hous enter pompey menec
assist deed justest men menecr know worthi pompey d
sue menecr ignor beg often harm wise powr deni us g
peopl love sea mine power crescent augur hope say w
```

```
will make war without door caesar get money lose he
```

This is where things start to get interesting. Next we create a [Term Document Matrix](#) (TDM) which reflects the number of times each word in the corpus is found in each of the documents.

```
> TDM <- TermDocumentMatrix(doc.corpus)
> TDM
A term-document matrix (18651 terms, 182 documents)
```

```
Non-/sparse entries: 182898/3211584
Sparsity           : 95%
Maximal term length: 31
Weighting          : term frequency (tf)
> inspect(TDM[1:10,1:10])
A term-document matrix (10 terms, 10 documents)
```

```
Non-/sparse entries: 1/99
Sparsity           : 99%
Maximal term length: 9
Weighting          : term frequency (tf)
```

	Docs									
Terms	1	2	3	4	5	6	7	8	9	10
aaron	0	0	0	0	0	0	0	0	0	0
abaissiez	0	0	0	0	0	0	0	0	0	0
abandon	0	0	0	0	0	0	0	0	0	0
abandond	0	1	0	0	0	0	0	0	0	0
abas	0	0	0	0	0	0	0	0	0	0
abashd	0	0	0	0	0	0	0	0	0	0
abat	0	0	0	0	0	0	0	0	0	0
abatfowl	0	0	0	0	0	0	0	0	0	0
abbess	0	0	0	0	0	0	0	0	0	0
abbey	0	0	0	0	0	0	0	0	0	0

The extract from the TDM shows, for example, that the word “abandond” occurred once in document number 2 but was not present in any of the other first ten documents. We could have generated the transpose of the DTM as well.

```
> DTM <- DocumentTermMatrix(doc.corpus)
> inspect(DTM[1:10,1:10])
A document-term matrix (10 documents, 10 terms)
```

```
Non-/sparse entries: 1/99
Sparsity           : 99%
Maximal term length: 9
Weighting          : term frequency (tf)
```

	Terms						
Docs	aaron	abaissiez	abandon	abandond	abas	abashd	ab
1	0	0	0	0	0	0	0
2	0	0	0	1	0	0	0
3	0	0	0	0	0	0	0
4	0	0	0	0	0	0	0
5	0	0	0	0	0	0	0
6	0	0	0	0	0	0	0
7	0	0	0	0	0	0	0
8	0	0	0	0	0	0	0
9	0	0	0	0	0	0	0
10	0	0	0	0	0	0	0

Which of these proves to be most convenient will depend on the relative number of documents and terms in your data.

Now we can start asking questions like: what are the most frequently occurring terms?

```
> findFreqTerms(TDM, 2000)
[1] "come" "enter" "good" "king" "let" "lord"
[14] "thi" "thou" "well" "will"
```

Each of these words occurred more than 2000 times.

What about associations between words? Let's have a look at what other words had a high association with “love”.

```
> findAssocs(TDM, "love", 0.8)
beauti eye
0.83 0.80
```

Well that's not too surprising!

From our first look at the TDM we know that there are many terms

which do not occur very often. It might make sense to simply remove these sparse terms from the analysis.

```
> TDM.common = removeSparseTerms(TDM, 0.1)
> dim(TDM)
[1] 18651 182
> dim(TDM.common)
[1] 71 182
```

From the 18651 terms that we started with, we are now left with a TDM which considers on 71 commonly occurring terms.

```
> inspect(TDM.common[1:10,1:10])
A term-document matrix (10 terms, 10 documents)
```

```
Non-/sparse entries: 94/6
Sparsity           : 6%
Maximal term length: 6
Weighting          : term frequency (tf)
```

Terms	1	2	3	4	5	6	7	8	9	10
act	1	4	7	9	6	3	2	14	1	0
art	53	0	9	3	5	3	2	17	0	6
away	18	5	8	4	2	10	5	13	1	7
call	17	1	4	2	2	1	6	17	3	7
can	44	8	12	5	10	6	10	24	1	5
come	19	9	16	17	12	15	14	89	9	15
day	43	2	2	4	1	5	3	17	2	3
enter	0	7	12	11	10	10	14	87	4	6
exeunt	0	3	8	8	5	4	7	49	1	4
exit	0	6	8	5	6	5	3	31	3	2

Finally we are going to put together a visualisation. The TDM is stored as a [sparse matrix](#). This was an apt representation for the initial TDM, but the reduced TDM containing only frequently occurring words is probably better stored as a normal matrix. We'll make the conversion and see.


```
> library(slam)
>
> TDM.dense <- as.matrix(TDM.common)
>
> TDM.dense
> object.size(TDM.common)
207872 bytes
> object.size(TDM.dense)
112888 bytes
```

So, as it turns out the sparse representation was actually wasting space! (This will generally not be true though: it will only apply for a matrix consisting of just the common terms). Anyway, we need the data as a normal matrix in order to produce the visualisation. The next step is to convert it into a tidy format.

```
> library(reshape2)
>
> TDM.dense = melt(TDM.dense, value.name = "count")
> head(TDM.dense)
  Terms Docs count
1   act    1     1
2   art    1    53
3  away    1    18
4  call    1    17
5   can    1    44
6  come    1    19
```

And finally generate the visualisation.

```
> library(ggplot2)
>
> ggplot(TDM.dense, aes(x = Docs, y = Terms, fill =
+   geom_tile(colour = "white") +
+   scale_fill_gradient(high="#FF0000" , low="#FFF
+   ylab("")) +
+   theme(panel.background = element_blank()) +
+   theme(axis.text.x = element_blank()), axis.tick
```



The colour scale indicates the number of times that each of the terms cropped up in each of the documents. I applied a logarithmic transform to the counts since there was a very large disparity in the numbers across terms and documents. The grey tiles correspond to terms which are not found in the corresponding document.

One can see that some terms, like “will” turn up frequently in most documents, while “love” is common in some and rare or absent in others.

That was interesting. Not sure that I would like to make any conclusions on the basis of the results above (Shakespeare is well outside my field of expertise!), but I now have a pretty good handle on how the tm package works. As always, feedback will be appreciated!

## References


- [Build a search engine in 20 minutes or less](#)
- Feinerer, I. (2013). Introduction to the tm Package: Text Mining in R.
- Feinerer, I., Hornik, K., & Meyer, D. (2008). Text Mining Infrastructure in R. Journal of Statistical Software, 25(5).

### Related




The mnemoneitoR

In "R bloggers"



Clustering the Words of William Shakespeare

In "R bloggers"



My Intro to Multiple Classification with Random Forests, Conditional Inference Trees, and Linear Discriminant Analysis

In "R bloggers"

To leave a comment for the author, please follow the link and comment on his blog: [Exegetic Analytics » R](#).

R-bloggers.com offers [daily e-mail updates](#) about R news and [tutorials](#) on topics such as: visualization ([ggplot2](#), [Boxplots](#), [maps](#), [animation](#)), programming ([RStudio](#), [Sweave](#), [LaTeX](#), [SQL](#), [Eclipse](#), [git](#), [hadoop](#), [Web Scraping](#)) statistics ([regression](#), [PCA](#), [time series](#), [trading](#)) and more...

If you got this far, why not **subscribe for updates** from the site?  
Choose your flavor: [e-mail](#), [twitter](#), [RSS](#), or [facebook](#)...

Comments are closed.

## Top 3 Posts from the past 2 days

- [R at Microsoft](#)
- [In-depth introduction to machine learning in 15 hours of expert videos](#)
- [The Single Most Important Skill for a Data Scientist](#)

Search & Hit Enter

## Top 9 articles of the week

1. [Scatterplots](#)
2. [In-depth introduction to machine learning in 15 hours of expert videos](#)
3. [Installing R packages](#)
4. [The Single Most Important Skill for a Data Scientist](#)
5. [Illustrated Guide to ROC and AUC](#)
6. [Using apply, sapply, lapply in R](#)
7. [Network analysis with igraph](#)
8. [R vs Python: Survival Analysis with Plotly](#)
9. [KDD Cup 2015: The story of how I built hundreds of predictive models....And got so close, yet so far away from 1st place!](#)

## Sponsors



### Highland Statistics Ltd

Zero Inflated Models & GLMM

Beginner's Guide to GAM

Beginner's Guide to GLM & GLMM

Beginner's Guide to GAMM



 Quantide: statistical consulting and training





**REVOLUTION**  
ANALYTICS



Werden Sie zum Expe[R]ten mit der  
R-Akademie von



Beratung | Software  
Training | Lösungen



[Plotly: collaborative, publication-quality graphing](#)



*use***R!**

June 30 - July 3, 2015 Aalborg, Denmark



**Data Analysis  
Modeling**  
Cloud Solution

**Free Sign Up!**



**STATWORX**


Consulting  
Schulung  
Data Mining

**Mehr erfahren**



**Try the FASTEST ML  
for R**

Click for a Free Trial



YOTTAMINE ANALYTICS

**CLIMB THE BIG DATA BANDWAGON!**

ALL-IN-ONE BIG DATA & ANALYTICS SUITE

3-DAYS MONEY BACK GUARANTEE

**\$399**

**simplilearn**

**ENROLL NOW**

**OPEN DATA SCIENCE CONFERENCE**

Meet Data Science Thought Leaders in Boston May 30th & 31st

Tickets Start at \$100 (early reg)

21 Workshops 76 Presentations #ODSC



**SIX**

**SIX**

**EARN YOUR MASTER'S IN  
DATA SCIENCE online.**

**Berkeley**  
UNIVERSITY OF CALIFORNIA

**LEARN MORE**

Accelerated R, Scientific Computing,  
Cloud, & AWS. **Easy to get started.**

**Start Your Free Trial**

**bitfusion.io**

Search & Hit Enter

[Full list of contributing R-bloggers](#)

[R-bloggers](#) was founded by [Tal Galili](#), with gratitude to the [R](#) community.

Is powered by [WordPress](#) using a [bavotasan.com](#) design.

Copyright © 2015 [R-bloggers](#). All Rights Reserved. [Terms and Conditions](#) for this website