

master

...

[tutorials](#) / [beginner\\_source](#) / [blitz](#) / [cifar10\\_tutorial.py](#) / <> Jump to ▾

malfet Modernize cifar10\_tutorial (#2086) ... ✓



24 contributors



+12

367 lines (293 sloc) | 12.5 KB

...

```
1  # -*- coding: utf-8 -*-
2  """
3  Training a Classifier
4  =====
5
6  This is it. You have seen how to define neural networks, compute loss and make
7  updates to the weights of the network.
8
9  Now you might be thinking,
10
11  What about data?
12  -----
13
14  Generally, when you have to deal with image, text, audio or video data,
15  you can use standard python packages that load data into a numpy array.
16  Then you can convert this array into a ``torch.*Tensor``.
17
18  - For images, packages such as Pillow, OpenCV are useful
19  - For audio, packages such as scipy and librosa
20  - For text, either raw Python or Cython based loading, or NLTK and
21    SpaCy are useful
22
23  Specifically for vision, we have created a package called
24  ``torchvision``, that has data loaders for common datasets such as
25  ImageNet, CIFAR10, MNIST, etc. and data transformers for images, viz.,
26  ``torchvision.datasets`` and ``torch.utils.data.DataLoader``.
27
28  This provides a huge convenience and avoids writing boilerplate code.
29
30  For this tutorial, we will use the CIFAR10 dataset.
31  It has the classes: 'airplane', 'automobile', 'bird', 'cat', 'deer',
32  'dog', 'frog', 'horse', 'ship', 'truck'. The images in CIFAR-10 are of
```

```

33 size 3x32x32, i.e. 3-channel color images of 32x32 pixels in size.
34
35 .. figure:: /_static/img/cifar10.png
36     :alt: cifar10
37
38     cifar10
39
40
41 Training an image classifier
42 -----
43
44 We will do the following steps in order:
45
46 1. Load and normalize the CIFAR10 training and test datasets using
47     ``torchvision``
48 2. Define a Convolutional Neural Network
49 3. Define a loss function
50 4. Train the network on the training data
51 5. Test the network on the test data
52
53 1. Load and normalize CIFAR10
54 ^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^
55
56 Using ``torchvision``, it's extremely easy to load CIFAR10.
57 """"
58 import torch
59 import torchvision
60 import torchvision.transforms as transforms
61
62 #####
63 # The output of torchvision datasets are PILImage images of range [0, 1].
64 # We transform them to Tensors of normalized range [-1, 1].
65
66 #####
67 # .. note::
68 #     If running on Windows and you get a BrokenPipeError, try setting
69 #     the num_worker of torch.utils.data.DataLoader() to 0.
70
71 transform = transforms.Compose(
72     [transforms.ToTensor(),
73      transforms.Normalize((0.5, 0.5, 0.5), (0.5, 0.5, 0.5))])
74
75 batch_size = 4
76
77 trainset = torchvision.datasets.CIFAR10(root='./data', train=True,
78                                         download=True, transform=transform)
79 trainloader = torch.utils.data.DataLoader(trainset, batch_size=batch_size,
80                                           shuffle=True, num_workers=2)
81
82 testset = torchvision.datasets.CIFAR10(root='./data', train=False,
83                                         download=True, transform=transform)
84 testloader = torch.utils.data.DataLoader(testset, batch size=batch size,
```

```

85         shuffle=False, num_workers=2)
86
87 classes = ('plane', 'car', 'bird', 'cat',
88            'deer', 'dog', 'frog', 'horse', 'ship', 'truck')
89
90 #####
91 # Let us show some of the training images, for fun.
92
93 import matplotlib.pyplot as plt
94 import numpy as np
95
96 # functions to show an image
97
98
99 def imshow(img):
100     img = img / 2 + 0.5     # unnormalize
101     npimg = img.numpy()
102     plt.imshow(np.transpose(npimg, (1, 2, 0)))
103     plt.show()
104
105
106 # get some random training images
107 dataiter = iter(trainloader)
108 images, labels = next(dataiter)
109
110 # show images
111 imshow(torchvision.utils.make_grid(images))
112 # print labels
113 print(' '.join(f'{classes[labels[j]]:5s}' for j in range(batch_size)))
114
115
116 #####
117 # 2. Define a Convolutional Neural Network
118 # ^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^
119 # Copy the neural network from the Neural Networks section before and modify it to
120 # take 3-channel images (instead of 1-channel images as it was defined).
121
122 import torch.nn as nn
123 import torch.nn.functional as F
124
125
126 class Net(nn.Module):
127     def __init__(self):
128         super().__init__()
129         self.conv1 = nn.Conv2d(3, 6, 5)
130         self.pool = nn.MaxPool2d(2, 2)
131         self.conv2 = nn.Conv2d(6, 16, 5)
132         self.fc1 = nn.Linear(16 * 5 * 5, 120)
133         self.fc2 = nn.Linear(120, 84)
134         self.fc3 = nn.Linear(84, 10)
135
136     def forward(self, x):

```

```

137         x = self.pool(F.relu(self.conv1(x)))
138         x = self.pool(F.relu(self.conv2(x)))
139         x = torch.flatten(x, 1) # flatten all dimensions except batch
140         x = F.relu(self.fc1(x))
141         x = F.relu(self.fc2(x))
142         x = self.fc3(x)
143         return x
144
145
146 net = Net()
147
148 #####
149 # 3. Define a Loss function and optimizer
150 # ^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^
151 # Let's use a Classification Cross-Entropy loss and SGD with momentum.
152
153 import torch.optim as optim
154
155 criterion = nn.CrossEntropyLoss()
156 optimizer = optim.SGD(net.parameters(), lr=0.001, momentum=0.9)
157
158 #####
159 # 4. Train the network
160 # ^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^^
161 #
162 # This is when things start to get interesting.
163 # We simply have to loop over our data iterator, and feed the inputs to the
164 # network and optimize.
165
166 for epoch in range(2): # loop over the dataset multiple times
167
168     running_loss = 0.0
169     for i, data in enumerate(trainloader, 0):
170         # get the inputs; data is a list of [inputs, labels]
171         inputs, labels = data
172
173         # zero the parameter gradients
174         optimizer.zero_grad()
175
176         # forward + backward + optimize
177         outputs = net(inputs)
178         loss = criterion(outputs, labels)
179         loss.backward()
180         optimizer.step()
181
182         # print statistics
183         running_loss += loss.item()
184         if i % 2000 == 1999:    # print every 2000 mini-batches
185             print(f'[epoch + 1], {i + 1:5d}] loss: {running_loss / 2000:.3f}')
186             running_loss = 0.0
187
188 print('Finished Training')

```



```

241 #####
242 # The results seem pretty good.
243 #
244 # Let us look at how the network performs on the whole dataset.
245
246 correct = 0
247 total = 0
248 # since we're not training, we don't need to calculate the gradients for our outputs
249 with torch.no_grad():
250     for data in testloader:
251         images, labels = data
252         # calculate outputs by running images through the network
253         outputs = net(images)
254         # the class with the highest energy is what we choose as prediction
255         _, predicted = torch.max(outputs.data, 1)
256         total += labels.size(0)
257         correct += (predicted == labels).sum().item()
258
259 print(f'Accuracy of the network on the 10000 test images: {100 * correct // total} %')
260
261 #####
262 # That looks way better than chance, which is 10% accuracy (randomly picking
263 # a class out of 10 classes).
264 # Seems like the network learnt something.
265 #
266 # Hmmm, what are the classes that performed well, and the classes that did
267 # not perform well:
268
269 # prepare to count predictions for each class
270 correct_pred = {classname: 0 for classname in classes}
271 total_pred = {classname: 0 for classname in classes}
272
273 # again no gradients needed
274 with torch.no_grad():
275     for data in testloader:
276         images, labels = data
277         outputs = net(images)
278         _, predictions = torch.max(outputs, 1)
279         # collect the correct predictions for each class
280         for label, prediction in zip(labels, predictions):
281             if label == prediction:
282                 correct_pred[classes[label]] += 1
283                 total_pred[classes[label]] += 1
284
285
286 # print accuracy for each class
287 for classname, correct_count in correct_pred.items():
288     accuracy = 100 * float(correct_count) / total_pred[classname]
289     print(f'Accuracy for class: {classname:5s} is {accuracy:.1f} %')
290
291 #####
292 # Okay, so what next?

```

```

293 #
294 # How do we run these neural networks on the GPU?
295 #
296 # Training on GPU
297 # -----
298 # Just like how you transfer a Tensor onto the GPU, you transfer the neural
299 # net onto the GPU.
300 #
301 # Let's first define our device as the first visible cuda device if we have
302 # CUDA available:
303
304 device = torch.device('cuda:0' if torch.cuda.is_available() else 'cpu')
305
306 # Assuming that we are on a CUDA machine, this should print a CUDA device:
307
308 print(device)
309
310 #####
311 # The rest of this section assumes that ``device`` is a CUDA device.
312 #
313 # Then these methods will recursively go over all modules and convert their
314 # parameters and buffers to CUDA tensors:
315 #
316 # .. code:: python
317 #
318 #     net.to(device)
319 #
320 #
321 # Remember that you will have to send the inputs and targets at every step
322 # to the GPU too:
323 #
324 # .. code:: python
325 #
326 #     inputs, labels = data[0].to(device), data[1].to(device)
327 #
328 # Why don't I notice MASSIVE speedup compared to CPU? Because your network
329 # is really small.
330 #
331 # **Exercise:** Try increasing the width of your network (argument 2 of
332 # the first ``nn.Conv2d``, and argument 1 of the second ``nn.Conv2d`` -
333 # they need to be the same number), see what kind of speedup you get.
334 #
335 # **Goals achieved**:

```

```
345 # Where do I go next?
346 # -----
347 #
348 # - :doc:`Train neural nets to play video games </intermediate/reinforcement_q_learning>`
349 # - `Train a state-of-the-art ResNet network on imagenet`_
350 # - `Train a face generator using Generative Adversarial Networks`_
351 # - `Train a word-level language model using Recurrent LSTM networks`_
352 # - `More examples`_
353 # - `More tutorials`_
354 # - `Discuss PyTorch on the Forums`_
355 # - `Chat with other users on Slack`_
356 #
357 # .. _Train a state-of-the-art ResNet network on imagenet: https://github.com/pytorch/examples,
358 # .. _Train a face generator using Generative Adversarial Networks: https://github.com/pytorch,
359 # .. _Train a word-level language model using Recurrent LSTM networks: https://github.com/pytor
360 # .. _More examples: https://github.com/pytorch/examples
361 # .. _More tutorials: https://github.com/pytorch/tutorials
362 # .. _Discuss PyTorch on the Forums: https://discuss.pytorch.org/
363 # .. _Chat with other users on Slack: https://pytorch.slack.com/messages/beginner/
364
365 # %%%INVISIBLE_CODE_BLOCK%%%
366 del dataiter
367 # %%%INVISIBLE_CODE_BLOCK%%%
```