


Programming Review



PyTorch | Automatic Differentiation



Table of Contents:

- [What is PyTorch AD ?](#)
 - [What PyTorch AD is not ?](#)
 - [How reverse mode AD works?](#)
- [Forward computational graph](#)
- [viz](#)
- [n AD](#)
- [ie deep learning](#)



PyTorch AD ?

Automatic Differentiation (AD) is a technique to compute the derivative of a function $f(x_1, \dots, x_n)$ at some point.



calculate the derivate. Symbolic math approach would be to use derivation rules. For example, if you have $f(x) = \frac{1}{x}$ then $f'(x) = -\frac{1}{x^2}$.

AD is also not numeric procedure to calculate the derivate. The numerical procedure to calculate the derivative of **tanh** function at point $x = 1$ would be:

```
def tanh(x):  
    y=np.exp(-x)  
    return (1.0-y)/(1.0+y)  
  
s=0.00001 # some small number  
x=1.0  
d=(tanh(x+s)-tanh(x))/s  
print(d)
```

Output:

```
0.39322295790622513
```

How reverse mode AD works?

PyTorch uses **reverse mode** AD. AD *forward mode* exists, but it is computationally more expensive.

Reverse mode AD works the following way.

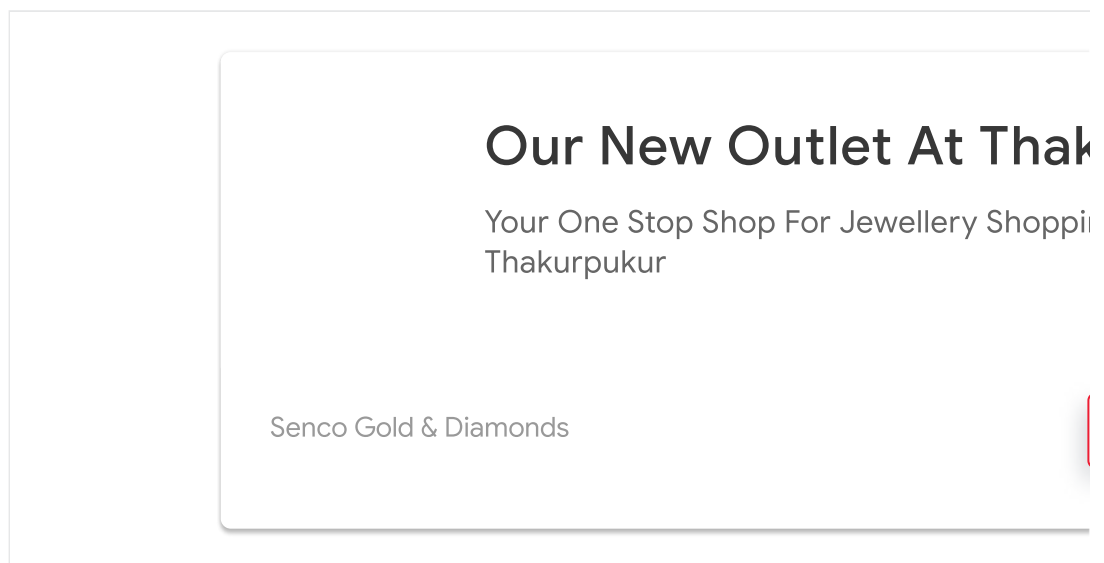


calculates the intermediate variables based on inputs.


The computational graph being calculated is like a tree. Inputs are *tree leaves* and each node in the graph corresponds to some operation (such as $+$), or to some function (such as \sin).

The output function is *the root* of the tree. Once we create the computational tree in the forward pass, together with the intermediate gradients, we then may get the gradients from the root to any of the leaves.

We say: we compute the gradients of a function with respect to the input variable x .



This pass when we compute the gradients is known as the backward pass and corresponds to PyTorch [grad](#) function.



PyTorch **grad** function is very cheap. It just traverses the computational graph and creates the sum of the intermediate gradient products to calculate the final gradient. The math behind calculating gradient is called the **chain rule**.

NOTE:

Note: There is one tool similar to Pytorch called [Chainer](#) just because of this chain rule principle.

Example:

To get a clue how PyTorch AD works the next code example we will create the computational graph for the function:

$$f(x_1, x_2) = \frac{1 + \sin(x_2)}{x_2 + e^{x_1}} + x_1 x_2$$

We will calculate the gradient of a function $f(x_1, x_2)$ with respect to x_2 .

```
import math  
class ADNumber:
```



```
def __truediv__(self,other):
    new = ADNumber(self._val / other._val)
    self._children.append((1.0/other._val,new))
    other._children.append((-self._val,new))
    return new
```

```
def __mul__(self,other):
    new = ADNumber(self._val*other._val)
    self._children.append((other._val,new))
    other._children.append((self._val,new))
    return new
```

```
def __add__(self,other):
    if isinstance(other, (int, float)):
        other = ADNumber(other, str(other))
    new = ADNumber(self._val+other._val)
    self._children.append((1.0,new))
    other._children.append((1.0,new))
    return new
```

```
def __sub__(self,other):
    new = ADNumber(self._val-other._val)
    self._children.append((1.0,new))
    other._children.append((-1.0,new))
    return new
```

```
@staticmethod
def exp(self):
    new = ADNumber(math.exp(self._val))
    self._children.append((self._val,new))
    return new
```

```
@staticmethod
def sin(self):
    new = ADNumber(math.sin(self._val))
    self._children.append((math.cos(self._val),new))
    return new
```



SENCO
GOLD & DIAMONDS

We are now at
Thakurpukur

310 BEHALA THAKURPUKUR, GROUND FL
NEAR THAKURPUKUR 3A BUS STAND
310 DIAMOND HARBOUR ROAD

```

        result=0.0
        for child in other._children:
            result+=child[0]*self.grad
        return result

A = ADNumber # shortcuts
sin = A.sin
exp = A.exp

def print_child(f, wrt): # with respect to
    for e in f._children:
        print("child:", wrt, "->" , e[1].name)
        print_child(e[1], e[1].name)

x1 = A(1.5, name="x1")
x2 = A(0.5, name="x2")
f=(sin(x2)+1)/(x2+exp(x1))+x1*x2

print_childs(x2,"x2")
print("\ncalculated gradient for the function f with wrt x2")

```

Out:

```

child: x2 -> sin(x2) grad: 0.8775825618903727
child: sin(x2) -> sin(x2)+1 grad: 1.0
child: sin(x2)+1 -> sin(x2)+1/x2+exp(x1) grad: 0.8775825618903727
child: sin(x2)+1/x2+exp(x1) -> sin(x2)+1/x2+exp(x1) grad: 0.8775825618903727
child: x2 -> x2+exp(x1) grad: 1.0
child: x2+exp(x1) -> sin(x2)+1/x2+exp(x1) grad: 0.8775825618903727
child: sin(x2)+1/x2+exp(x1) -> sin(x2)+1/x2+exp(x1) grad: 0.8775825618903727
child: x2 -> x1*x2 grad: 1.5
child: x1*x2 -> sin(x2)+1/x2+exp(x1)+x1*x2 grad: 0.8775825618903727

calculated gradient for the function f with wrt x2

```





Out:

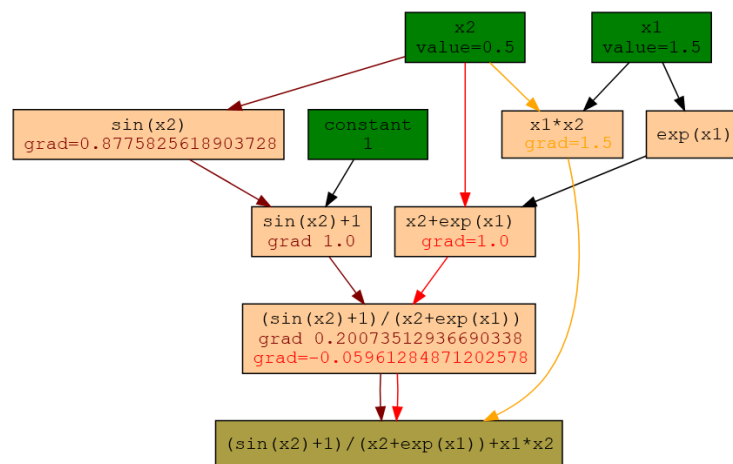
1.6165488003791768

The next image shows the computational graph for the example function:

$$f(x_1, x_2) = \frac{1 + \sin(x_2)}{x_2 + e^{x_1}} + x_1 x_2$$

where

$$x_1 = 1.5, x_2 = 0.5$$



Each node in the tree graph is either a leaf node (green) or the root node (brown) or something in between.

From the input x_2 in forward pass we identify three paths leading to the root. The arrows in dark red, red and orange denote these paths. We can ignore black





We are now at
Thakurpukur



310 BEHALA THAKURPUKUR, GROUND FL
NEAR THAKURPUKUR 3A BUS STAND
310 DIAMOND HARBOUR ROAD

To compute the final gradient for our function `f` with respect to the `x2` we need to multiply the gradient values along the paths and finally to sum them up.

The calculus is as follows:

```
1.5 + (0.8775825618903728 * 1.0 * 0.20073!)  
# 1.6165488003791768
```



This is exactly what our function `grad` will do if we print `f.grad(x2)` the result will be 1.6165488003791766.

Let's show the numerical procedure will provide the same result.

```
import math  
def f(x1, x2):  
    return (math.sin(x2)+1)/(x2+math.exp(x1))  
  
e=0.0001 # some small e  
x1 = 1.5  
x2 = 0.5
```

```
grad = (f(x1, x2+e)-f(x1, x2))/e
```



computational graph using torchviz

```
# !pip install torchviz
from torchviz import make_dot

# Create tensors
x1 = torch.tensor(1.5, requires_grad=True)
x2 = torch.tensor(0.5, requires_grad=True)
c = torch.tensor(1., requires_grad=True)

# Build a computational graph
y=(torch.sin(x2)+c)/(x2+torch.exp(x1))+x1
y.backward() # compute gradients

print(x1.grad)
print(x2.grad)
print(c.grad)


params = {'x1': x1, 'x2':x2, 'c': c}
param_map = {id(v): k for k, v in params.items()}
param_map

make_dot(y, {'x1': x1, 'x2':x2, 'c': c})
```

Out:

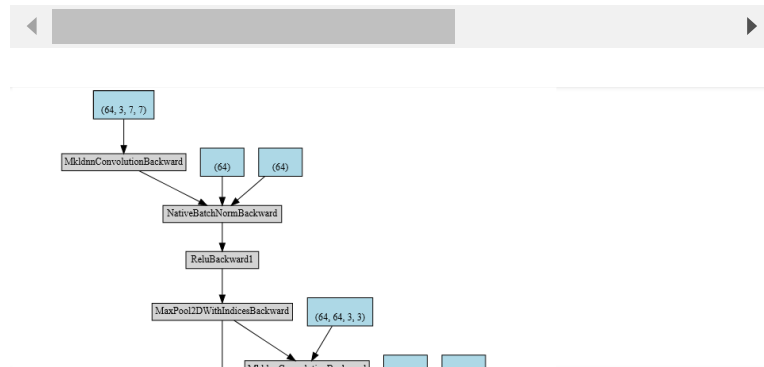
```
tensor(0.2328)
tensor(1.6165)
tensor(0.2007)
```





Example: Create resnet18 computational graph

```
import torch
import torchvision.models as models
resnet18 = models.resnet18()
x = torch.zeros(1, 3, 224, 224, dtype=torch.float)
out = resnet18(x)
make_dot(out)
```



Example: Using hiddenlayer

```
import torch
import hiddenlayer as hl
import torchvision.models as models
resnet18 = models.resnet18()
x = torch.zeros(1, 3, 224, 224, dtype=torch.float)

transforms = [ hl.transforms.Prune('Constant') ]
# resnet18 from torchvision and x is input
graph = hl.build_graph(resnet18, x, transforms)
graph.theme = hl.graph.THEMES['blue'].copy()
# graph.save('rnn_hiddenlayer' format='png')
```

SENCO
GOLD & DIAMONDS

We are now at
Thakurpukur


310 BEHALA THAKURPUKUR, GROUND FL
NEAR THAKURPUKUR 3A BUS STAND
310 DIAMOND HARBOUR ROAD

Save and zoom to check the details

Detach from AD

Here is one computational graph.

```
from torchviz import make_dot
x=torch.ones(2, requires_grad=True)
y=2*x
z=3+x
r=(y+z).sum()
make_dot(r)
```



It is possible to `detach()` the tensor from the AD computational graph.

```
from torchviz import make_dot
x=torch.ones(2, requires_grad=True)
y=2*x
z=3+x.detach()
r=(y+z).sum()
make_dot(r)
```






We are now at
Thakurpukur



310 BEHALA THAKURPUKUR, GROUND FL
NEAR THAKURPUKUR 3A BUS STAND
310 DIAMOND HARBOUR ROAD



NOTE:

`x.detach()` is the same as `x.data`.

```
from torchviz import make_dot
x=torch.ones(2, requires_grad=True)
y=2*x
z=3+x.data
r=(y+z).sum()
make_dot(r)
```

You can use the `with torch.no_grad()` class (context manager). Whatever is created inside that block, will end as `requires_grad=False`. The next example will show just that. Tensor `x` that `requires_grad=True` will create tensor `y`, but that tensor will have `requires_grad=False`.

```
x=torch.tensor(2., requires_grad=True)
print(x)
with torch.no_grad():
    y = x * 2
print(y, y.requires_grad)
```

Out:


```
tensor(2., requires_grad=True)
tensor(4.) False
```



Bonus define deep learning

In essence, for the deep learning you need to have deep models. By definition, shallow models have just one hidden layer:

```
model = torch.nn.Sequential(  
    torch.nn.Linear(D_in, H),  
    torch.nn.Linear(H, D_out),  
)
```



Input Layer $\in \mathbb{R}^2$

Hidden Layer $\in \mathbb{R}^3$

Output Layer $\in \mathbb{R}^1$

Deep models have 2 or more hidden layers.

```
model = torch.nn.Sequential(  
    torch.nn.Linear(D_in, H1),  
    torch.nn.Linear(H1, H2),  
    torch.nn.Linear(H2, D_out),  
)
```






We are now at
Thakurpukur



310 BEHALA THAKURPUKUR, GROUND FL
NEAR THAKURPUKUR 3A BUS STAND
310 DIAMOND HARBOUR ROAD



In other words, to do some deep learning you need to have at least three linear layers. The dimension H is called the hidden dimension. Instead of `nn.Linear` layers you may use convolution layers.

...

tags: *ad - pytorch automatic differentiation -
pytorch ad - automatic differentiation -
computational graph - backward computational
graph - reverse mode ad - derivation rule &*
category: *pytorch*

© [PROGRAMMING REVIEW 2010-2021](#)

We are now at
Thakurpukur

310 BEHALA THAKURPUKUR, GROUND FL
NEAR THAKURPUKUR 3A BUS STAND
310 DIAMOND HARBOUR ROAD