

# BCELOSS

CLASS `torch.nn.BCELoss(weight=None, size_average=None, reduce=None, reduction='mean')` [\[SOURCE\]](#)

Creates a criterion that measures the Binary Cross Entropy between the target and the input probabilities:

The unreduced (i.e. with `reduction` set to `'none'`) loss can be described as:

$$\ell(x, y) = L = \{l_1, \dots, l_N\}^\top, \quad l_n = -w_n [y_n \cdot \log x_n + (1 - y_n) \cdot \log(1 - x_n)],$$

where  $N$  is the batch size. If `reduction` is not `'none'` (default `'mean'`), then

$$\ell(x, y) = \begin{cases} \text{mean}(L), & \text{if reduction} = \text{'mean'}; \\ \text{sum}(L), & \text{if reduction} = \text{'sum'}. \end{cases}$$

This is used for measuring the error of a reconstruction in for example an auto-encoder. Note that the targets  $y$  should be numbers between 0 and 1.

Notice that if  $x_n$  is either 0 or 1, one of the log terms would be mathematically undefined in the above loss equation. PyTorch chooses to set  $\log(0) = -\infty$ , since  $\lim_{x \rightarrow 0} \log(x) = -\infty$ . However, an infinite term in the loss equation is not desirable for several reasons.

For one, if either  $y_n = 0$  or  $(1 - y_n) = 0$ , then we would be multiplying 0 with infinity. Secondly, if we have an infinite loss value, then we would also have an infinite term in our gradient, since  $\lim_{x \rightarrow 0} \frac{d}{dx} \log(x) = \infty$ . This would make BCELoss's backward method nonlinear with respect to  $x_n$ , and using it for things like linear regression would not be straight-forward.

Our solution is that BCELoss clamps its log function outputs to be greater than or equal to -100. This way, we can always have a finite loss value and a linear backward method.

## Parameters

- **weight** (*Tensor, optional*) – a manual rescaling weight given to the loss of each batch element. If given, has to be a Tensor of size *nbatch*.
- **size\_average** (*bool, optional*) – Deprecated (see `reduction`). By default, the losses are averaged over each loss element in the batch. Note that for some losses, there are multiple elements per sample. If the field `size_average` is set to `False`, the losses are instead summed for each minibatch. Ignored when `reduce` is `False`. Default: `True`
- **reduce** (*bool, optional*) – Deprecated (see `reduction`). By default, the losses are averaged or summed over observations for each minibatch depending on `size_average`. When `reduce` is `False`, returns a loss per batch element instead and ignores `size_average`. Default: `True`
- **reduction** (*string, optional*) – Specifies the reduction to apply to the output: `'none'` | `'mean'` | `'sum'`. `'none'`: no reduction will be applied, `'mean'`: the sum of the output will be divided by the number of elements in the output, `'sum'`: the output will be summed. Note: `size_average` and `reduce` are in the process of being deprecated, and in the meantime, specifying either of those two args will override `reduction`. Default: `'mean'`

## Shape:

- Input:  $(*)$ , where  $*$  means any number of dimensions.
- Target:  $(*)$ , same shape as the input.
- Output: scalar. If `reduction` is `'none'`, then  $(*)$ , same shape as input.

## Examples:

```
>>> m = nn.Sigmoid()
>>> loss = nn.BCELoss()
>>> input = torch.randn(3, requires_grad=True)
>>> target = torch.empty(3).random_(2)
>>> output = loss(m(input), target)
>>> output.backward()
```

### Docs

Access comprehensive developer documentation for  
PyTorch

[View Docs](#)

### Tutorials

Get in-depth tutorials for beginners and advanced  
developers

[View Tutorials](#)

### Resources

Find development resources and get your questions  
answered

[View Resources](#)

PyTorch

Get Started

Features

Ecosystem

Blog

Contributing

Resources

Tutorials

Docs

Discuss

Github Issues

Brand Guidelines

Stay Connected