

trax.models

atari_cnn

Simple net for playing Atari games using PPO.

trax.models.atari_cnn.AtariCnn (<i>n_frames=4, hidden_sizes=(32, 32), output_size=128, mode='train'</i>)
An Atari CNN.
trax.models.atari_cnn.AtariCnnBody (<i>n_frames=4, hidden_sizes=(32, 64, 64), output_size=512, mode='train', kernel_initializer=None, padding='VALID'</i>)
An Atari CNN.
trax.models.atari_cnn.FrameStackMLP (<i>n_frames=4, hidden_sizes=(64,), output_size=64, mode='train'</i>)
MLP operating on a fixed number of last frames.

mlp

mlp – functions that assemble “multilayer perceptron” networks.

trax.models.mlp.PureMLP (<i>layer_widths=(128, 64), activation_fn=<function Relu>, out_activation=False, flatten=True, mode='train'</i>)
A “multilayer perceptron” (MLP) network.
This is a classic fully connected feedforward network, with one or more layers and a (nonlinear) activation function between each layer. For historical reasons, such networks are often called multilayer perceptrons; but they are more accurately described as multilayer networks, where each layer + activation function is a perceptron-like unit (see, e.g., [https://en.wikipedia.org/wiki/Multilayer_perceptron#Terminology]).
<div>Parameters:</div> <div><ul style="list-style-type: none">layer_widths – Tuple of ints telling the number of layers and the width of each layer. For example, setting <i>layer_widths</i>=(128, 64, 32) would yield 3 layers with successive widths of 128, 64, and 32.activation_fn – Type of activation function between pairs of fully connected layers; must be an activation-type subclass of <i>Layer</i>.out_activation – If True, include a copy of the activation function as the last layer in the network.flatten – If True, insert a layer at the head of the network to flatten the input tensor into a matrix of shape (batch_size, -1).mode – Ignored.</div>
<div>Returns:</div> <div>An assembled MLP network with the specified layers. This network can either be initialized and trained as a full model, or can be used as a building block in a larger network.</div>

trax.models.mlp.MLP (<i>d_hidden=512, n_hidden_layers=2, activation_fn=<function Relu>, n_output_classes=10, flatten=True, mode='train'</i>)
An MLP network, with a final layer for n-way classification.

neural_gpu

Implementation of the improved Neural GPU (NGPU).

trax.models.neural_gpu.SaturationCost (<i>x, limit=0.9</i>)
trax.models.neural_gpu.DiagonalGate ()
Split channels in 3 parts. Shifts 1st and 3rd sections to left/right.
trax.models.neural_gpu.ConvDiagonalGRU (<i>units, kernel_size=(3, 3)</i>)
Build convolutional GRU with diagonal gating as in ImprovedNGPU.
trax.models.neural_gpu.NeuralGPU (<i>d_feature=96, steps=16, vocab_size=2, mode='train'</i>)
Implementation of Neural GPU: https://arxiv.org/abs/1702.08727 .
<div>Parameters:</div> <div><ul style="list-style-type: none">d_feature – Number of memory channels (dimensionality of feature embedding).steps – Number of times depthwise recurrence steps.vocab_size – Vocabulary size.mode – Whether we are training or evaluating or doing inference.</div>
<div>Returns:</div> <div>A NeuralGPU Stax model.</div>

resnet

ResNet.

trax.models.resnet.ConvBlock (<i>kernel_size, filters, strides, norm, non_linearity, mode='train'</i>)
ResNet convolutional striding block.
trax.models.resnet.IdentityBlock (<i>kernel_size, filters, norm, non_linearity, mode='train'</i>)
ResNet identical size block.

trax.models.resnet.Resnet50 (<i>d_hidden=64, n_output_classes=1001, mode='train', norm=<sphinx.ext.autodoc.importer._MockObject object>, non_linearity=<function Relu></i>)	
ResNet.	
Parameters:	<ul style="list-style-type: none">d_hidden – Dimensionality of the first hidden layer (multiplied later).n_output_classes – Number of distinct output classes.mode – Whether we are training or evaluating or doing inference.norm – <i>Layer</i> used for normalization, Ex: BatchNorm or FilterResponseNorm.non_linearity – <i>Layer</i> used as a non-linearity, Ex: If norm is BatchNorm then this is a Relu, otherwise for FilterResponseNorm this should be ThresholdedLinearUnit.
Returns:	The list of layers comprising a ResNet model with the given parameters.

trax.models.resnet.WideResnetBlock(*channels, strides=(1, 1), bn_momentum=0.9, mode='train'*)

WideResnet convolutional block.

trax.models.resnet.WideResnetGroup(*n, channels, strides=(1, 1), bn_momentum=0.9, mode='train'*)

trax.models.resnet.WideResnet(*n_blocks=3, widen_factor=1, n_output_classes=10, bn_momentum=0.9, mode='train'*)

WideResnet from <https://arxiv.org/pdf/1605.07146.pdf>.

Parameters:	<ul style="list-style-type: none">n_blocks – int, number of blocks in a group. total layers = 6n + 4.widen_factor – int, widening factor of each group. k=1 is vanilla resnet.n_output_classes – int, number of distinct output classes.bn_momentum – float, momentum in BatchNorm.mode – Whether we are training or evaluating or doing inference.
Returns:	The list of layers comprising a WideResnet model with the given parameters.

rl

Policy networks.

trax.models.rl.Policy(*policy_distribution, body=None, normalizer=None, head_init_range=None, batch_axes=None, mode='train'*)

Attaches a policy head to a model body.

trax.models.rl.Value(*body=None, normalizer=None, inject_actions=False, inject_actions_n_layers=1, inject_actions_dim=64, batch_axes=None, mode='train', is_discrete=False, vocab_size=2, multiplicative_action_injection=False*)

Attaches a value head to a model body.

trax.models.rl.PolicyAndValue(*policy_distribution, body=None, policy_top=<function Policy>, value_top=<function Value>, normalizer=None, head_init_range=None, mode='train'*)

Attaches policy and value heads to a model body.

trax.models.rl.Quality(*body=None, normalizer=None, batch_axes=None, mode='train', n_actions=2*)

The network takes as input an observation and outputs values of actions.

rnn

RNNs (recursive neural networks).

trax.models.rnn.RNNLM(*vocab_size, d_model=512, n_layers=2, rnn_cell=<sphinx.ext.autodoc.importer._MockObject object>, rnn_cell_d_state_multiplier=2, dropout=0.1, mode='train'*)

Returns an RNN language model.

This model performs autoregressive language modeling:

- input: rank 2 tensor representing a batch of text strings via token IDs plus padding markers; shape is (batch_size, sequence_length). The tensor elements are integers in *range(vocab_size)*, and 0 values mark padding positions.
- output: rank 3 tensor representing a batch of log-probability distributions for each sequence position over possible token IDs; shape is (batch_size, sequence_length, *vocab_size*).

Parameters:	<ul style="list-style-type: none">vocab_size – Input vocabulary size – each element of the input tensor should be an integer in <i>range(vocab_size)</i>. These integers typically represent token IDs from a vocabulary-based tokenizer.d_model – Embedding depth throughout the model.n_layers – Number of RNN layers.rnn_cell – Type of RNN cell; must be a subclass of <i>Layer</i>.rnn_cell_d_state_multiplier – Multiplier for feature depth of RNN cell state.dropout – Stochastic rate (probability) for dropping an activation value when applying dropout.mode – If <i>'predict'</i>, use fast inference; if <i>'train'</i> apply dropout.
Returns:	An RNN language model as a layer that maps from a tensor of tokens to activations over a vocab set.

trax.models.rnn.GRULM(*vocab_size=256, d_model=512, n_layers=1, mode='train'*)

Returns a GRU (gated recurrent unit) language model.

This model performs autoregressive language modeling:

- input: rank 2 tensor representing a batch of text strings via token IDs plus padding markers; shape is (batch_size, sequence_length). The tensor elements are integers in *range(vocab_size)*, and 0 values mark padding positions.
- output: rank 3 tensor representing a batch of log-probability distributions for each sequence position over possible token IDs; shape is (batch_size, sequence_length, *vocab_size*).

Parameters:	<ul style="list-style-type: none">• vocab_size – Input vocabulary size – each element of the input tensor should be an integer in <i>range(vocab_size)</i>. These integers typically represent token IDs from a vocabulary-based tokenizer.• d_model – Embedding depth throughout the model.• n_layers – Number of GRU layers.• mode – If <i>'predict'</i>, use fast inference (and omit the right shift).
Returns:	A GRU language model as a layer that maps from a tensor of tokens to activations over a vocab set.

```
trax.models.rnn.LSTMSeq2SeqAttn(input_vocab_size=256, target_vocab_size=256, d_model=512,
n_encoder_layers=2, n_decoder_layers=2, n_attention_heads=1, attention_dropout=0.0, mode='train')
```

Returns an LSTM sequence-to-sequence model with attention.

This model is an encoder-decoder that performs tokenized string-to-string (“source”-to-“target”) transduction:

- inputs (2):
 - source: rank 2 tensor representing a batch of text strings via token IDs plus padding markers; shape is (batch_size, sequence_length). The tensor elements are integers in *range(input_vocab_size)*, and 0 values mark padding positions.
 - target: rank 2 tensor representing a batch of text strings via token IDs plus padding markers; shape is (batch_size, sequence_length). The tensor elements are integers in *range(output_vocab_size)*, and 0 values mark padding positions.
- output: rank 3 tensor representing a batch of log-probability distributions for each sequence position over possible token IDs; shape is (batch_size, sequence_length, *vocab_size*).

An example use would be to translate (tokenized) sentences from English to German.

The model works as follows:

- Input encoder runs on the input tokens and creates activations that are used as both keys and values in attention.
- Pre-attention decoder runs on the targets and creates activations that are used as queries in attention.
- Attention runs on the queries, keys and values masking out input padding.
- Decoder runs on the result, followed by a cross-entropy loss.

Parameters:	<ul style="list-style-type: none">• input_vocab_size – Input vocabulary size – each element of the input tensor should be an integer in <i>range(vocab_size)</i>. These integers typically represent token IDs from a vocabulary-based tokenizer.• target_vocab_size – Target vocabulary size.• d_model – Final dimension of tensors at most points in the model, including the initial embedding output.• n_encoder_layers – Number of LSTM layers in the encoder.• n_decoder_layers – Number of LSTM layers in the decoder after attention.• n_attention_heads – Number of attention heads.• attention_dropout – Stochastic rate (probability) for dropping an activation value when applying dropout within an attention block.• mode – If <i>'predict'</i>, use fast inference. If <i>'train'</i>, each attention block will include dropout; else, it will pass all values through unaltered.
Returns:	An LSTM sequence-to-sequence model as a layer that maps from a source-target tokenized text pair to activations over a vocab set.

transformer

Transformer models: encoder, decoder, language model, and encoder-decoder.

The “Transformer” name and network architecture were introduced in the paper [Attention Is All You Need](<https://arxiv.org/abs/1706.03762>).

```
trax.models.transformer.TransformerEncoder(vocab_size, n_classes=10, d_model=512, d_ff=2048,
n_layers=6, n_heads=8, max_len=2048, dropout=0.1, dropout_shared_axes=None, mode='train', ff_activation=
<function Relu>)
```

Returns a Transformer encoder merged with an N-way categorization head.

This model performs text categorization:

- input: rank 2 tensor representing a batch of text strings via token IDs plus padding markers; shape is (batch_size, sequence_length). The tensor elements are integers in *range(vocab_size)*, and 0 values mark padding positions.
- output: rank 2 tensor representing a batch of log-probability distributions over N categories; shape is (batch_size, *n_classes*).

Parameters:	<ul style="list-style-type: none">• vocab_size – Input vocabulary size – each element of the input tensor should be an integer in <i>range(vocab_size)</i>. These integers typically represent token IDs from a vocabulary-based tokenizer.• n_classes – Final dimension of the output tensors, representing N-way classification.• d_model – Final dimension of tensors at most points in the model, including the initial embedding output.• d_ff – Size of special dense layer in the feed-forward part of each encoder block.• n_layers – Number of encoder blocks. Each block includes attention, dropout, residual, feed-forward (<i>Dense</i>), and activation layers.• n_heads – Number of attention heads.• max_len – Maximum symbol length for positional encoding.• dropout – Stochastic rate (probability) for dropping an activation value when applying dropout within an encoder block.• dropout_shared_axes – Tensor axes on which to share a dropout mask. Sharing along batch and sequence axes (<i>dropout_shared_axes=(0,1)</i>) is a useful way to save memory and apply consistent masks to activation vectors at different sequence positions.• mode – If <i>'train'</i>, each encoder block will include dropout; else, it will pass all values through unaltered.• ff_activation – Type of activation function at the end of each encoder block; must be an activation-type subclass of <i>Layer</i>.
Returns:	A Transformer model that maps strings (conveyed via token IDs) to probability-like activations over a range of output classes.

```
trax.models.transformer.TransformerDecoder(vocab_size=None, d_model=512, d_ff=2048, n_layers=6,
n_heads=8, max_len=2048, dropout=0.1, dropout_shared_axes=None, mode='train', ff_activation=<function Relu>)
```

Returns a Transformer decoder.

This model maps sequential inputs to sequential outputs:

- input if *vocab_size* is specified: rank 2 tensor representing a batch of text strings via token IDs plus padding markers; shape is (batch_size, sequence_length). The tensor elements are integers in *range(vocab_size)*, and 0 values mark padding positions.
- input if *vocab_size* is None: rank 2 tensor representing a batch of activation vectors; shape is (batch_size, sequence_length, *d_model*).
- output: rank 3 tensor with shape (batch_size, sequence_length, *d_model*).

The model uses causal attention and does *not* shift the input to the right. Thus, the output for position *t* is based on inputs up to and including position *t*.

Parameters:	<ul style="list-style-type: none">• vocab_size – If specified, gives the input vocabulary size – each element of the input tensor should be an integer in <i>range(vocab_size)</i>. If None, indicates that the model expects as input floating point vectors, each with <i>d_model</i> components.• d_model – Final dimension of tensors at most points in the model, including the initial embedding output.• d_ff – Size of special dense layer in the feed-forward part of each decoder block.• n_layers – Number of decoder blocks. Each block includes attention, dropout, residual, feed-forward (<i>Dense</i>), and activation layers.• n_heads – Number of attention heads.• max_len – Maximum symbol length for positional encoding.• dropout – Stochastic rate (probability) for dropping an activation value when applying dropout within a decoder block.• dropout_shared_axes – Tensor axes on which to share a dropout mask. Sharing along batch and sequence axes (<i>dropout_shared_axes=(0,1)</i>) is a useful way to save memory and apply consistent masks to activation vectors at different sequence positions.• mode – If <i>'train'</i>, each decoder block will include dropout; else, it will pass all values through unaltered.• ff_activation – Type of activation function at the end of each decoder block; must be an activation-type subclass of <i>Layer</i>.
Returns:	<p>a Transformer model that maps strings (conveyed via token IDs) to sequences of activation vectors.</p> <p>If <i>vocab_size</i> is None: a Transformer model that maps sequences of activation vectors to sequences of activation vectors.</p>
Return type:	If <i>vocab_size</i> is defined

```
trax.models.transformer.TransformerLM(vocab_size, d_model=512, d_ff=2048, n_layers=6, n_heads=8,
max_len=2048, dropout=0.1, dropout_shared_axes=None, mode='train', ff_activation=<function Relu>)
```

Returns a Transformer language model.

This model performs autoregressive language modeling:

- input: rank 2 tensor representing a batch of text strings via token IDs plus padding markers; shape is (batch_size, sequence_length). The tensor elements are integers in *range(vocab_size)*, and 0 values mark padding positions.
- output: rank 3 tensor representing a batch of log-probability distributions for each sequence position over possible token IDs; shape is (batch_size, sequence_length, *vocab_size*).

This model uses only the decoder part of the overall Transformer.

Parameters:	<ul style="list-style-type: none">• vocab_size – Input vocabulary size – each element of the input tensor should be an integer in <i>range(vocab_size)</i>. These integers typically represent token IDs from a vocabulary-based tokenizer.• d_model – Final dimension of tensors at most points in the model, including the initial embedding output.• d_ff – Size of special dense layer in the feed-forward part of each encoder block.• n_layers – Number of encoder blocks. Each block includes attention, dropout, residual, feed-forward (<i>Dense</i>), and activation layers.• n_heads – Number of attention heads.• max_len – Maximum symbol length for positional encoding.• dropout – Stochastic rate (probability) for dropping an activation value when applying dropout within an encoder block.• dropout_shared_axes – Tensor axes on which to share a dropout mask. Sharing along batch and sequence axes (<i>dropout_shared_axes=(0,1)</i>) is a useful way to save memory and apply consistent masks to activation vectors at different sequence positions.• mode – If <i>'predict'</i>, use fast inference. If <i>'train'</i>, each encoder block will include dropout; else, it will pass all values through unaltered.• ff_activation – Type of activation function at the end of each encoder block; must be an activation-type subclass of <i>Layer</i>.
Returns:	A Transformer language model as a layer that maps from a tensor of tokens to activations over a vocab set.

trax.models.transformer.Transformer(*input_vocab_size, output_vocab_size=None, d_model=512, d_ff=2048, n_encoder_layers=6, n_decoder_layers=6, n_heads=8, max_len=2048, dropout=0.1, dropout_shared_axes=None, mode='train', ff_activation=<function Relu>*)

Returns a full Transformer model.

This model is an encoder-decoder that performs tokenized string-to-string ("source"-to-"target") transduction:

- inputs (2):
 - source: rank 2 tensor representing a batch of text strings via token IDs plus padding markers; shape is (batch_size, sequence_length). The tensor elements are integers in *range(input_vocab_size)*, and 0 values mark padding positions.
 - target: rank 2 tensor representing a batch of text strings via token IDs plus padding markers; shape is (batch_size, sequence_length). The tensor elements are integers in *range(output_vocab_size)*, and 0 values mark padding positions.
- output: rank 3 tensor representing a batch of log-probability distributions for each sequence position over possible token IDs; shape is (batch_size, sequence_length, vocab_size).

An example use would be to translate (tokenized) sentences from English to German.

Parameters:	<ul style="list-style-type: none">• input_vocab_size – Input vocabulary size – each element of the input tensor should be an integer in <i>range(vocab_size)</i>. These integers typically represent token IDs from a vocabulary-based tokenizer.• output_vocab_size – If specified, gives the vocabulary size for the targets; if None, then input and target integers (token IDs) are assumed to come from the same vocabulary.• d_model – Final dimension of tensors at most points in the model, including the initial embedding output.• d_ff – Size of special dense layer in the feed-forward part of each encoder and decoder block.• n_encoder_layers – Number of encoder blocks.• n_decoder_layers – Number of decoder blocks.• n_heads – Number of attention heads.• max_len – Maximum symbol length for positional encoding.• dropout – Stochastic rate (probability) for dropping an activation value when applying dropout within an encoder/decoder block.• dropout_shared_axes – Tensor axes on which to share a dropout mask. Sharing along batch and sequence axes (<i>dropout_shared_axes=(0,1)</i>) is a useful way to save memory and apply consistent masks to activation vectors at different sequence positions.• mode – If <i>'predict'</i>, use fast inference. If <i>'train'</i>, each encoder/decoder block will include dropout; else, it will pass all values through unaltered.• ff_activation – Type of activation function at the end of each encoder/decoder block; must be an activation-type subclass of <i>Layer</i>.
Returns:	A Transformer model as a layer that maps from a source-target tokenized text pair to activations over a vocab set.

reformer.reformer

Reformer Models.

trax.models.reformer.reformer.FeedForward(*d_model, d_ff, dropout, activation, act_dropout, mode*)

Feed-forward block with layer normalization at start.

trax.models.reformer.reformer.ChunkedFeedForward(*d_model, d_ff, dropout, activation, act_dropout, chunk_size, mode*)

Chunked feed-forward block with layer normalization at start.

trax.models.reformer.reformer.FeedForwardWithOptions(*d_model, d_ff, dropout, ff_activation, ff_dropout, ff_chunk_size, ff_use_sru, ff_sparsity, mode*)

Feed-Forward block with all the options.

trax.models.reformer.reformer.DecoderBlock(*d_model, d_ff, d_attention_key, d_attention_value,*

n_heads, attention_type, dropout, ff_activation, ff_dropout, ff_use_sru, ff_chunk_size, ff_sparsity, mode)

Reversible transformer decoder layer.

- Parameters:
- **d_model** – int: depth of embedding
 - **d_ff** – int: depth of feed-forward layer
 - **d_attention_key** – int: depth of key vector for each attention head
 - **d_attention_value** – int: depth of value vector for each attention head
 - **n_heads** – int: number of attention heads
 - **attention_type** – subclass of `tl.BaseCausalAttention`: attention class to use
 - **dropout** – float: dropout rate (how much to drop out)
 - **ff_activation** – the non-linearity in feed-forward layer
 - **ff_dropout** – the dropout rate in feed-forward layer
 - **ff_use_sru** – int; if > 0, we use this many SRU layers instead of feed-forward
 - **ff_chunk_size** – int; if > 0, chunk feed-forward into this-sized chunks
 - **ff_sparsity** – int, if > 0 use sparse feed-forward block with this sparsity
 - **mode** – str: 'train' or 'eval'

Returns: the layer.

trax.models.reformer.reformer.PositionalEncoding(*mode, dropout=None, max_len=None, axial_pos_shape=None, d_axial_pos_embs=None*)

Returns the positional encoding layer depending on the arguments.

trax.models.reformer.reformer.ReformerLM(*vocab_size, d_model=512, d_ff=2048, d_attention_key=64, d_attention_value=64, n_layers=6, n_heads=8, dropout=0.1, max_len=2048, attention_type=<sphinx.ext.autodoc.importer._MockObject object>, axial_pos_shape=(), d_axial_pos_embs=None, ff_activation=<function FastGelu>, ff_use_sru=0, ff_chunk_size=0, ff_sparsity=0, mode='train')*)

Reversible transformer language model (only uses a decoder, no encoder).

- Parameters:
- **vocab_size** – int: vocab size
 - **d_model** – int: depth of *each half* of the two-part features
 - **d_ff** – int: depth of feed-forward layer
 - **d_attention_key** – int: depth of key vector for each attention head
 - **d_attention_value** – int: depth of value vector for each attention head
 - **n_layers** – int: number of decoder layers
 - **n_heads** – int: number of attention heads
 - **dropout** – float: dropout rate (how much to drop out)
 - **max_len** – int: maximum symbol length for positional encoding
 - **attention_type** – class: attention class to use, such as `SelfAttention`.
 - **axial_pos_shape** – tuple of ints: input shape to use for the axial position encoding. If unset, axial position encoding is disabled.
 - **d_axial_pos_embs** – tuple of ints: depth of position embedding for each axis. Tuple length must match `axial_pos_shape`, and values must sum to `d_model`.
 - **ff_activation** – the non-linearity in feed-forward layer
 - **ff_use_sru** – int; if > 0, we use this many SRU layers instead of feed-forward
 - **ff_chunk_size** – int; if > 0, chunk feed-forward into this-sized chunks
 - **ff_sparsity** – int, if > 0 use sparse feed-forward block with this sparsity
 - **mode** – str: 'train', 'eval', or 'predict'

Returns: the layer.

trax.models.reformer.reformer.ReformerShortenLM(*vocab_size, shorten_factor=1, d_embedding=256, d_model=512, d_ff=2048, d_attention_key=64, d_attention_value=64, n_layers=6, n_heads=8, dropout=0.1, max_len=2048, attention_type=<sphinx.ext.autodoc.importer._MockObject object>, axial_pos_shape=(), d_axial_pos_embs=None, ff_activation=<function FastGelu>, ff_use_sru=0, ff_chunk_size=0, ff_sparsity=0, mode='train')*)

Reversible transformer language model with shortening.

When `shorten_factor` is `F` and processing an input of shape `[batch, length]`, we embed the (shifted-right) input and then group each `F` elements (on `length`) into a single vector – so that in the end we process a tensor of shape

[batch, length // F, d_model]

almost until the end – at the end it's un-shortend and a SRU is applied. This reduces the length processed inside the main model body, effectively making the model faster but possibly slightly less accurate.

- Parameters:
- **vocab_size** – int: vocab size
 - **shorten_factor** – by how much to shorten, see above
 - **d_embedding** – the depth of the embedding layer and final logits
 - **d_model** – int: depth of *each half* of the two-part features
 - **d_ff** – int: depth of feed-forward layer
 - **d_attention_key** – int: depth of key vector for each attention head
 - **d_attention_value** – int: depth of value vector for each attention head
 - **n_layers** – int: number of decoder layers
 - **n_heads** – int: number of attention heads
 - **dropout** – float: dropout rate (how much to drop out)
 - **max_len** – int: maximum symbol length for positional encoding
 - **attention_type** – class: attention class to use, such as `SelfAttention`.
 - **axial_pos_shape** – tuple of ints: input shape to use for the axial position encoding. If unset, axial position encoding is disabled.
 - **d_axial_pos_embs** – tuple of ints: depth of position embedding for each axis. Tuple length must match `axial_pos_shape`, values must sum to `d_embedding`.
 - **ff_activation** – the non-linearity in feed-forward layer
 - **ff_use_sru** – int; if > 0, we use this many SRU layers instead of feed-forward
 - **ff_chunk_size** – int; if > 0, chunk feed-forward into this-sized chunks
 - **ff_sparsity** – int, if > 0 use sparse feed-forward block with this sparsity
 - **mode** – str: 'train' or 'eval'

Returns: the layer.

`trax.models.reformer.reformer.EncoderBlock(d_model, d_ff, n_heads, attention_type, dropout, ff_activation, ff_dropout, ff_use_sru=0, ff_chunk_size=0, ff_sparsity=0, mode='train')`

Returns a list of layers that implements a Reformer encoder block.

The input to the layer is a pair, (activations, mask), where the mask was created from the original source tokens to prevent attending to the padding part of the input.

- Parameters:
- `d_model` – int: depth of embedding
 - `d_ff` – int: depth of feed-forward layer
 - `n_heads` – int: number of attention heads
 - `attention_type` – subclass of `tl.BaseCausalAttention`: attention class to use
 - `dropout` – float: dropout rate (how much to drop out)
 - `ff_activation` – the non-linearity in feed-forward layer
 - `ff_dropout` – the dropout rate in feed-forward layer
 - `ff_use_sru` – int; if > 0, we use this many SRU layers instead of feed-forward
 - `ff_chunk_size` – int; if > 0, chunk feed-forward into this-sized chunks
 - `ff_sparsity` – int, if > 0 use sparse feed-forward block with this sparsity
 - `mode` – str: 'train' or 'eval'

Returns: A list of layers that maps (activations, mask) to (activations, mask).

`trax.models.reformer.reformer.EncoderDecoderBlock(d_model, d_ff, n_heads, dropout, ff_activation, ff_dropout, mode)`

Reversible transformer decoder layer.

- Parameters:
- `d_model` – int: depth of embedding
 - `d_ff` – int: depth of feed-forward layer
 - `n_heads` – int: number of attention heads
 - `dropout` – float: dropout rate (how much to drop out)
 - `ff_activation` – the non-linearity in feed-forward layer
 - `ff_dropout` – float: (optional) separate dropout rate for feed-forward layer
 - `mode` – str: 'train' or 'eval'

Returns: the layer.

`trax.models.reformer.reformer.Reformer(input_vocab_size, output_vocab_size=None, d_model=512, d_ff=2048, n_encoder_layers=6, n_decoder_layers=6, n_heads=8, dropout=0.1, max_len=2048, ff_activation=<function Relu>, ff_dropout=None, mode='train')`

Reversible transformer encoder-decoder model.

This model expects an input pair: target, source.

At the moment, this model supports dot-product attention only. For the attention types in the Reformer paper, see ReformerLM.

- Parameters:
- `input_vocab_size` – int: vocab size of the source.
 - `output_vocab_size` – int (optional): vocab size of the target. If None, the source and target are assumed to have the same vocab.
 - `d_model` – int: depth of embedding
 - `d_ff` – int: depth of feed-forward layer
 - `n_encoder_layers` – int: number of encoder layers
 - `n_decoder_layers` – int: number of decoder layers
 - `n_heads` – int: number of attention heads
 - `dropout` – float: dropout rate (how much to drop out)
 - `max_len` – int: maximum symbol length for positional encoding
 - `ff_activation` – the non-linearity in feed-forward layer
 - `ff_dropout` – float: (optional) separate dropout rate at feed-forward nonlinearity. This is called `relu_dropout` in T2T.
 - `mode` – str: 'train' or 'eval'

Returns: A Reformer model as a layer that maps from a target, source pair to activations over a vocab set.

`trax.models.reformer.reformer.Reformer2(input_vocab_size, output_vocab_size=None, d_model=512, d_ff=2048, d_attention_key=None, d_attention_value=None, n_encoder_layers=6, n_decoder_layers=6, n_heads=8, dropout=0.1, max_len=2048, encoder_attention_type=<sphinx.ext.autodoc.importer._MockObject object>, encoder_decoder_attention_type=<sphinx.ext.autodoc.importer._MockObject object>, axial_pos_shape='fixed-base', d_axial_pos_embs=None, ff_activation=<function Relu>, ff_use_sru=0, ff_chunk_size=0, ff_dropout=None, ff_sparsity=0, n_layers_forget=0, mode='train')`

Reversible transformer encoder-decoder model.

This model expects an input pair: source, target.

At the moment, this model supports dot-product attention only. For the attention types in the Reformer paper, see ReformerLM.

Parameters:	<ul style="list-style-type: none">• input_vocab_size – int: vocab size of the source.• output_vocab_size – int (optional): vocab size of the target. If None, the source and target are assumed to have the same vocab.• d_model – int: depth of embedding• d_ff – int: depth of feed-forward layer• d_attention_key – int: depth of key vector for each attention head• d_attention_value – int: depth of value vector for each attention head• n_encoder_layers – int: number of encoder layers• n_decoder_layers – int: number of decoder layers• n_heads – int: number of attention heads• dropout – float: dropout rate (how much to drop out)• max_len – int: maximum symbol length for positional encoding• encoder_attention_type – class: attention class to use, such as SelfAttention• encoder_decoder_attention_type – class: attention class to use, such as SelfAttention• axial_pos_shape – tuple of ints: input shape to use for the axial position encoding. If unset, axial position encoding is disabled.• d_axial_pos_embs – tuple of ints: depth of position embedding for each axis. Tuple length must match axial_pos_shape, and values must sum to d_model.• ff_activation – the non-linearity in feed-forward layer• ff_use_sru – int; if > 0, we use this many SRU layers instead of feed-forward• ff_chunk_size – int; if > 0, chunk feed-forward into this-sized chunks• ff_dropout – float: (optional) separate dropout rate at feed-forward nonlinearity. This is called relu_dropout in T2T.• ff_sparsity – int, if > 0 use sparse feed-forward block with this sparsity• n_layers_forget – how often to have a forgetting block between layers• mode – str: 'train' or 'eval'
Returns:	A Reformer model as a layer that maps from a target, source pair to activations over a vocab set.

research.bert

BERT.

class <code>trax.models.research.bert.AddBias</code> <i>(n_in=1, n_out=1, name=None, sublayers_to_print=None)</i>	
Bases:	<code>trax.layers.base.Layer</code>
forward <i>(inputs)</i>	Computes this layer's output as part of a forward pass through the model. <p>Authors of new layer subclasses should override this method to define the forward computation that their layer performs. Use <i>self.weights</i> to access trainable weights of this layer. If you need to use local non-trainable state or randomness, use <i>self.rng</i> for the random seed (no need to set it) and use <i>self.state</i> for non-trainable state (and set it to the new value).</p>
Parameters:	inputs – Zero or more input tensors, packaged as described in the <i>Layer</i> class docstring.
Returns:	Zero or more output tensors, packaged as described in the <i>Layer</i> class docstring.
init_weights_and_state <i>(input_signature)</i>	Initializes weights and state for inputs with the given signature. <p>Authors of new layer subclasses should override this method if their layer uses trainable weights or non-trainable state. To initialize trainable weights, set <i>self.weights</i> and to initialize non-trainable state, set <i>self.state</i> to the intended value.</p>
Parameters:	input_signature – A <i>ShapeDtype</i> instance (if this layer takes one input) or a list/tuple of <i>ShapeDtype</i> instances; signatures of inputs.
trax.models.research.bert.BERTClassifierHead <i>(n_classes)</i>	
trax.models.research.bert.BERTRegressionHead <i>()</i>	
trax.models.research.bert.BERT <i>(d_model=768, vocab_size=30522, max_len=512, type_vocab_size=2, n_heads=12, d_ff=3072, n_layers=12, head=None, init_checkpoint=None, mode='eval')</i>	
BERT (default hparams are for bert-base-uncased).	
class <code>trax.models.research.bert.PretrainedBERT</code> <i>(*sublayers, init_checkpoint=None)</i>	
Bases:	<code>trax.layers.combinators.Serial</code>
Wrapper that always initializes weights from a pre-trained checkpoint.	
__init__ <i>(*sublayers, init_checkpoint=None)</i>	Creates a partially initialized, unconnected layer instance.
Parameters:	<ul style="list-style-type: none">• n_in – Number of inputs expected by this layer.• n_out – Number of outputs promised by this layer.• name – Class-like name for this layer; for use when printing this layer.• sublayers_to_print – Sublayers to display when printing out this layer; By default (when None) we display all sublayers.
new_weights <i>(input_signature)</i>	

research.skipping_transformer