

tf.keras.layers.Conv2D



[TensorFlow](#)
[1 version](#)

[_\(/versions/r1.15/api_docs/python/tf/keras/layers/Conv2D\)](https://www.tensorflow.org/api_docs/python/tf/keras/layers/Conv2D)



[View](#)

[source_\(https://github.com/tensorflow/tensorflow\)](https://github.com/tensorflow/tensorflow)

[on L678](#)

[GitHub](#)

2D convolution layer (e.g. spatial convolution over images).

Inherits From: [Layer](#) (https://www.tensorflow.org/api_docs/python/tf/keras/layers/Layer), [Module](#) (https://www.tensorflow.org/api_docs/python/tf/Module)

[+ View aliases](#)

Main aliases

[tf.keras.layers.Convolution2D](#) (https://www.tensorflow.org/api_docs/python/tf/keras/layers/Conv2D)

Compat aliases for migration

See [Migration guide](#) (<https://www.tensorflow.org/guide/migrate>) for more details.

[tf.compat.v1.keras.layers.Conv2D](#) (https://www.tensorflow.org/api_docs/python/tf/keras/layers/Conv2D),

[tf.compat.v1.keras.layers.Convolution2D](#) (https://www.tensorflow.org/api_docs/python/tf/keras/layers/Conv2D)

```
tf.keras.layers.Conv2D(  
    filters, kernel_size, strides=(1, 1), padding='valid',  
    data_format=None, dilation_rate=(1, 1), groups=1, activation=None,  
    use_bias=True, kernel_initializer='glorot_uniform',  
    bias_initializer='zeros', kernel_regularizer=None,  
    bias_regularizer=None, activity_regularizer=None, kernel_constraint=None,  
    bias_constraint=None, **kwargs  
)
```

Used in the notebooks

Used in the guide

Used in the tutorials

Used in the guide	Used in the tutorials
<ul style="list-style-type: none"> • The Functional API (https://www.tensorflow.org/guide/keras/functional) • The Sequential model (https://www.tensorflow.org/guide/keras/sequential_model) • Migrate your TensorFlow 1 code to TensorFlow 2 (https://www.tensorflow.org/guide/migrate) • Eager execution (https://www.tensorflow.org/guide/eager) • Customize what happens in Model.fit (https://www.tensorflow.org/guide/keras/customizing_what_happens_in_fit) 	<ul style="list-style-type: none"> • Custom layers (https://www.tensorflow.org/tutorials/custom) • Image classification (https://www.tensorflow.org/tutorials/image) • Data augmentation (https://www.tensorflow.org/tutorials/image) • Intro to Autoencoders (https://www.tensorflow.org/tutorials/generative) • Pix2Pix (https://www.tensorflow.org/tutorials/image)

This layer creates a convolution kernel that is convolved with the layer input to produce a tensor of outputs. If `use_bias` is True, a bias vector is created and added to the outputs. Finally, if `activation` is not None, it is applied to the outputs as well.

When using this layer as the first layer in a model, provide the keyword argument `input_shape` (tuple of integers or None, does not include the sample axis), e.g. `input_shape=(128, 128, 3)` for 128x128 RGB pictures in `data_format="channels_last"`. You can use None when a dimension has variable size.

Examples: 

```
>>> # The inputs are 28x28 RGB images with `channels_last` and the batch
>>> # size is 4.
>>> input_shape = (4, 28, 28, 3)
>>> x = tf.random.normal(input_shape)
>>> y = tf.keras.layers.Conv2D(
... 2, 3, activation='relu', input_shape=input_shape[1:])(x)
>>> print(y.shape)
(4, 26, 26, 2)

>>> # With `dilation_rate` as 2.
>>> input_shape = (4, 28, 28, 3)
>>> x = tf.random.normal(input_shape)
>>> y = tf.keras.layers.Conv2D(
... 2, 3, activation='relu', dilation_rate=2, input_shape=input_shape[1:])(x)
>>> print(y.shape)
(4, 24, 24, 2)

>>> # With `padding` as "same".
>>> input_shape = (4, 28, 28, 3)
>>> x = tf.random.normal(input_shape)
>>> y = tf.keras.layers.Conv2D(
... 2, 3, activation='relu', padding="same", input_shape=input_shape[1:])(x)
```

```
>>> print(y.shape)
(4, 28, 28, 2)

>>> # With extended batch shape [4, 7]:
>>> input_shape = (4, 7, 28, 28, 3)
>>> x = tf.random.normal(input_shape)
>>> y = tf.keras.layers.Conv2D(
... 2, 3, activation='relu', input_shape=input_shape[2:])(x)
>>> print(y.shape)
(4, 7, 26, 26, 2)
```

Args

filters	Integer, the dimensionality of the output space (i.e. the number of output filters in the convolution).
kernel_size	An integer or tuple/list of 2 integers, specifying the height and width of the 2D convolution window. Can be a single integer to specify the same value for all spatial dimensions.
strides	An integer or tuple/list of 2 integers, specifying the strides of the convolution along the height and width. Can be a single integer to specify the same value for all spatial dimensions. Specifying any stride value != 1 is incompatible with specifying any dilation_rate value != 1.
padding	one of "valid" or "same" (case-insensitive). "valid" means no padding. "same" results in padding with zeros evenly to the left/right or up/down of the input such that output has the same height/width dimension as the input.
data_format	A string, one of channels_last (default) or channels_first . The ordering of the dimensions in the inputs. channels_last corresponds to inputs with shape (batch_size, height, width, channels) while channels_first corresponds to inputs with shape (batch_size, channels, height, width). It defaults to the image_data_format value found in your Keras config file at ~/.keras/keras.json. If you never set it, then it will be channels_last .
dilation_rate	an integer or tuple/list of 2 integers, specifying the dilation rate to use for dilated convolution. Can be a single integer to specify the same value for all spatial dimensions. Currently, specifying any dilation_rate value != 1 is incompatible with specifying any stride value != 1.
groups	A positive integer specifying the number of groups in which the input is split along the channel axis. Each group is convolved separately with filters / groups filters. The output is the concatenation of all the groups results along the channel axis. Input channels and filters must both be divisible by groups .
activation	Activation function to use. If you don't specify anything, no activation is applied (see <u>keras.activations</u> (https://www.tensorflow.org/api_docs/python/tf/keras/activations)).
use_bias	Boolean, whether the layer uses a bias vector.
kernel_initializer	Initializer for the kernel weights matrix (see <u>keras.initializers</u> (https://www.tensorflow.org/api_docs/python/tf/keras/initializers)). Defaults to 'glorot_uniform'.

bias_initializer	Initializer for the bias vector (see keras.initializers (https://www.tensorflow.org/api_docs/python/tf/keras/initializers)). Defaults to 'zeros'.
kernel_regularizer	Regularizer function applied to the kernel weights matrix (see keras.regularizers (https://www.tensorflow.org/api_docs/python/tf/keras/regularizers)).
bias_regularizer	Regularizer function applied to the bias vector (see keras.regularizers (https://www.tensorflow.org/api_docs/python/tf/keras/regularizers)).
activity_regularizer	Regularizer function applied to the output of the layer (its "activation") (see keras.regularizers (https://www.tensorflow.org/api_docs/python/tf/keras/regularizers)).
kernel_constraint	Constraint function applied to the kernel matrix (see keras.constraints (https://www.tensorflow.org/api_docs/python/tf/keras/constraints)).
bias_constraint	Constraint function applied to the bias vector (see keras.constraints (https://www.tensorflow.org/api_docs/python/tf/keras/constraints)).

Input shape:

4+D tensor with shape: `batch_shape + (channels, rows, cols)` if `data_format='channels_first'` or 4+D tensor with shape: `batch_shape + (rows, cols, channels)` if `data_format='channels_last'`.

Output shape:

4+D tensor with shape: `batch_shape + (filters, new_rows, new_cols)` if `data_format='channels_first'` or 4+D tensor with shape: `batch_shape + (new_rows, new_cols, filters)` if `data_format='channels_last'`. rows and cols values might have changed due to padding.

Returns

A tensor of rank 4+ representing `activation(conv2d(inputs, kernel) + bias)`.

Raises

ValueError if `padding` is "causal".

ValueError when both `strides > 1` and `dilation_rate > 1`.

Except as otherwise noted, the content of this page is licensed under the [Creative Commons Attribution 4.0 License](https://creativecommons.org/licenses/by/4.0/) (<https://creativecommons.org/licenses/by/4.0/>), and code samples are licensed under the [Apache 2.0 License](https://www.apache.org/licenses/LICENSE-2.0) (<https://www.apache.org/licenses/LICENSE-2.0>). For details, see the [Google Developers Site Policies](https://developers.google.com/site-policies) (<https://developers.google.com/site-policies>). Java is a registered trademark of Oracle and/or its affiliates. Some content is licensed under the [numpy license](https://numpy.org/doc/stable/license.html) (<https://numpy.org/doc/stable/license.html>).

Last updated 2021-05-14 UTC.