

# Docker

Docker (<https://docs.docker.com/install/>) uses *containers* to create virtual environments that isolate a TensorFlow installation from the rest of the system. TensorFlow programs are run *within* this virtual environment that can share resources with its host machine (access directories, use the GPU, connect to the Internet, etc.). The TensorFlow Docker images (<https://hub.docker.com/r/tensorflow/tensorflow/>) are tested for each release.

Docker is the easiest way to enable TensorFlow GPU support (<https://www.tensorflow.org/install/gpu>) on Linux since only the NVIDIA® GPU driver (<https://github.com/NVIDIA/nvidia-docker/wiki/Frequently-Asked-Questions#how-do-i-install-the-nvidia-driver>) is required on the *host* machine (the *NVIDIA® CUDA® Toolkit* does not need to be installed).

## TensorFlow Docker requirements

1. Install Docker (<https://docs.docker.com/install/>) on your local *host* machine.
2. For GPU support on Linux, install NVIDIA Docker support (<https://github.com/NVIDIA/nvidia-docker>).
  - Take note of your Docker version with `docker -v`. Versions **earlier than** 19.03 require `nvidia-docker2` and the `--runtime=nvidia` flag. On versions **including and after** 19.03, you will use the `nvidia-container-toolkit` package and the `--gpus all` flag. Both options are documented on the page linked above.

To run the **docker** command without **sudo**, create the **docker** group and add your user. For details, see the post-installation steps (<https://docs.docker.com/install/linux/linux-postinstall/>).

## Download a TensorFlow Docker image

The official TensorFlow Docker images are located in the tensorflow/tensorflow (<https://hub.docker.com/r/tensorflow/tensorflow/>) Docker Hub repository. Image releases are tagged (<https://hub.docker.com/r/tensorflow/tensorflow/tags/>) using the following format:

Tag	Description
<b>latest</b>	The latest release of TensorFlow CPU binary image. Default.
<b>nightly</b>	Nightly builds of the TensorFlow image. (Unstable.)
<b>version</b>	Specify the <i>version</i> of the TensorFlow binary image, for example: <i>2.1.0</i>
<b>devel</b>	Nightly builds of a TensorFlow <b>master</b> development environment. Includes TensorFlow source code.

Tag	Description
custom-op	Special experimental image for developing TF custom ops. More info <a href="https://github.com/tensorflow/custom-op">here</a> ( <a href="https://github.com/tensorflow/custom-op">https://github.com/tensorflow/custom-op</a> ).

Each base *tag* has variants that add or change functionality:

Tag Variants	Description
tag-gpu	The specified <i>tag</i> release with GPU support. ( <a href="#">See below</a> (#gpu_support))
tag-jupyter	The specified <i>tag</i> release with Jupyter (includes TensorFlow tutorial notebooks)

You can use multiple variants at once. For example, the following downloads TensorFlow release images to your machine:

```
$ docker pull tensorflow/tensorflow           # latest stable release
$ docker pull tensorflow/tensorflow:devel-gpu # nightly dev release w/ GPU support
$ docker pull tensorflow/tensorflow:latest-gpu-jupyter # latest release w/ GPU support and Jupyter
```

## Start a TensorFlow Docker container

To start a TensorFlow-configured container, use the following command form:

```
$ docker run [-it] [--rm] [-p hostPort:containerPort] tensorflow/tensorflow[:tag] [command]
```

For details, see the [docker run reference](https://docs.docker.com/engine/reference/run/) (<https://docs.docker.com/engine/reference/run/>).

### Examples using CPU-only images

Let's verify the TensorFlow installation using the `latest` tagged image. Docker downloads a new TensorFlow image the first time it is run:

```
$ docker run -it --rm tensorflow/tensorflow \
  python -c "import tensorflow as tf; print(tf.reduce_sum(tf.random.normal([1000, 1000])))"
```

**ss:** TensorFlow is now installed. Read the [tutorials](https://www.tensorflow.org/tutorials) (<https://www.tensorflow.org/tutorials>) to get started.

Let's demonstrate some more TensorFlow Docker recipes. Start a **bash** shell session within a TensorFlow-configured container:

```
$ docker run -it tensorflow/tensorflow bash
```

Within the container, you can start a `python` session and import TensorFlow.

To run a TensorFlow program developed on the *host* machine within a container, mount the host directory and change the container's working directory (`-v hostDir:containerDir -w workDir`):

```
$ docker run -it --rm -v $PWD:/tmp -w /tmp tensorflow/tensorflow python ./script.py
```

Permission issues can arise when files created within a container are exposed to the host. It's usually best to edit files on the host system.

Start a Jupyter Notebook (<https://jupyter.org/>) server using TensorFlow's nightly build:

```
$ docker run -it -p 8888:8888 tensorflow/tensorflow:nightly-jupyter
```

Follow the instructions and open the URL in your host web browser: `http://127.0.0.1:8888/?token=...`

## GPU support

Docker is the easiest way to run TensorFlow on a GPU since the *host* machine only requires the NVIDIA® driver (<https://github.com/NVIDIA/nvidia-docker/wiki/Frequently-Asked-Questions#how-do-i-install-the-nvidia-driver>) (the *NVIDIA® CUDA® Toolkit* is not required).

Install the Nvidia Container Toolkit (<https://github.com/NVIDIA/nvidia-docker/blob/master/README.md#quickstart>) to add NVIDIA® GPU support to Docker. `nvidia-container-runtime` is only available for Linux. See the nvidia-container-runtime platform support FAQ (<https://github.com/NVIDIA/nvidia-docker/wiki/Frequently-Asked-Questions#platform-support>) for details.

Check if a GPU is available:

```
$ lspci | grep -i nvidia
```

Verify your `nvidia-docker` installation:

```
$ docker run --gpus all --rm nvidia/cuda nvidia-smi
```

`nvidia-docker` v2 uses `--runtime=nvidia` instead of `--gpus all`. `nvidia-docker` v1 uses the `nvidia-docker` alias, rather than `time=nvidia` or `--gpus all` command line flags.

## Examples using GPU-enabled images

Download and run a GPU-enabled TensorFlow image (may take a few minutes):

```
$ docker run --gpus all -it --rm tensorflow/tensorflow:latest-gpu \  
python -c "import tensorflow as tf; print(tf.reduce_sum(tf.random.normal([1000, 1000])))"
```

It can take a while to set up the GPU-enabled image. If repeatedly running GPU-based scripts, you can use `docker exec` to reuse a container.

Use the latest TensorFlow GPU image to start a **bash** shell session in the container:

```
$ docker run --gpus all -it tensorflow/tensorflow:latest-gpu bash
```

**ss:** TensorFlow is now installed. Read the [tutorials](https://www.tensorflow.org/tutorials) (https://www.tensorflow.org/tutorials) to get started.

Except as otherwise noted, the content of this page is licensed under the [Creative Commons Attribution 4.0 License](https://creativecommons.org/licenses/by/4.0/) (https://creativecommons.org/licenses/by/4.0/), and code samples are licensed under the [Apache 2.0 License](https://www.apache.org/licenses/LICENSE-2.0) (https://www.apache.org/licenses/LICENSE-2.0). For details, see the [Google Developers Site Policies](https://developers.google.com/site-policies) (https://developers.google.com/site-policies). Java is a registered trademark of Oracle and/or its affiliates.

Last updated 2021-01-28 UTC.