

# Capstone Project

April 17, 2021

## 1 Capstone Project

### 1.1 Image classifier for the SVHN dataset

#### 1.1.1 Instructions

In this notebook, you will create a neural network that classifies real-world images digits. You will use concepts from throughout this course in building, training, testing, validating and saving your Tensorflow classifier model.

This project is peer-assessed. Within this notebook you will find instructions in each section for how to complete the project. Pay close attention to the instructions as the peer review will be carried out according to a grading rubric that checks key parts of the project instructions. Feel free to add extra cells into the notebook as required.

#### 1.1.2 How to submit

When you have completed the Capstone project notebook, you will submit a pdf of the notebook for peer review. First ensure that the notebook has been fully executed from beginning to end, and all of the cell outputs are visible. This is important, as the grading rubric depends on the reviewer being able to view the outputs of your notebook. Save the notebook as a pdf (File -> Download as -> PDF via LaTeX). You should then submit this pdf for review.

#### 1.1.3 Let's get started!

We'll start by running some imports, and loading the dataset. For this project you are free to make further imports throughout the notebook as you wish.

```
In [1]: import tensorflow as tf
        from scipy.io import loadmat

        import numpy as np
        import matplotlib.pyplot as plt
        import math

        from tensorflow.keras.models import Sequential
        from tensorflow.keras.layers import Dense, Flatten, Conv2D, MaxPooling2D, BatchNormaliz
        from tensorflow.keras import regularizers
        from tensorflow.keras.callbacks import ModelCheckpoint, ReduceLROnPlateau, EarlyStoppin
```



For the capstone project, you will use the [SVHN dataset](#). This is an image dataset of over 600,000 digit images in all, and is a harder dataset than MNIST as the numbers appear in the context of natural scene images. SVHN is obtained from house numbers in Google Street View images.

- Y. Netzer, T. Wang, A. Coates, A. Bissacco, B. Wu and A. Y. Ng. “Reading Digits in Natural Images with Unsupervised Feature Learning”. NIPS Workshop on Deep Learning and Unsupervised Feature Learning, 2011.

Your goal is to develop an end-to-end workflow for building, training, validating, evaluating and saving a neural network that classifies a real-world image into one of ten classes.

In [2]: # Run this cell to load the dataset

```
train = loadmat('data/train_32x32.mat')
test = loadmat('data/test_32x32.mat')
```

Both `train` and `test` are dictionaries with keys `X` and `y` for the input images and labels respectively.

## 1.2 1. Inspect and preprocess the dataset

- Extract the training and testing images and labels separately from the train and test dictionaries loaded for you.
- Select a random sample of images and corresponding labels from the dataset (at least 10), and display them in a figure.
- Convert the training and test images to grayscale by taking the average across all colour channels for each pixel. *Hint: retain the channel dimension, which will now have size 1.*
- Select a random sample of the grayscale images and corresponding labels from the dataset (at least 10), and display them in a figure.

In [3]: """

```
print(train.keys()) #dict_keys(['__header__', '__version__', '__globals__', 'X', 'y'])
"""
```

```

train_features = train["X"]
train_labels    = train["y"]

test_features   = test["X"]
test_labels     = test["y"]

"""
print(train_features.shape) #(32, 32, 3, 73257)
print(train_labels.shape) #(73257, 1)
"""

train_features = np.moveaxis(train_features,3,0) #move last axis on first place for it
test_features  = np.moveaxis(test_features,3,0) #move last axis on first place for itera
"""
print(train_features.shape) #(73257, 32, 32, 3)
print(train_labels.shape) #(73257, 1)
"""

```

Out[3]: '\nprint(train\_features.shape) #(73257, 32, 32, 3)\nprint(train\_labels.shape) #(73257,

```

In [4]: def plot2square(features, labels, N):
        #pick number of samples to show
        n_samples = 9
        random_inx = np.random.choice(features.shape[0], n_samples)


        #random samples with their labels
        random_test_images = features[random_inx] #move last axis on first place for itera
        random_test_labels = labels[random_inx]

        #make square like plot field
        x_len = int(math.sqrt(n_samples))
        y_len = math.ceil(n_samples/x_len)

        #prepare fig
        fig, axes = plt.subplots(x_len, y_len, figsize=(8, 8))
        fig.subplots_adjust(hspace=0.4, wspace= 0.4)
        for i, (image, label) in enumerate(zip(random_test_images, random_test_labels)):
            axes[int(i / y_len), i % y_len].imshow(np.squeeze(image))
            axes[int(i / y_len), i % y_len].set_title(f"Label: {label}")

```


In [5]: plot2square(train\_features, train\_labels, 9)



```
In [6]: #matlab's (NTSC/PAL) implementation:
def rgb2gray(features):
    RGB_weights = (0.2989, 0.5870, 0.1140)
    return np.average(features, axis=3, weights=RGB_weights).reshape((features.shape[0],
    features.shape[1], features.shape[2]))

train_features_gray = rgb2gray(train_features)

#image test
plot2square(train_features_gray, train_labels, 9)
```



### 1.3 2. MLP neural network classifier

- Build an MLP classifier model using the Sequential API. Your model should use only Flatten and Dense layers, with the final layer having a 10-way softmax output.
- You should design and build the model yourself. Feel free to experiment with different MLP architectures. *Hint: to achieve a reasonable accuracy you won't need to use more than 4 or 5 layers.*
- Print out the model summary (using the summary() method)
- Compile and train the model (we recommend a maximum of 30 epochs), making use of both training and validation sets during the training run.
- Your model should track at least one appropriate metric, and use at least two callbacks during training, one of which should be a ModelCheckpoint callback.
- As a guide, you should aim to achieve a final categorical cross entropy training loss of less than 1.0 (the validation loss might be higher).

- Plot the learning curves for loss vs epoch and accuracy vs epoch for both training and validation sets.
- Compute and display the loss and accuracy of the trained model on the test set.

```
In [7]: def get_model(shape, wd1, wd2):
        return Sequential([
            Flatten(input_shape=shape),
            Dense(512, activation='relu', kernel_regularizer=regularizers.l1_l2(wd1, wd2),
            Dense(256, activation='relu', kernel_regularizer=regularizers.l1_l2(wd1, wd2),
            Dense(128, activation='relu', kernel_regularizer=regularizers.l1_l2(wd1, wd2),
            Dense(64, activation='relu', kernel_regularizer=regularizers.l1_l2(wd1, wd2),
            Dense(10, activation='softmax')
        ])

        mlp_model = get_model(train_features_gray[0].shape, 1e-5, 1e-3)
        mlp_model.summary()
```

Model: "sequential"

Layer (type)	Output Shape	Param #
flatten (Flatten)	(None, 1024)	0
dense (Dense)	(None, 512)	524800
dense_1 (Dense)	(None, 256)	131328
dense_2 (Dense)	(None, 128)	32896
dense_3 (Dense)	(None, 64)	8256
dense_4 (Dense)	(None, 10)	650
Total params: 697,930		
Trainable params: 697,930		
Non-trainable params: 0		

```
In [8]: opt = tf.keras.optimizers.Adam(learning_rate=0.0025)
        mlp_model.compile(optimizer=opt, loss='sparse_categorical_crossentropy', metrics=['acc'])

In [9]: reduce_lr = ReduceLROnPlateau(monitor='val_loss', factor=0.2,
                                       patience=2, min_lr=0.0001, verbose=2)
        early_stopping = EarlyStopping(monitor='accuracy', patience=3)
        save_best_mlp = ModelCheckpoint(filepath='./checkpoints/mlp/model_best_epoch', monitor='val_loss')

        history = mlp_model.fit(train_features_gray, train_labels-1, epochs=30,
                                validation_split=0.2, batch_size=128, verbose=2,
```

```
callbacks=[reduce_lr, early_stopping, save_best_mlp]
)
```

Train on 58605 samples, validate on 14652 samples

Epoch 1/30

Epoch 00001: loss improved from inf to 33.53691, saving model to ./checkpoints/mlp/model\_best\_58605/58605 - 41s - loss: 33.5369 - accuracy: 0.1514 - val\_loss: 5.8613 - val\_accuracy: 0.1292  
Epoch 2/30

Epoch 00002: loss improved from 33.53691 to 3.76455, saving model to ./checkpoints/mlp/model\_best\_58605/58605 - 37s - loss: 3.7645 - accuracy: 0.2858 - val\_loss: 3.5614 - val\_accuracy: 0.3615  
Epoch 3/30

Epoch 00003: loss improved from 3.76455 to 3.13827, saving model to ./checkpoints/mlp/model\_best\_58605/58605 - 37s - loss: 3.1383 - accuracy: 0.4746 - val\_loss: 3.2197 - val\_accuracy: 0.4808  
Epoch 4/30

Epoch 00004: loss improved from 3.13827 to 2.87923, saving model to ./checkpoints/mlp/model\_best\_58605/58605 - 37s - loss: 2.8792 - accuracy: 0.5454 - val\_loss: 2.6728 - val\_accuracy: 0.6068  
Epoch 5/30

Epoch 00005: loss improved from 2.87923 to 2.81623, saving model to ./checkpoints/mlp/model\_best\_58605/58605 - 37s - loss: 2.8162 - accuracy: 0.5428 - val\_loss: 2.6388 - val\_accuracy: 0.5775  
Epoch 6/30

Epoch 00006: loss improved from 2.81623 to 2.64597, saving model to ./checkpoints/mlp/model\_best\_58605/58605 - 37s - loss: 2.6460 - accuracy: 0.5638 - val\_loss: 2.4996 - val\_accuracy: 0.5937  
Epoch 7/30

Epoch 00007: loss improved from 2.64597 to 2.41388, saving model to ./checkpoints/mlp/model\_best\_58605/58605 - 37s - loss: 2.4139 - accuracy: 0.6013 - val\_loss: 2.3858 - val\_accuracy: 0.5868  
Epoch 8/30

Epoch 00008: loss improved from 2.41388 to 2.25281, saving model to ./checkpoints/mlp/model\_best\_58605/58605 - 37s - loss: 2.2528 - accuracy: 0.6241 - val\_loss: 2.1773 - val\_accuracy: 0.6356  
Epoch 9/30

Epoch 00009: loss improved from 2.25281 to 2.14230, saving model to ./checkpoints/mlp/model\_best\_58605/58605 - 37s - loss: 2.1423 - accuracy: 0.6287 - val\_loss: 2.0962 - val\_accuracy: 0.6225  
Epoch 10/30

Epoch 00010: loss improved from 2.14230 to 2.02571, saving model to ./checkpoints/mlp/model\_best\_58605/58605 - 38s - loss: 2.0257 - accuracy: 0.6407 - val\_loss: 1.9734 - val\_accuracy: 0.6372  
Epoch 11/30

Epoch 00011: loss improved from 2.02571 to 1.92652, saving model to ./checkpoints/mlp/model\_best\_58605/58605 - 37s - loss: 1.9265 - accuracy: 0.6420 - val\_loss: 1.8744 - val\_accuracy: 0.6353

Epoch 12/30

Epoch 00012: loss improved from 1.92652 to 1.76827, saving model to ./checkpoints/mlp/model\_best\_58605/58605 - 37s - loss: 1.7683 - accuracy: 0.6618 - val\_loss: 1.7071 - val\_accuracy: 0.6695  
Epoch 13/30

Epoch 00013: loss improved from 1.76827 to 1.69317, saving model to ./checkpoints/mlp/model\_best\_58605/58605 - 36s - loss: 1.6932 - accuracy: 0.6573 - val\_loss: 1.6439 - val\_accuracy: 0.6637  
Epoch 14/30

Epoch 00014: loss improved from 1.69317 to 1.59199, saving model to ./checkpoints/mlp/model\_best\_58605/58605 - 36s - loss: 1.5920 - accuracy: 0.6642 - val\_loss: 1.5367 - val\_accuracy: 0.6740  
Epoch 15/30

Epoch 00015: loss improved from 1.59199 to 1.50660, saving model to ./checkpoints/mlp/model\_best\_58605/58605 - 36s - loss: 1.5066 - accuracy: 0.6702 - val\_loss: 1.5764 - val\_accuracy: 0.6329  
Epoch 16/30

Epoch 00016: loss improved from 1.50660 to 1.45933, saving model to ./checkpoints/mlp/model\_best\_58605/58605 - 36s - loss: 1.4593 - accuracy: 0.6691 - val\_loss: 1.4280 - val\_accuracy: 0.6752  
Epoch 17/30

Epoch 00017: loss improved from 1.45933 to 1.43822, saving model to ./checkpoints/mlp/model\_best\_58605/58605 - 36s - loss: 1.4382 - accuracy: 0.6652 - val\_loss: 1.5226 - val\_accuracy: 0.6374  
Epoch 18/30

Epoch 00018: loss improved from 1.43822 to 1.38218, saving model to ./checkpoints/mlp/model\_best\_58605/58605 - 37s - loss: 1.3822 - accuracy: 0.6733 - val\_loss: 1.3928 - val\_accuracy: 0.6683  
Epoch 19/30

Epoch 00019: loss improved from 1.38218 to 1.34554, saving model to ./checkpoints/mlp/model\_best\_58605/58605 - 37s - loss: 1.3455 - accuracy: 0.6795 - val\_loss: 1.3878 - val\_accuracy: 0.6600  
Epoch 20/30

Epoch 00020: loss improved from 1.34554 to 1.30072, saving model to ./checkpoints/mlp/model\_best\_58605/58605 - 37s - loss: 1.3007 - accuracy: 0.6819 - val\_loss: 1.3114 - val\_accuracy: 0.6833  
Epoch 21/30

Epoch 00021: loss improved from 1.30072 to 1.28560, saving model to ./checkpoints/mlp/model\_best\_58605/58605 - 37s - loss: 1.2856 - accuracy: 0.6863 - val\_loss: 1.3636 - val\_accuracy: 0.6567  
Epoch 22/30

Epoch 00022: loss improved from 1.28560 to 1.25140, saving model to ./checkpoints/mlp/model\_best\_58605/58605 - 36s - loss: 1.2514 - accuracy: 0.6882 - val\_loss: 1.2847 - val\_accuracy: 0.6860  
Epoch 23/30

Epoch 00023: loss did not improve from 1.25140  
58605/58605 - 36s - loss: 1.2701 - accuracy: 0.6863 - val\_loss: 1.2037 - val\_accuracy: 0.7069



Epoch 24/30

Epoch 00024: loss improved from 1.25140 to 1.24867, saving model to ./checkpoints/mlp/model\_best.pth  
58605/58605 - 36s - loss: 1.2487 - accuracy: 0.6919 - val\_loss: 1.2987 - val\_accuracy: 0.6775  
Epoch 25/30

Epoch 00025: ReduceLROnPlateau reducing learning rate to 0.0004999999888241291.

Epoch 00025: loss did not improve from 1.24867  
58605/58605 - 36s - loss: 1.2492 - accuracy: 0.6907 - val\_loss: 1.2344 - val\_accuracy: 0.6982  
Epoch 26/30

Epoch 00026: loss improved from 1.24867 to 1.06543, saving model to ./checkpoints/mlp/model\_best.pth  
58605/58605 - 36s - loss: 1.0654 - accuracy: 0.7393 - val\_loss: 1.0745 - val\_accuracy: 0.7314  
Epoch 27/30

Epoch 00027: loss improved from 1.06543 to 0.99996, saving model to ./checkpoints/mlp/model\_best.pth  
58605/58605 - 36s - loss: 1.0000 - accuracy: 0.7457 - val\_loss: 1.0362 - val\_accuracy: 0.7273  
Epoch 28/30

Epoch 00028: loss improved from 0.99996 to 0.96631, saving model to ./checkpoints/mlp/model\_best.pth  
58605/58605 - 36s - loss: 0.9663 - accuracy: 0.7480 - val\_loss: 1.0228 - val\_accuracy: 0.7316  
Epoch 29/30

Epoch 00029: loss improved from 0.96631 to 0.94263, saving model to ./checkpoints/mlp/model\_best.pth  
58605/58605 - 36s - loss: 0.9426 - accuracy: 0.7524 - val\_loss: 1.0036 - val\_accuracy: 0.7358  
Epoch 30/30

Epoch 00030: loss improved from 0.94263 to 0.93334, saving model to ./checkpoints/mlp/model\_best.pth  
58605/58605 - 36s - loss: 0.9333 - accuracy: 0.7483 - val\_loss: 1.0017 - val\_accuracy: 0.7275

```
In [10]: mlp_model.evaluate(rgb2gray(test_features), test_labels-1, verbose=2)
```

26032/1 - 13s - loss: 1.2140 - accuracy: 0.7071

```
Out[10]: [1.0998752706242838, 0.7070529]
```

```
In [11]: def plot_history(history):  
    fig = plt.figure(figsize=(12, 5))  
  
    fig.add_subplot(121)  
  
    plt.plot(history.history['accuracy'])  
    plt.plot(history.history['val_accuracy'])  
    plt.title('Accuracy vs. epochs')  
    plt.ylabel('Accuracy')  
    plt.xlabel('Epoch')
```

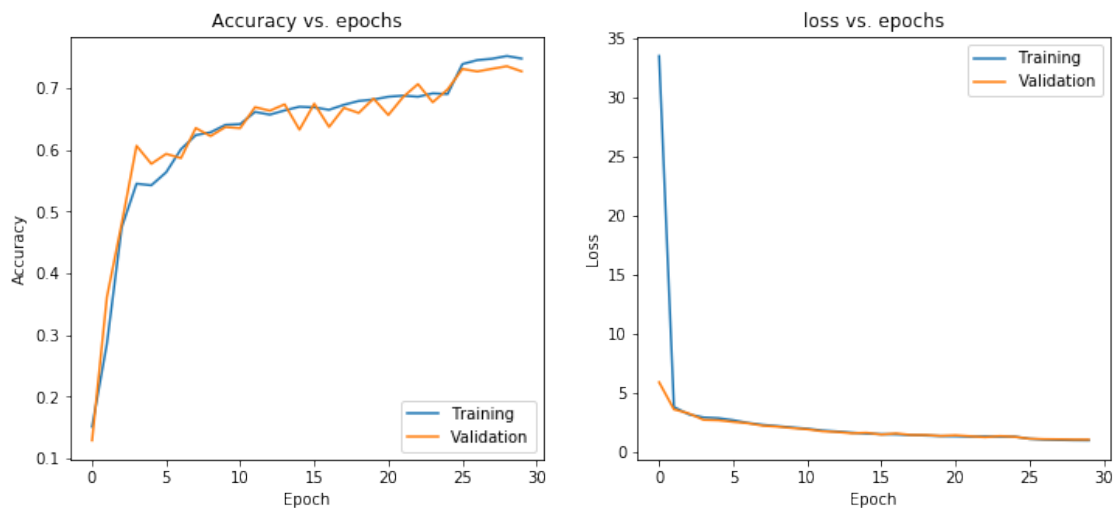
```
plt.legend(['Training', 'Validation'], loc='lower right')

fig.add_subplot(122)

plt.plot(history.history['loss'])
plt.plot(history.history['val_loss'])
plt.title('loss vs. epochs')
plt.ylabel('Loss')
plt.xlabel('Epoch')
plt.legend(['Training', 'Validation'], loc='upper right')

plt.show()
```

In [12]: plot\_history(history)



### 1.4 3. CNN neural network classifier

- Build a CNN classifier model using the Sequential API. Your model should use the Conv2D, MaxPool2D, BatchNormalization, Flatten, Dense and Dropout layers. The final layer should again have a 10-way softmax output.
- You should design and build the model yourself. Feel free to experiment with different CNN architectures. *Hint: to achieve a reasonable accuracy you won't need to use more than 2 or 3 convolutional layers and 2 fully connected layers.*
- The CNN model should use fewer trainable parameters than your MLP model.
- Compile and train the model (we recommend a maximum of 30 epochs), making use of both training and validation sets during the training run.
- Your model should track at least one appropriate metric, and use at least two callbacks during training, one of which should be a ModelCheckpoint callback.
- You should aim to beat the MLP model performance with fewer parameters!

- Plot the learning curves for loss vs epoch and accuracy vs epoch for both training and validation sets.
- Compute and display the loss and accuracy of the trained model on the test set.

```
In [13]: def get_model(shape, dropout_rate):
        return Sequential([
            Conv2D(32, (4,4), activation="relu", input_shape=(32,32,1), padding="same"),
            MaxPooling2D((4,4)),
            BatchNormalization(),
            Flatten(),
            Dense(256, activation="relu", kernel_initializer='he_uniform'),
            Dropout(dropout_rate),
            Dense(128, activation="relu", kernel_initializer='he_uniform'),
            Dense(10, activation="softmax")
        ])
```

```
In [14]: cnn_model = get_model(train_features_gray[0].shape, 0.3)
        cnn_model.summary()
```

Model: "sequential\_1"

Layer (type)	Output Shape	Param #
conv2d (Conv2D)	(None, 32, 32, 32)	544
max_pooling2d (MaxPooling2D)	(None, 8, 8, 32)	0
batch_normalization (Batch Normalization)	(None, 8, 8, 32)	128
flatten_1 (Flatten)	(None, 2048)	0
dense_5 (Dense)	(None, 256)	524544
dropout (Dropout)	(None, 256)	0
dense_6 (Dense)	(None, 128)	32896
dense_7 (Dense)	(None, 10)	1290
Total params: 559,402		
Trainable params: 559,338		
Non-trainable params: 64		

```
In [15]: save_best_cnn = ModelCheckpoint(filepath='./checkpoints/cnn/model_best_epoch', monitor=
        cnn_model.compile(optimizer="adam", loss="sparse_categorical_crossentropy", metrics=[
```

```
In [16]: history = cnn_model.fit(train_features_gray, train_labels-1, epochs=30,
                                validation_split=0.2, batch_size=128, verbose=2,
                                callbacks=[early_stopping, save_best_cnn]
                                )
```

Train on 58605 samples, validate on 14652 samples

Epoch 1/30

58605/58605 - 154s - loss: 1.2151 - accuracy: 0.5998 - val\_loss: 0.6731 - val\_accuracy: 0.7875

Epoch 2/30

58605/58605 - 149s - loss: 0.6820 - accuracy: 0.7875 - val\_loss: 0.6115 - val\_accuracy: 0.8143

Epoch 3/30

58605/58605 - 149s - loss: 0.5849 - accuracy: 0.8200 - val\_loss: 0.4933 - val\_accuracy: 0.8529

Epoch 4/30

58605/58605 - 148s - loss: 0.5347 - accuracy: 0.8344 - val\_loss: 0.5413 - val\_accuracy: 0.8339

Epoch 5/30

58605/58605 - 146s - loss: 0.4916 - accuracy: 0.8494 - val\_loss: 0.5197 - val\_accuracy: 0.8479

Epoch 6/30

58605/58605 - 146s - loss: 0.4648 - accuracy: 0.8550 - val\_loss: 0.4894 - val\_accuracy: 0.8546

Epoch 7/30

58605/58605 - 147s - loss: 0.4419 - accuracy: 0.8623 - val\_loss: 0.4531 - val\_accuracy: 0.8647

Epoch 8/30

58605/58605 - 147s - loss: 0.4185 - accuracy: 0.8693 - val\_loss: 0.4600 - val\_accuracy: 0.8675

Epoch 9/30

58605/58605 - 148s - loss: 0.4030 - accuracy: 0.8750 - val\_loss: 0.4520 - val\_accuracy: 0.8733

Epoch 10/30

58605/58605 - 148s - loss: 0.3897 - accuracy: 0.8781 - val\_loss: 0.5633 - val\_accuracy: 0.8333

Epoch 11/30

58605/58605 - 149s - loss: 0.3719 - accuracy: 0.8848 - val\_loss: 0.4203 - val\_accuracy: 0.8812

Epoch 12/30

58605/58605 - 152s - loss: 0.3598 - accuracy: 0.8874 - val\_loss: 0.4222 - val\_accuracy: 0.8799

Epoch 13/30

58605/58605 - 150s - loss: 0.3522 - accuracy: 0.8906 - val\_loss: 0.5682 - val\_accuracy: 0.8481

Epoch 14/30

58605/58605 - 149s - loss: 0.3372 - accuracy: 0.8956 - val\_loss: 0.4205 - val\_accuracy: 0.8831

Epoch 15/30

58605/58605 - 148s - loss: 0.3270 - accuracy: 0.8969 - val\_loss: 0.4026 - val\_accuracy: 0.8883

Epoch 16/30

58605/58605 - 145s - loss: 0.3134 - accuracy: 0.9021 - val\_loss: 0.4206 - val\_accuracy: 0.8825

Epoch 17/30

58605/58605 - 147s - loss: 0.3114 - accuracy: 0.9026 - val\_loss: 0.4322 - val\_accuracy: 0.8758

Epoch 18/30

58605/58605 - 147s - loss: 0.2978 - accuracy: 0.9055 - val\_loss: 0.4213 - val\_accuracy: 0.8810

Epoch 19/30

58605/58605 - 146s - loss: 0.2960 - accuracy: 0.9060 - val\_loss: 0.4536 - val\_accuracy: 0.8769

Epoch 20/30

58605/58605 - 153s - loss: 0.2838 - accuracy: 0.9102 - val\_loss: 0.4214 - val\_accuracy: 0.8833

Epoch 21/30


58605/58605 - 157s - loss: 0.2830 - accuracy: 0.9106 - val\_loss: 0.4107 - val\_accuracy: 0.8841

```

Epoch 22/30
58605/58605 - 148s - loss: 0.2748 - accuracy: 0.9127 - val_loss: 0.4188 - val_accuracy: 0.8877
Epoch 23/30
58605/58605 - 146s - loss: 0.2668 - accuracy: 0.9154 - val_loss: 0.4258 - val_accuracy: 0.8866
Epoch 24/30
58605/58605 - 149s - loss: 0.2638 - accuracy: 0.9172 - val_loss: 0.4116 - val_accuracy: 0.8917
Epoch 25/30
58605/58605 - 149s - loss: 0.2551 - accuracy: 0.9189 - val_loss: 0.4495 - val_accuracy: 0.8778
Epoch 26/30
58605/58605 - 147s - loss: 0.2445 - accuracy: 0.9221 - val_loss: 0.4093 - val_accuracy: 0.8920
Epoch 27/30
58605/58605 - 148s - loss: 0.2454 - accuracy: 0.9215 - val_loss: 0.4073 - val_accuracy: 0.8939
Epoch 28/30
58605/58605 - 147s - loss: 0.2439 - accuracy: 0.9211 - val_loss: 0.4579 - val_accuracy: 0.8811
Epoch 29/30
58605/58605 - 150s - loss: 0.2341 - accuracy: 0.9248 - val_loss: 0.4183 - val_accuracy: 0.8933
Epoch 30/30
58605/58605 - 148s - loss: 0.2259 - accuracy: 0.9273 - val_loss: 0.4136 - val_accuracy: 0.8952

```

In [17]: `plot_history(history)`



In [ ]:

## 1.5 4. Get model predictions

- Load the best weights for the MLP and CNN models that you saved during the training run.
- Randomly select 5 images and corresponding labels from the test set and display the images with their labels.
- Alongside the image and label, show each model's predictive distribution as a bar chart, and the final model prediction given by the label with maximum probability.

```

In [18]: mlp_model.load_weights(tf.train.latest_checkpoint('./checkpoints/mlp/'))
         cnn_model.load_weights(tf.train.latest_checkpoint('./checkpoints/cnn/'))

Out[18]: <tensorflow.python.training.tracking.util.CheckpointLoadStatus at 0x7f4ed8375b70>

In [19]: images = []
         labels = []
         predictions_mlp = []
         predictions_cnn = []
         n_tests = 10

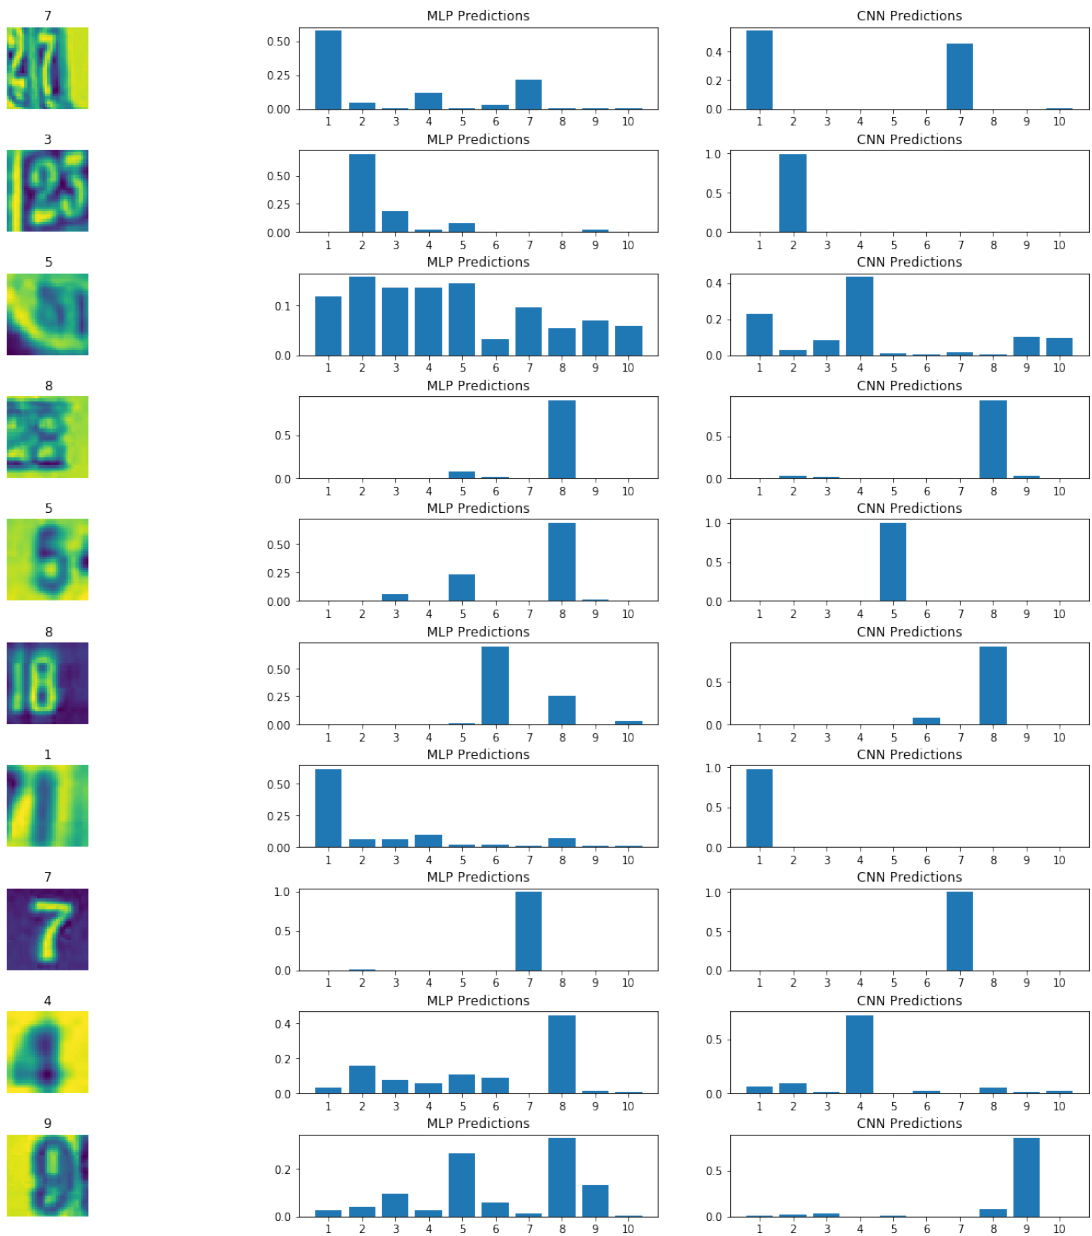
         for i in range(0, n_tests):
             rdm_num = np.random.randint(rgb2gray(test_features).shape[0], size=1)
             img = rgb2gray(test_features)[rdm_num, :, :, :]
             images.append(img)
             labels.append(str(test_labels[rdm_num]).strip('[]'))
             predictions_mlp.append(mlp_model.predict(img))
             predictions_cnn.append(cnn_model.predict(img))

In [20]: fig, axs = plt.subplots(nrows=n_tests, ncols=3, figsize=(20,20))
         fig.subplots_adjust(hspace=0.5, wspace=0.2)

         x = np.arange(1,11)

         for i, (image, label, mlp_prediction, cnn_prediction) in enumerate(zip(images, labels
             axs[i, 0].imshow(np.squeeze(image))
             axs[i, 0].set_title(label)
             axs[i, 0].axis('off')
             axs[i, 1].bar(x, mlp_prediction.reshape(10,))
             axs[i, 1].set_title('MLP Predictions')
             axs[i, 1].set_xticks(x)
             axs[i, 2].bar(x, cnn_prediction.reshape(10,))
             axs[i, 2].set_title('CNN Predictions')
             axs[i, 2].set_xticks(x)

```



In [ ]:

In [ ]:

In [ ]: